

# Terraform Training



# Pre-requisites (Tools to install)

- Adobe Acrobat Reader or equivalent
- [Docker Desktop](#)
- [Terraform](#)
- [Azure CLI](#) or ["Az" PowerShell module](#)
- Visual Studio Code
- VS Code Plugins:
  - [Azure Terraform](#)
  - [HashiCorp Terraform](#)
  - [JSON Formatter](#)

# Agenda


- | Introduction to IaC
- | Getting started with Terraform
- | Terraform in details
- | Collaboration & pipeline
- | Testing & deployment






# 1. Introduction to IaC


# What is infrastructure ?



# What is infrastructure ?



# What is infrastructure ?



# DevOps ?

- | Culture
- | **Automation**
- | Measurement
- | Sharing





# What is Infrastructure as Code (IaC) ?

## | Definition


- | Process of automating infrastructure changes in a codified manner to achieve scalability, reliability, security and sustainability.

## | Challenges

## | Practices

# Why use infrastructure as code ?


- | Change management
- | Return on time investment
- | Knowledge sharing
- | Security



# Principles


## Reproducible

Use a configuration to create a new environment with the same specification




## Idempotent

Repeatedly run the automation on the same code and yield the same result




## Composable

Create an infrastructure system using a set of building blocks

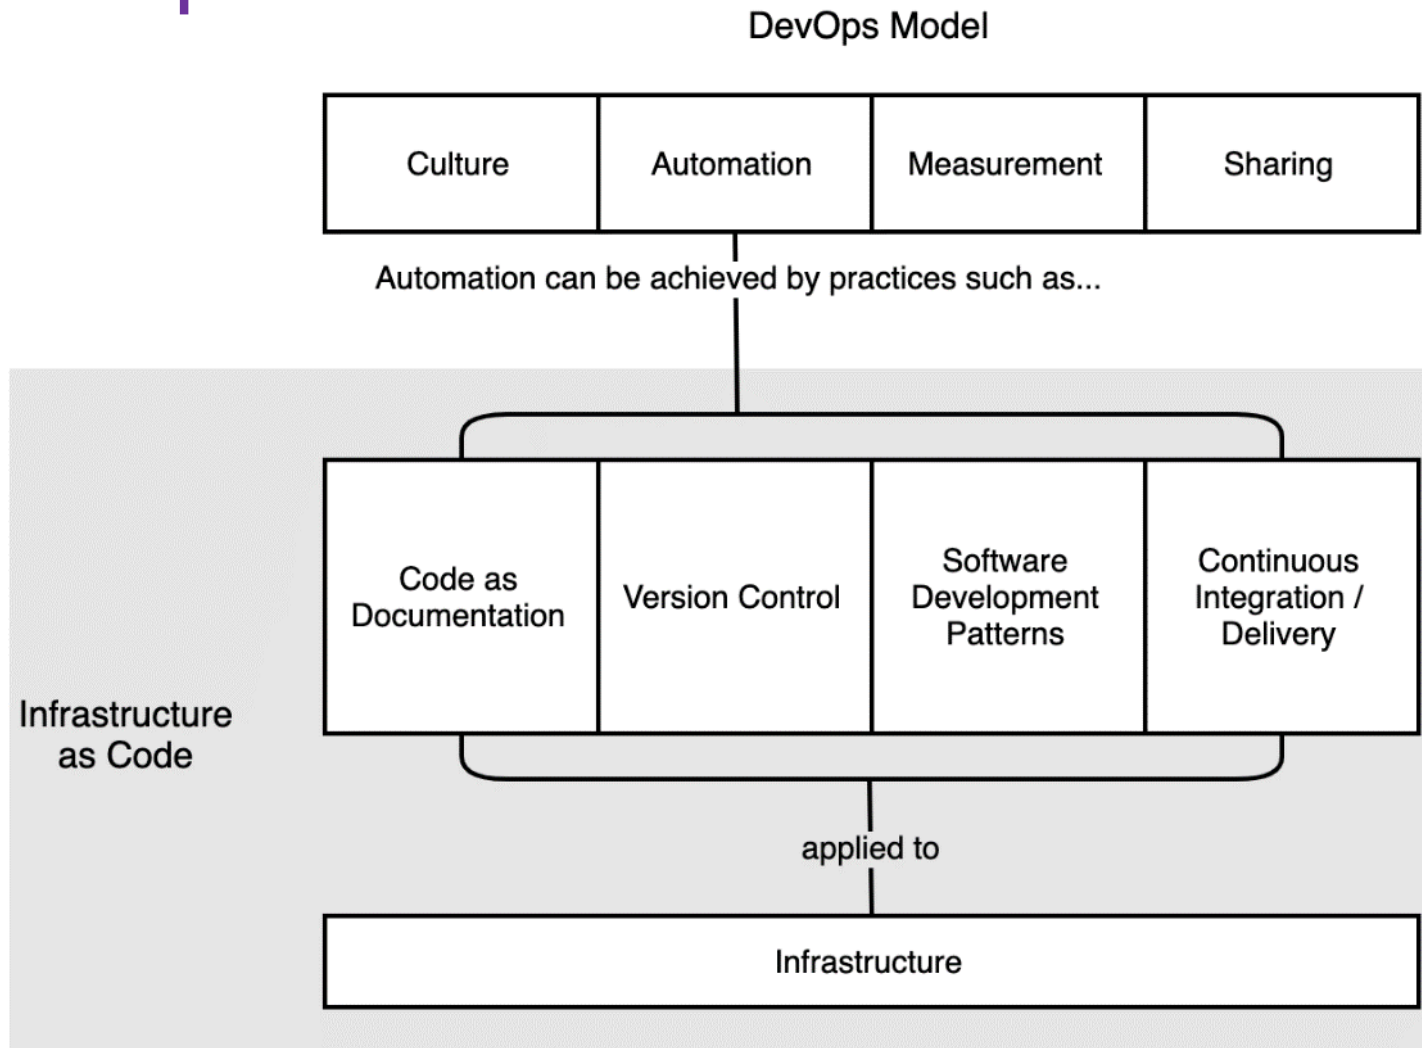


## Evolvable

Change part of the system with minimal disruption to other infrastructure



# Quick recap



# Tools

## Provisioning

Tool	Provider
Azure Resource Manager	Microsoft Azure
HashiCorp Terraform	Various
Pulumi SDK	Various
AWS Cloud Development Kit	Amazon Web Services
Kubernetes Manifests	Kubernetes (Container Orchestrator)

## Configuration Management

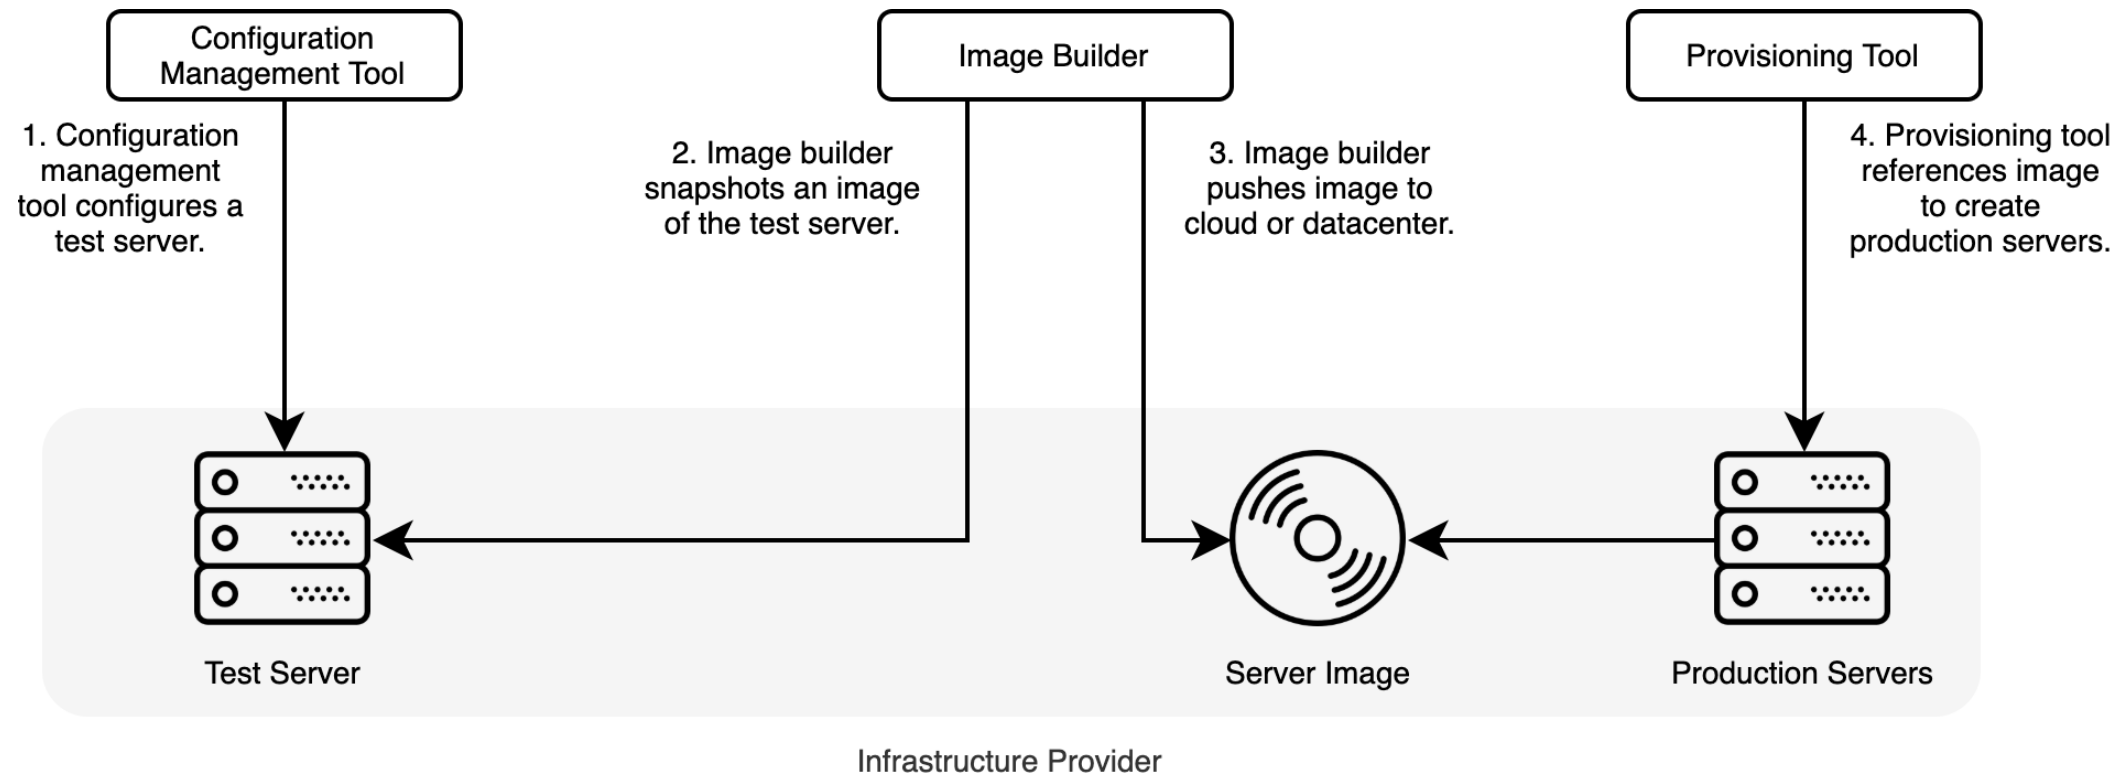
| Chef, Puppet, Ansible,  
SaltStack, CFEngine, ...

## Image building

Tool	Runtime Environment	Build Target
HashiCorp Packer	Containers & Servers	Various
Docker	Containers	Container Registries
Amazon EC2 Image Builder	Servers	Amazon Web Services
Azure VM Image Builder	Servers	Microsoft Azure



# Deployment process in a nutshell



# Quiz - Solution

- | Infrastructure can be software, platform, or hardware that delivers or deploys applications to production.
- | Infrastructure as code is a DevOps practice of automating infrastructure to achieve reliability, scalability, and security.
- | The principles of infrastructure as code are reproducibility, idempotency, composability, and evolvability.
- | By following the principles of infrastructure as code, you can improve change management processes, lower time spent on fixing failed systems in the long term, better share knowledge and context, and build security into your infrastructure.
- | Infrastructure as code tools include three types: provisioning tools, configuration management tools, and image builders.

# IaC - Summary

- | Infrastructure can be software, platform, or hardware that delivers or deploys applications to production.
- | Infrastructure as code is a DevOps practice of automating infrastructure to achieve reliability, scalability, and security.
- | The principles of infrastructure as code are reproducibility, idempotency, composability, and evolvability.
- | By following the principles of infrastructure as code, you can improve change management processes, lower time spent on fixing failed systems in the long term, better share knowledge and context, and build security into your infrastructure.
- | Infrastructure as code tools include three types: provisioning tools, configuration management tools, and image builders.

## 2. Getting Started with Terraform



# Content



Generating and applying execution plans



Analysing when function hooks are triggered by Terraform



Utilizing the Local provider to create and manage files



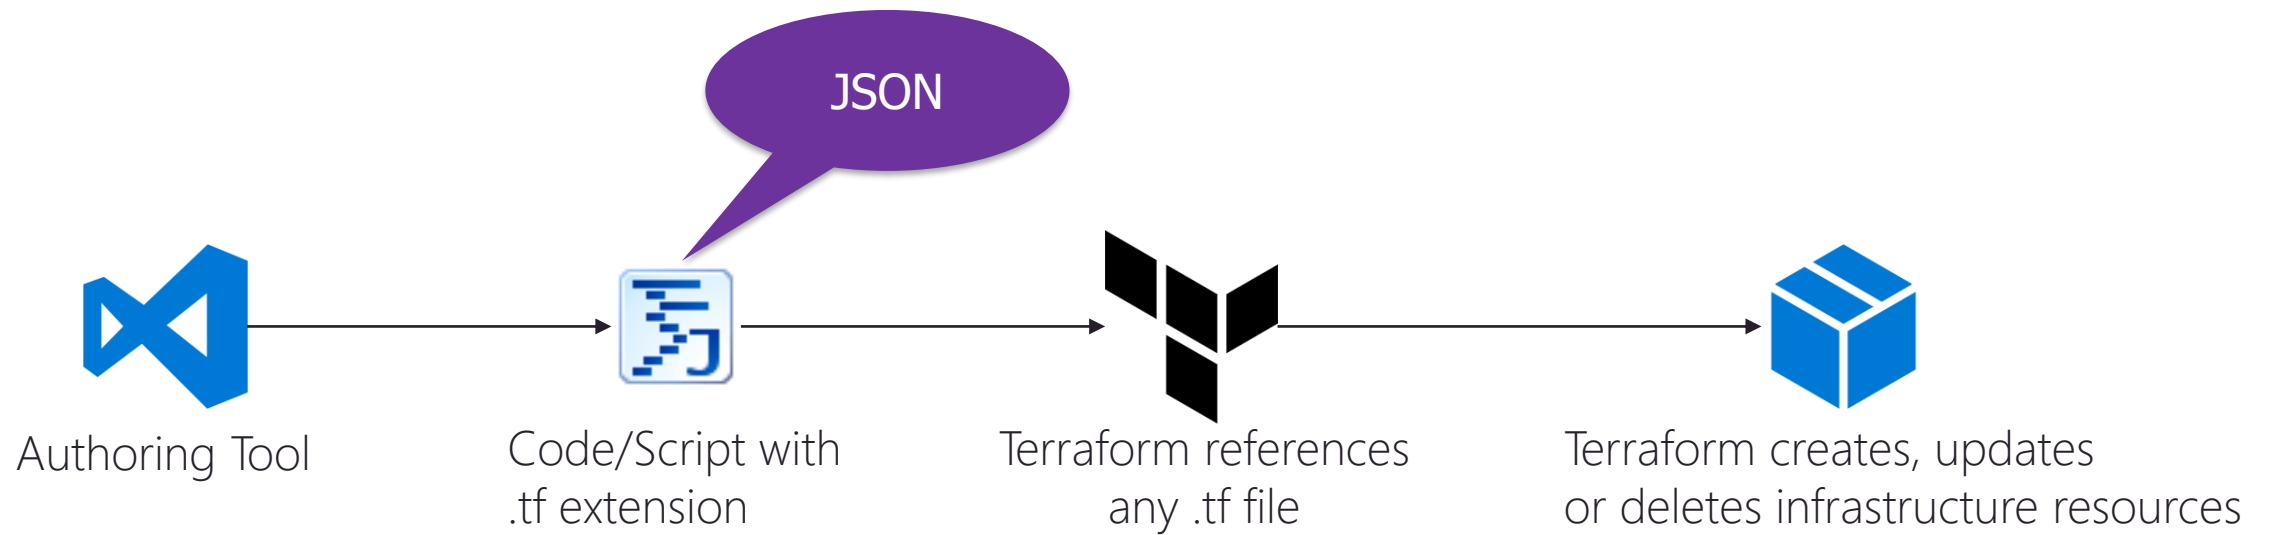
Simulating, detecting, and correcting for configuration drift



Understanding the basics of Terraform state management



# In a nutshell



# JSON ?

## | What is it ?

- | JSON stands for JavaScript Object Notation
- | JSON is a text format for storing and transporting data

## | Why do we need it ?

- | JSON is "self-describing" and easy to understand
- | JSON is a lightweight data-interchange format
- | JSON is language independent

# JSON Syntax

| JSON syntax is derived from JavaScript object notation syntax:

- | Data is in name/value pairs
- | Data is separated by commas
- | Curly braces hold objects
- | Square brackets hold arrays

```
{"employees":[  
    { "firstName":"John", "lastName":"Doe" },  
    { "firstName":"Anna", "lastName":"Smith" },  
    { "firstName":"Peter", "lastName":"Jones" }  
]}
```

# JSON Example

- | A site represents a set of buildings
- | A building has
  - | An address/location
  - | An identifier
  - | A name
  - | A collection of rooms
- | A building is assigned to a department
- | A room has a name and a max number of persons it can host

```
{
  "site": [
    "B20": {
      "address": "Nowhere street, 20",
      "identifier": 123456798,
      "name": "building-20",
      "member-of": [
        "IT",
        "Support"
      ],
      "rooms": [
        "room#1": {
          "name": "R2.2.1",
          "capacity": 4
        },
        "room#2": {
          "name": "R2.2.1",
          "capacity": 4
        }
      ]
    },
    "B21": {}
  ]
}
```

## Exercise 2.01 – JSON File

Check README file



# Terraform ?

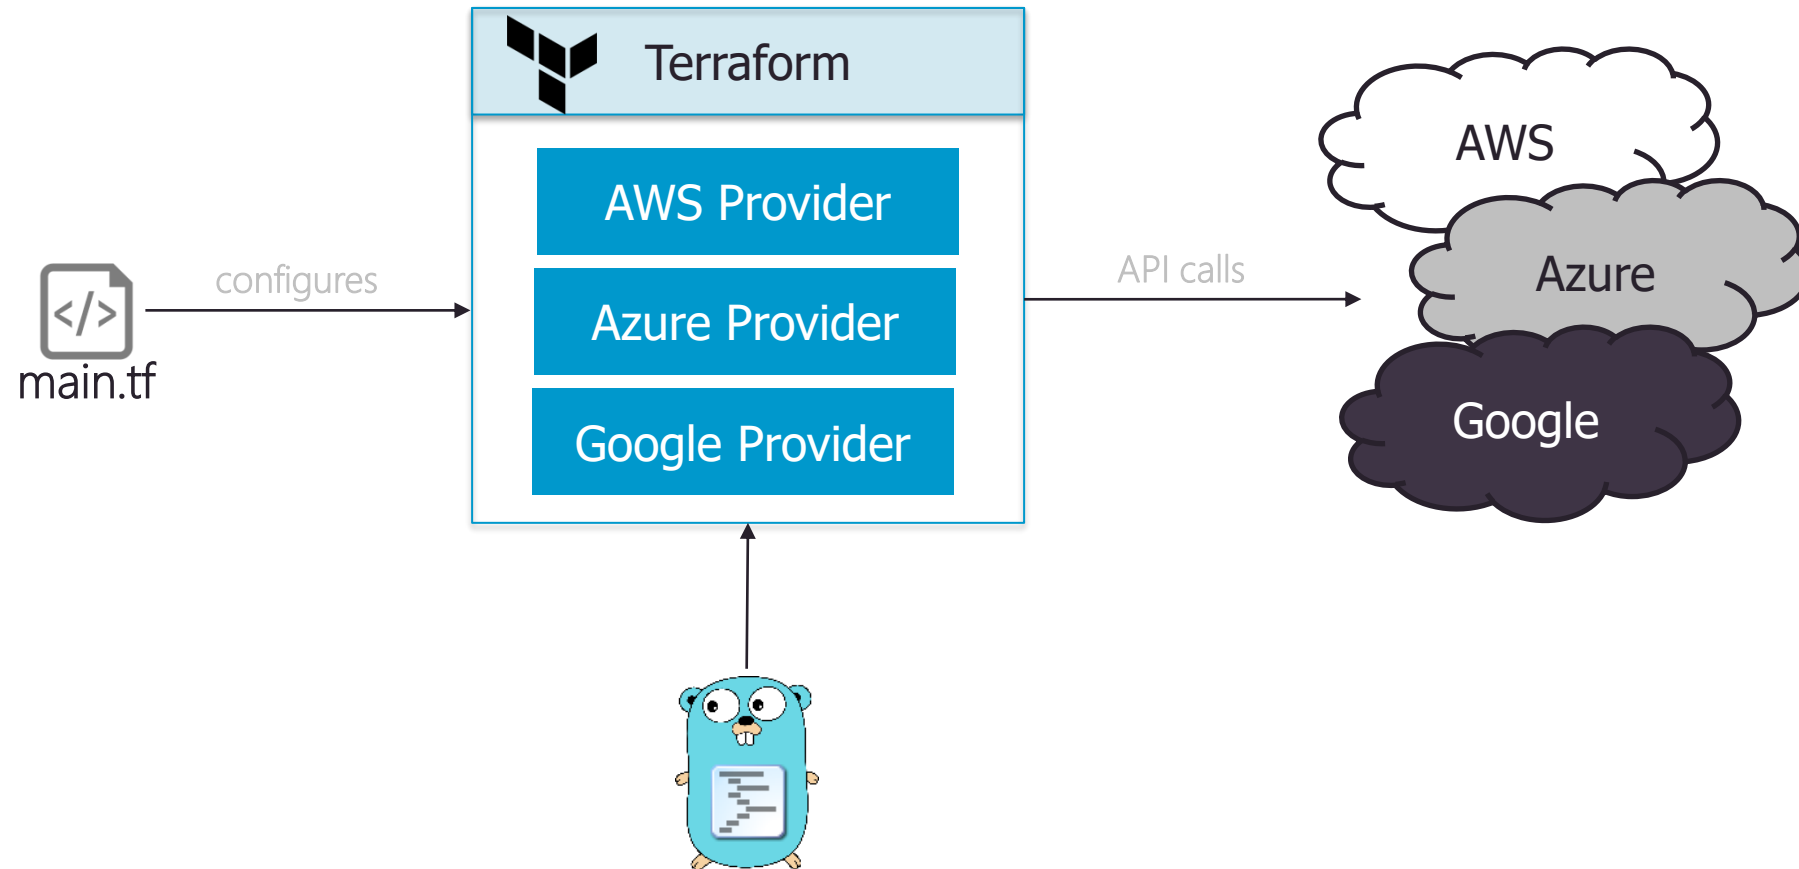
| Terraform is a deployment technology to provision and manage IaC.



# Terraform ?

- | The provisioning and management is done through the execution of commands on
  - | Workspace
  - | State
  - | Infrastructure
- | An Terraform infrastructure is represented by
  - | A set of resources [1..n]
  - | A set of providers [1..n]
  - | A set of data [0..n]

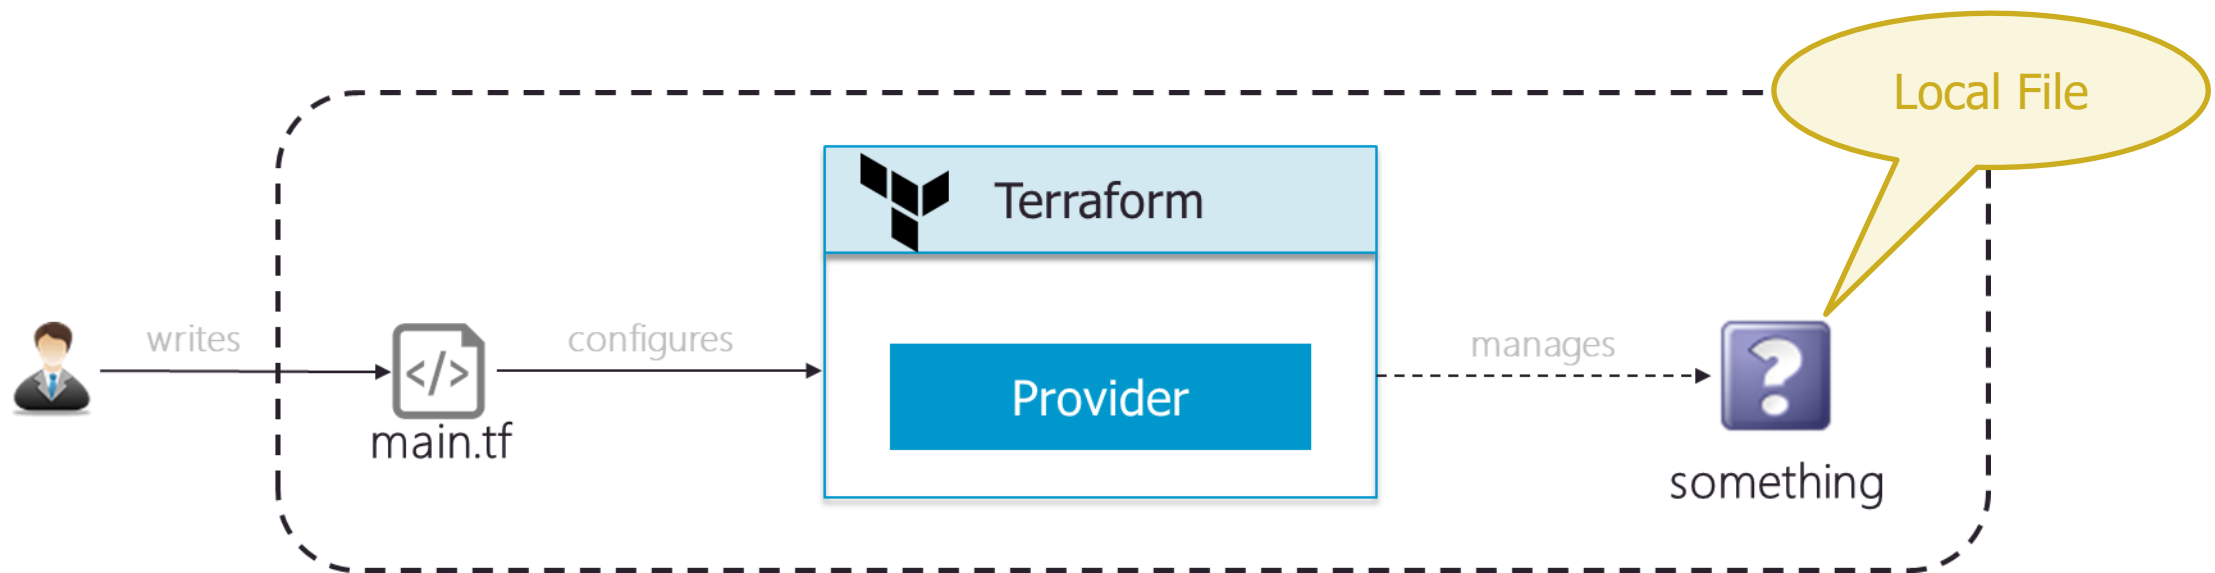
# Terraform ?



# Note: Local-only resources

| There are different sorts of resources: local-only resources.

| Examples ?



# Why Terraform ?

## | Key characteristics:

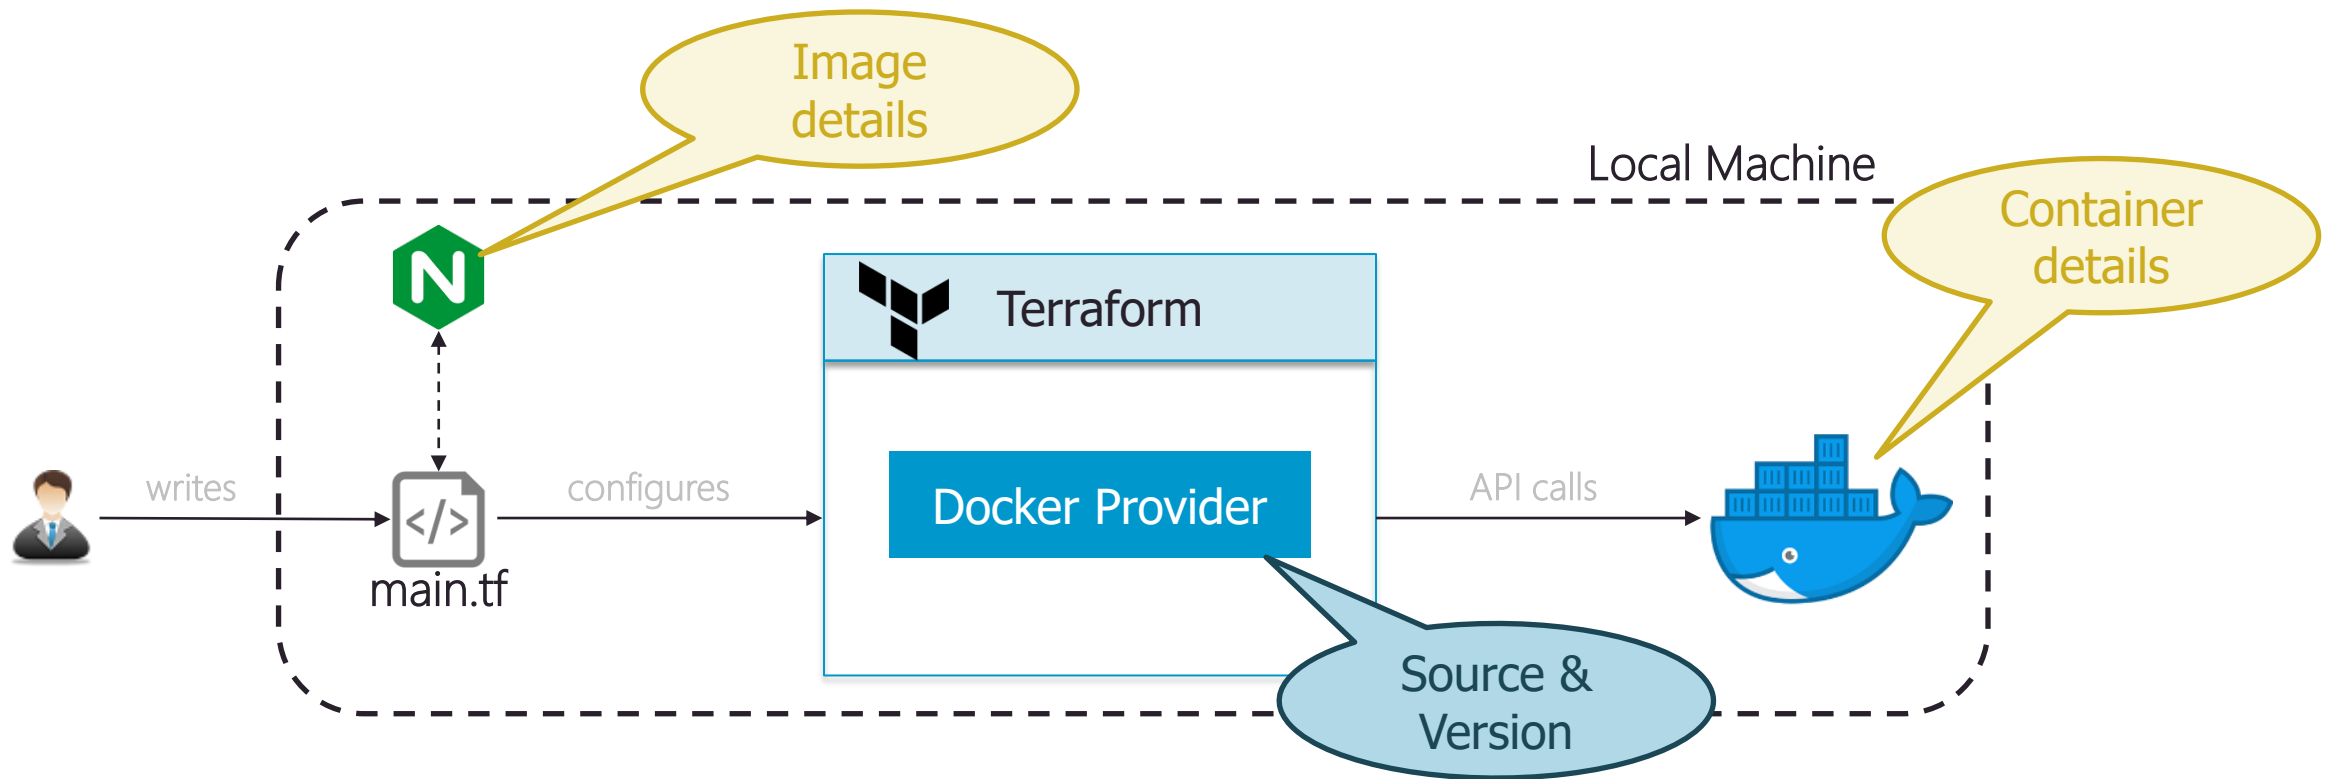
- | **Provisioning tool:** Deploys infrastructure, not just applications
- | **Easy to use:** For most of us (ie: non geniuses)
- | **Free and Open Source:** Who doesn't like free?
- | **Declarative:** Say what you want, not how to do it
- | **Cloud agnostic:** Deploy to any cloud using the same tool
- | **Expressive and extendable**



# Terraform vs other IaC tools

	Provisioning tool	Easy to use	Free and Open Source	Declarative	Cloud Agnostic	Expressive and extendable
Ansible		X	X		X	X
Chef			X	X	X	X
Puppet			X	X	X	X
SaltStack		X	X	X	X	X
Terraform	X	X	X	X	X	X
Pulumi	X		X		X	X
AWS CloudFormation	X	X		X		
GCP Deployment Manager	X	X		X		
Azure Resource Manager	X			X		

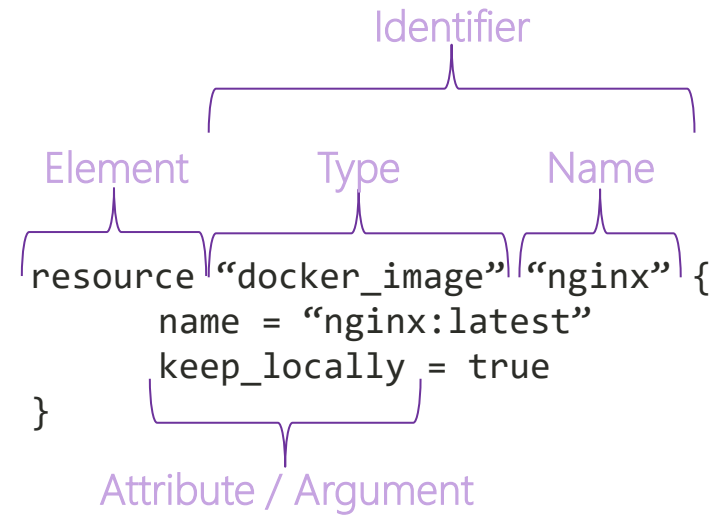

# "Hello World !"



# Terraform resource

## | Image

- | Name: nginx
- | Keep\_locally: true




# Terraform provider

| Provider: docker

| Source: kreuzwerker/docker

| Version: ~> 2.13.0



```
terraform {  
  required_providers {  
    docker = {  
      source = "kreuzwerker/docker"  
      version = "~> 2.13.0"  
    }  
  }  
}
```

# Exercise 2.02 – Docker – Hello world!

Check README file

## | Terraform

### | Provider

| Source: kreuzwerker/docker

| Version: ~> 2.13.0

## | Docker resources

### | Image

### | Container:

| Image: docker\_image.nginx.latest

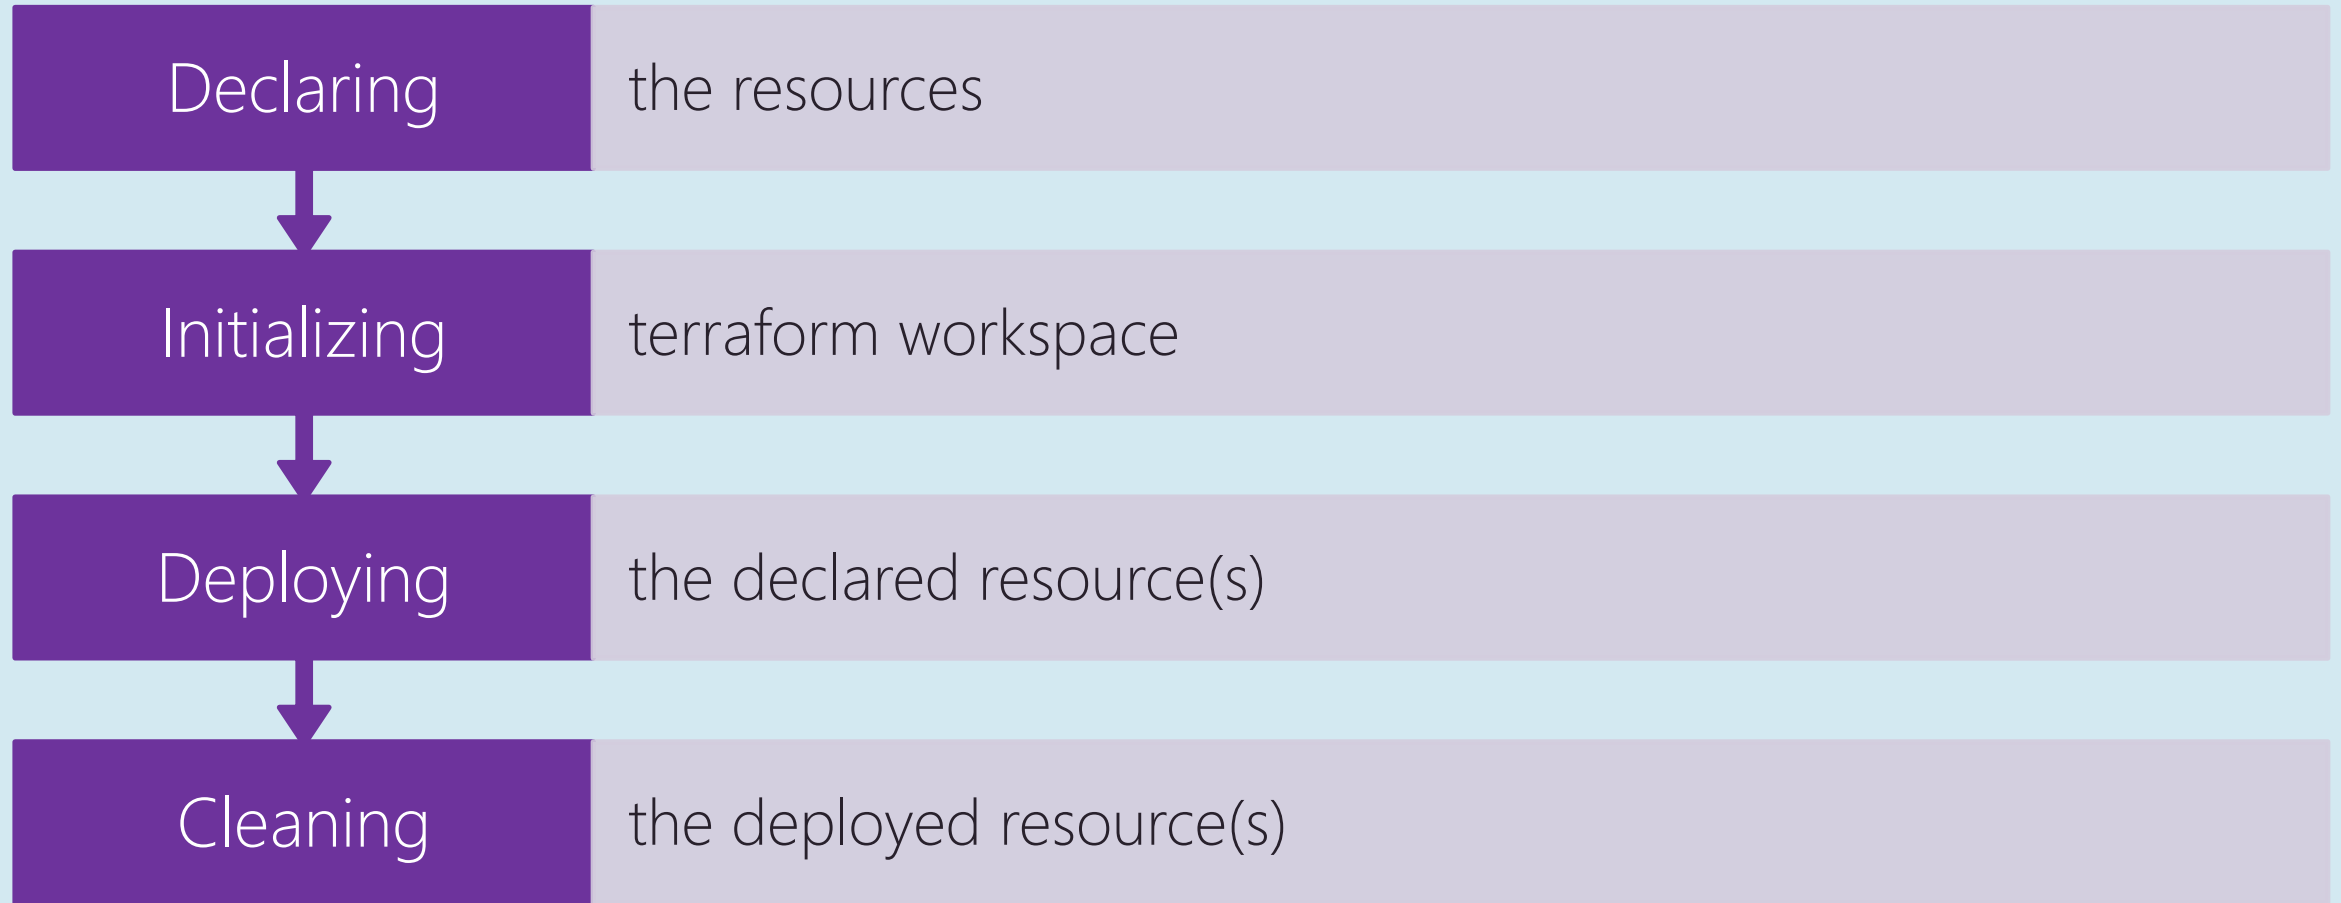
| Name: helloW

| Ports:

| Internal: 80

| External: 8000

# Deployment process in a nutshell (demo)



## Exercise 2.03 – Azure – Hello world!

Check README file

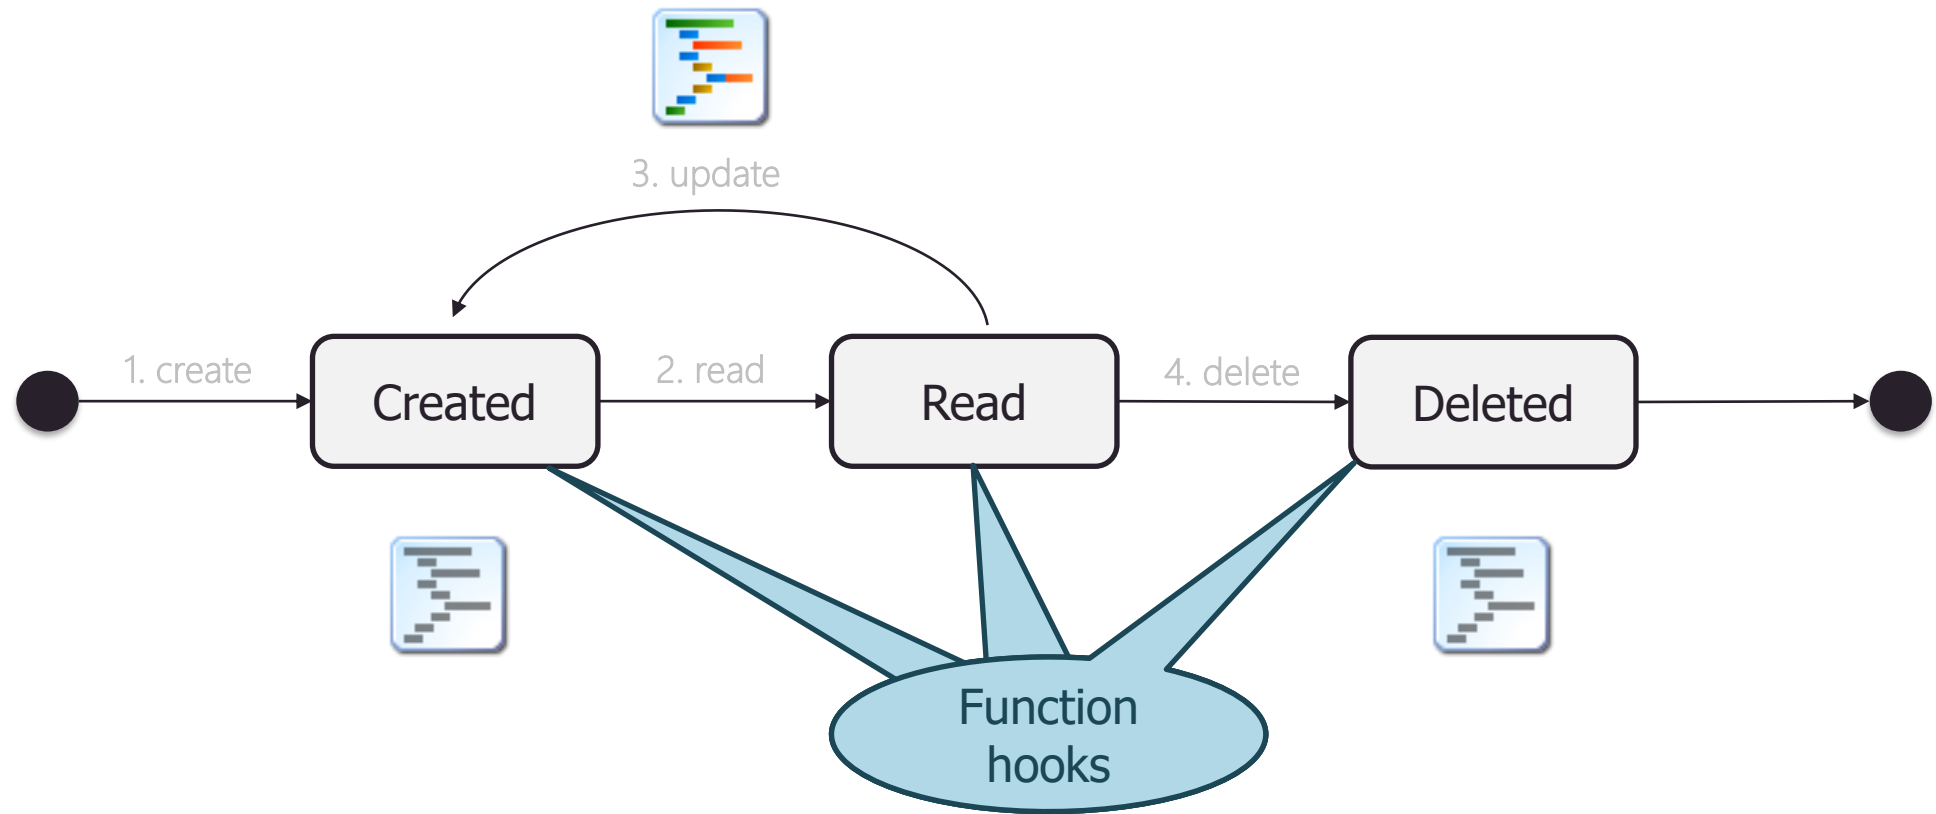
| Assumption: familiar with Azure concepts & particularly resource groups

## So far

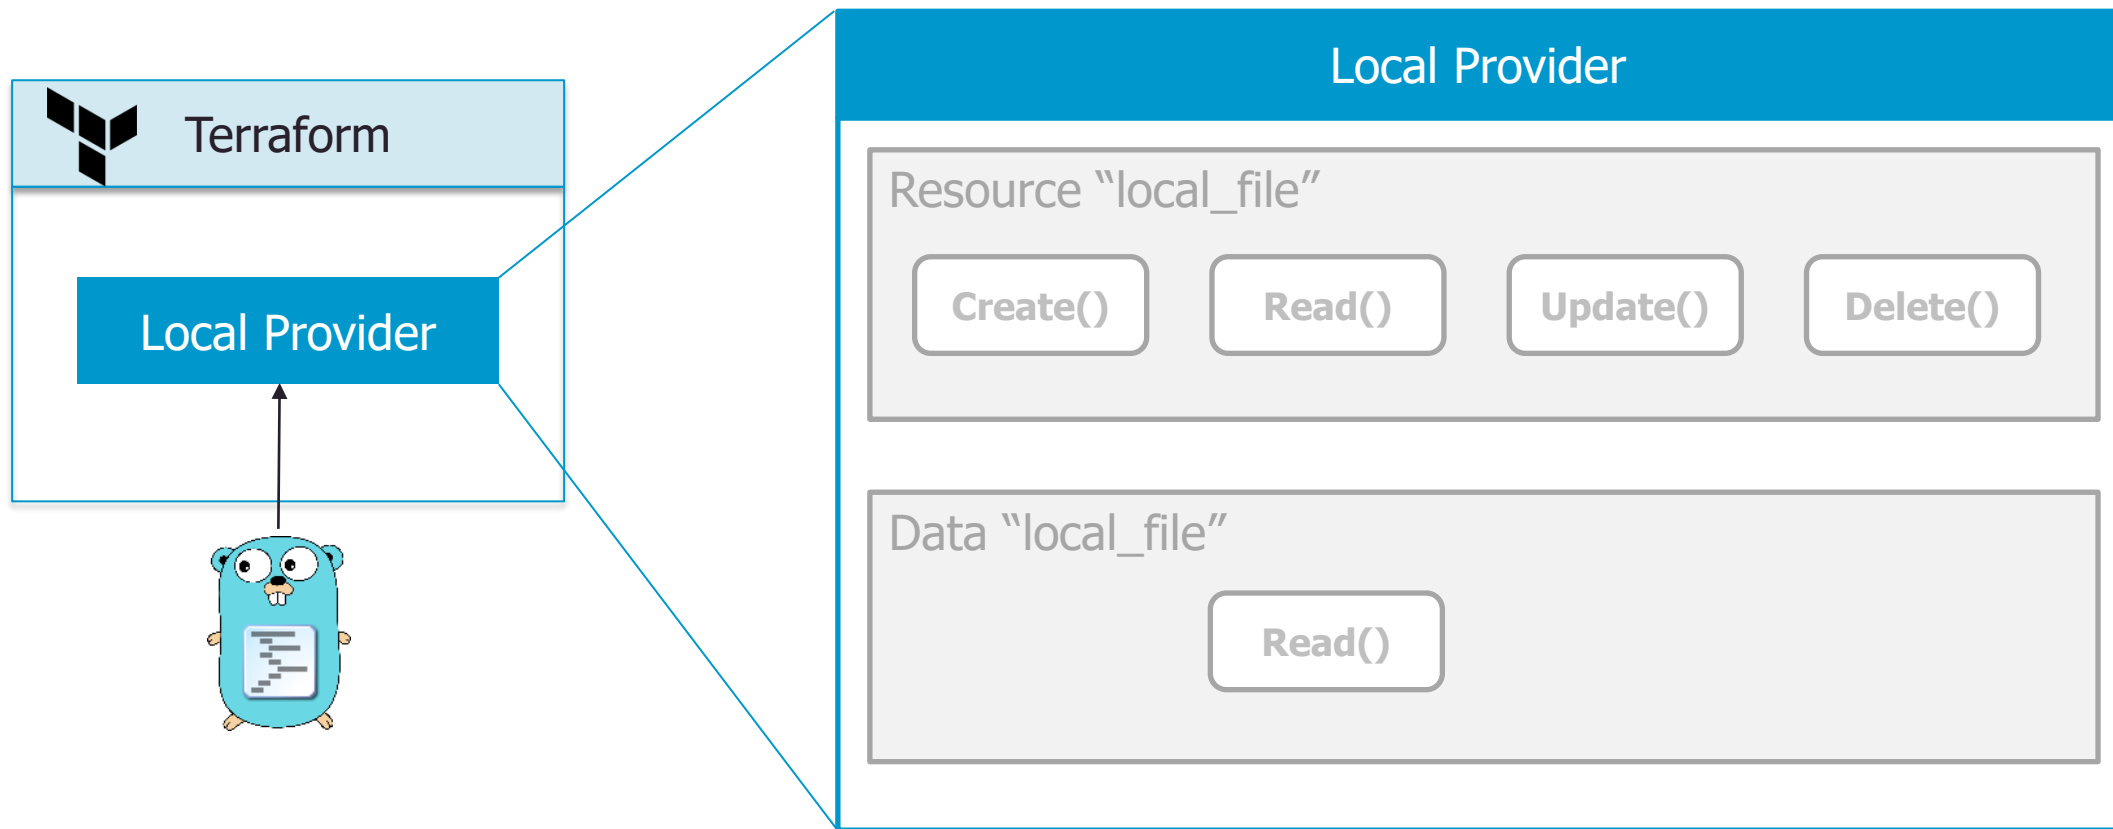
- | Terraform is a declarative, provisioning state management tool that performs CRUD operations on managed resources deployed onto any public or private cloud.
- | The major elements of Terraform are resources, data sources and providers.
- | To deploy a Terraform project you must first write configuration code, then configure providers and other input variables, initialize Terraform and finally apply changes. Clean-up is done with a destroy run.



# Resource lifecycle (States)



# Note: Function hook



# Terraform commands

## | Infrastructure (data, resource & provider)

- | Initialise: `init`, `get`
- | Provision: `plan`, `apply`, `destroy`
- | Inspect: `graph`, `output`, `show`
- | Import: `import`

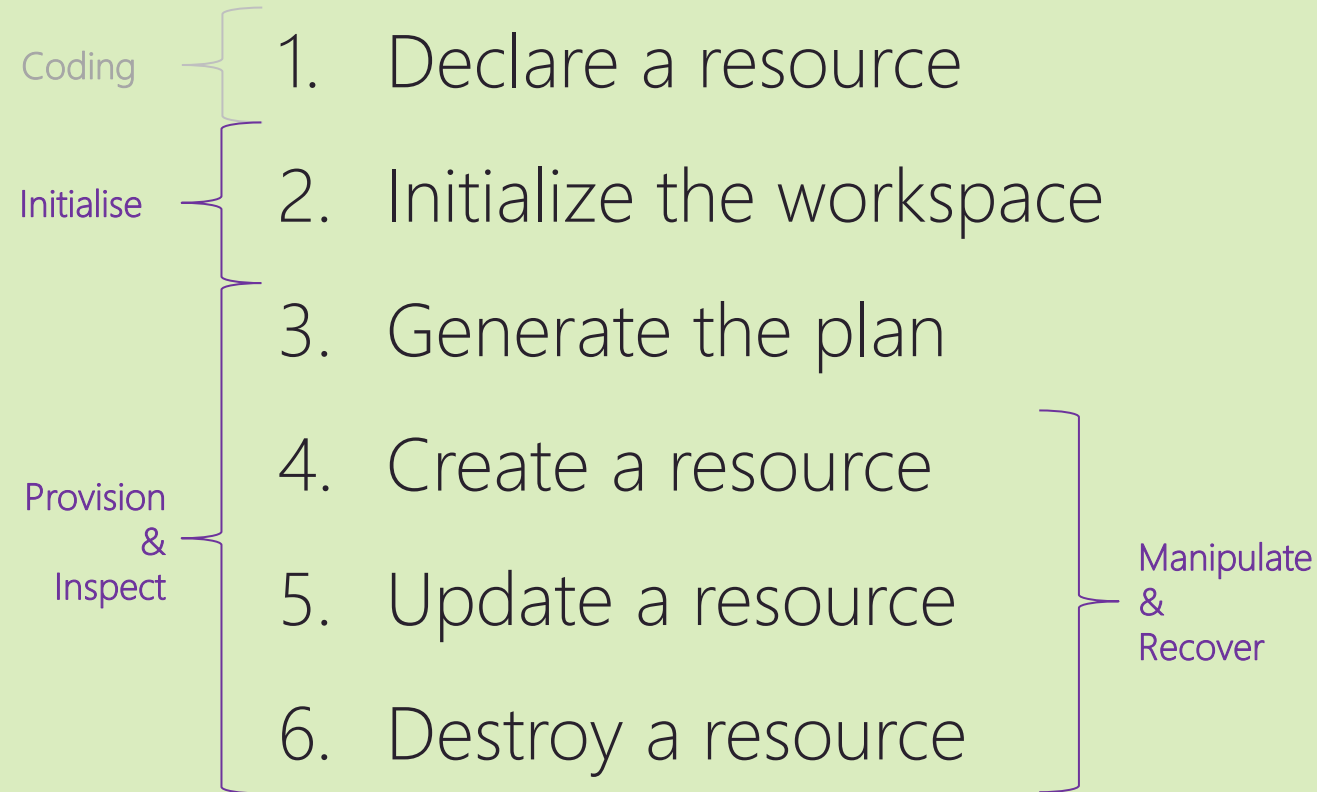
## | State

- | Manipulate: `list`, `show`, `refresh`, ...
- | Recover: `pull`, `push`, `force-unlock`, ...

## | Workspace

- | Manage: `list`, `select`, `new`, `delete`, `show`

# Example 2.04 - Lifecycle



# Declare a local file resource

Block 1

```
terraform {  
  required_version = ">= 0.15"  
  required_providers {  
    local = {  
      source = "hashicorp/local"  
      version = "~> 2.0"  
    }  
  }  
}
```

Block 2

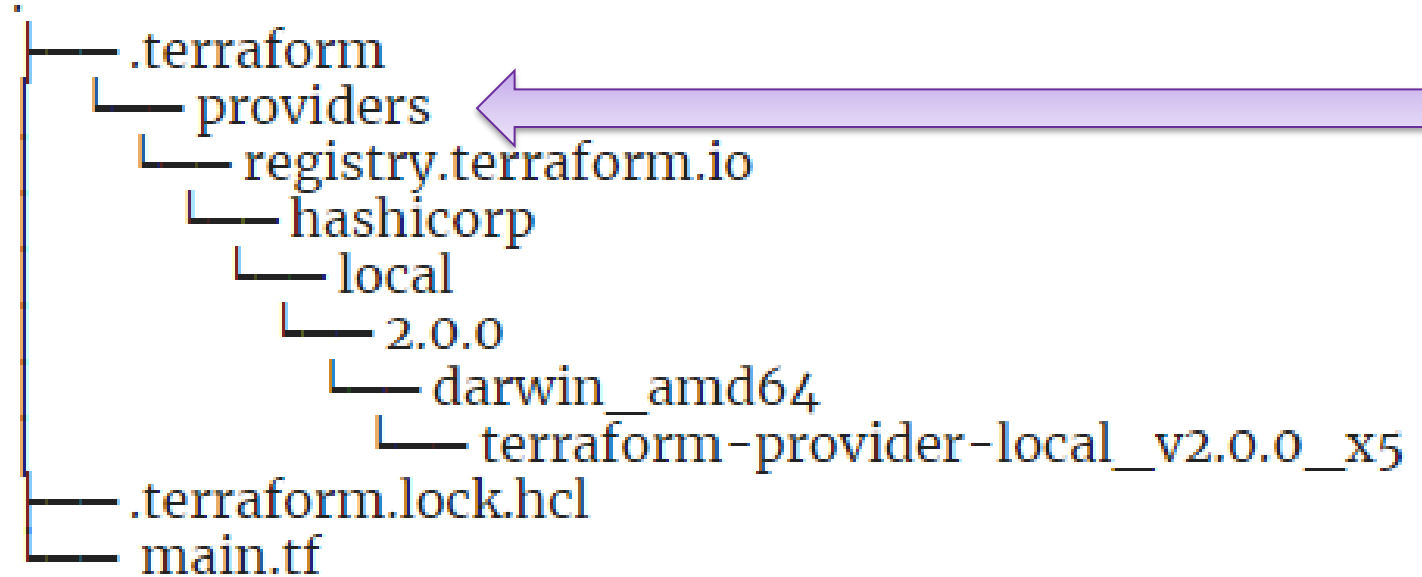
```
resource "local_file" "Britney" {  
  filename = "I_did_it_again.txt"  
  content = <<-EOF  
  'Cause to lose all my senses  
  
  That is just so typically me  
  You see, my problem is this, I'm dreaming away  
  Can't you see I'm a fool in so many ways?  
  EOF  
}
```



HEREDOC  
SYNTAX

# Initializing the workspace

```
| terraform init
```



# Format the script

```
| terraform fmt -recursive
```

```
| terraform fmt -write=false -diff=true
```

```
1
2  # Module Networking
3  module "networking" {
4    source = "./modules/networking"
5    location = var.location
6    suffix = var.suffix
7    costalloc = "it-hq"
8  }
9
10 # Module Database
    1 reference
11 module "database" {
12   source = "./modules/database"
13   rg = module.networking.rg
14   suffix = var.suffix
15   costalloc = "it-dba"
16 }
```

```
1
2  # Module Networking
3  module "networking" {
4    source      = "./modules/networking"
5    location    = var.location
6    suffix      = var.suffix
7    costalloc   = "it-hq"
8  }
9
10 # Module Database
    1 reference
11 module "database" {
12   source      = "./modules/database"
13   rg          = module.networking.rg
14   suffix      = var.suffix
15   costalloc   = "it-dba"
16 }
```

# Validate the script

```
| terraform validate
```

or

```
| terraform validate -json
```

## Error: Unsupported attribute

```
on modules/database/main.tf line 18, in resource "azurerm_sql_server" "training":  
18:   administrator_login      = random_pet.login.result
```

This object has no argument, nested block, or exported attribute named "result".



# Generating the execution plan

- | Plan is a static code analysis, read-only, idempotent and "*robust*"
- | Always run a terraform plan before deploying

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:

+ create

Terraform will perform the following actions:

```
# local_file.turing will be created
+ resource "local_file" "turing" {
  + content                = <<-EOT
    On Computable Numbers with an application to the entscheidungsproblem

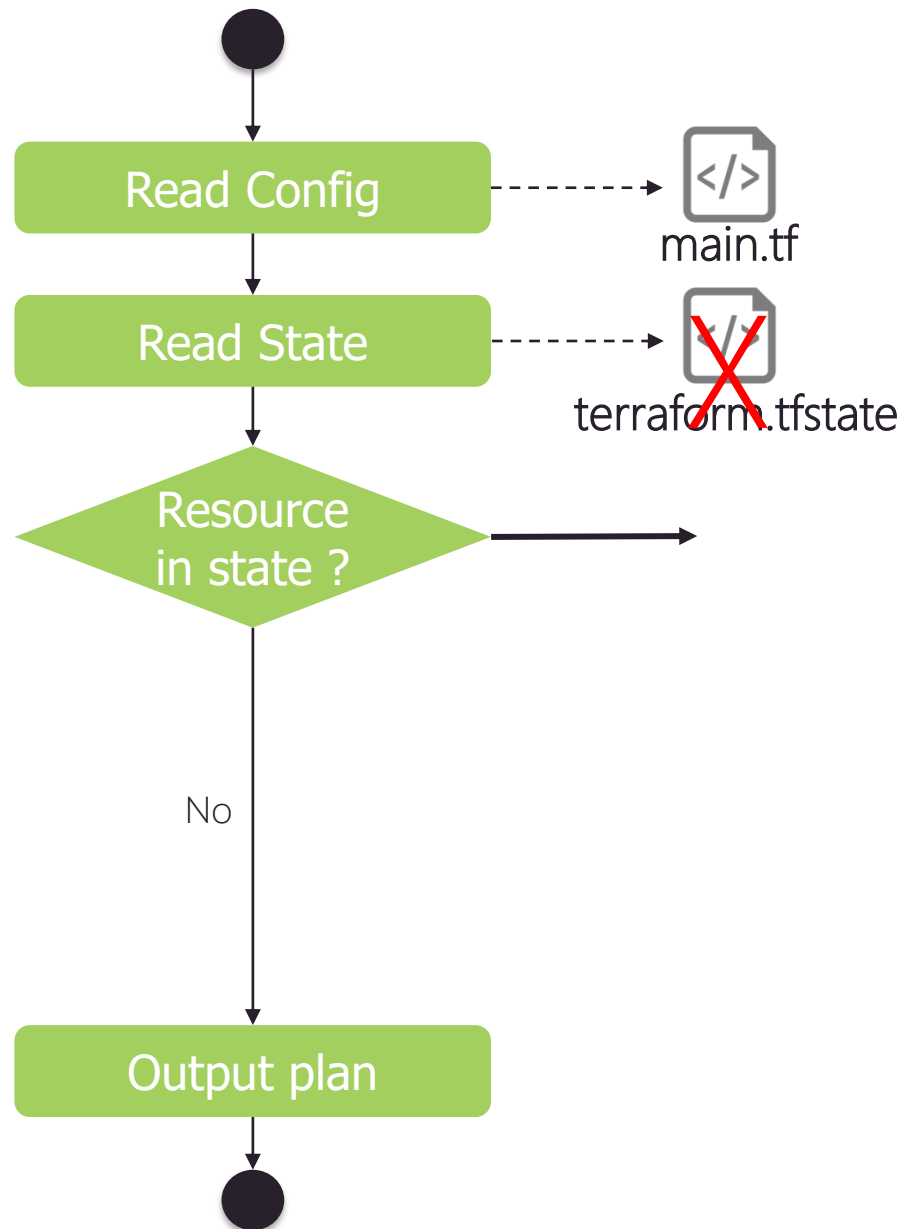
    The computable numbers may be described briefly as the real numbers
    whose expressions as a decimal are calculable by finite means.
  EOT
  + directory_permission = "0777"
  + file_permission      = "0777"
  + filename              = "OnComputableNumbers.txt"
  + id                   = (known after apply)
}
```

Plan: 1 to add, 0 to change, 0 to destroy.

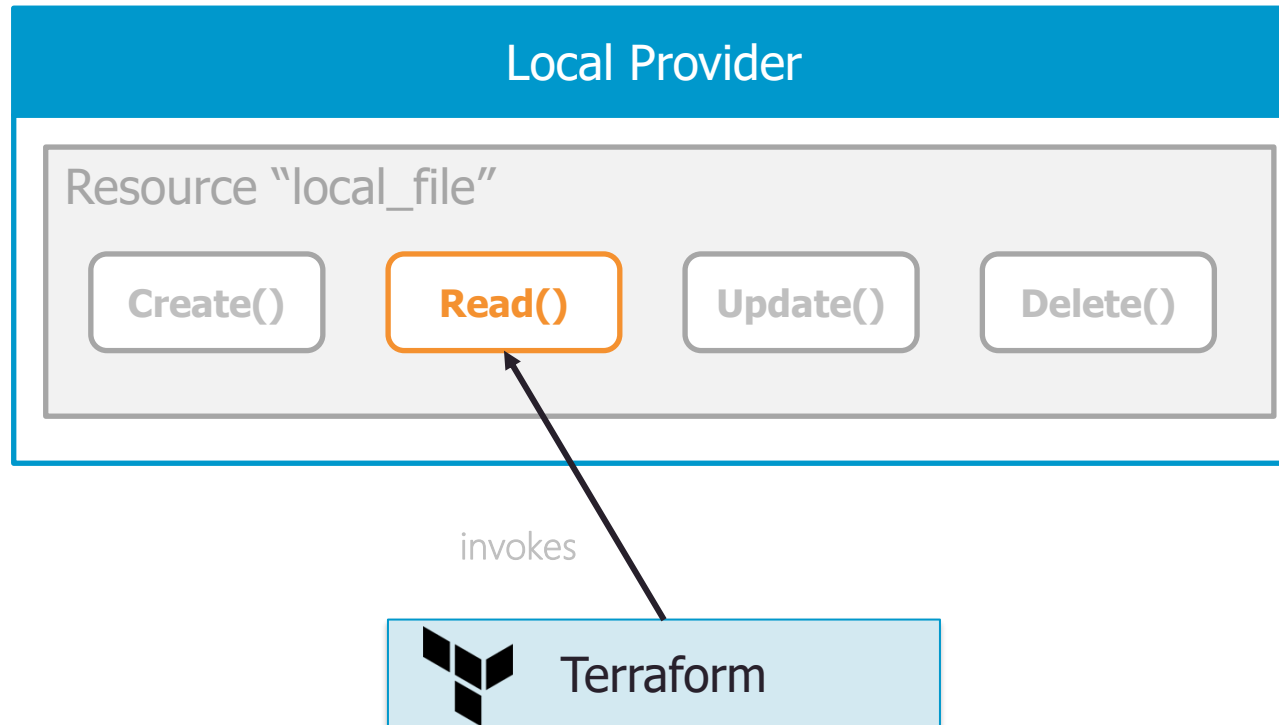


PLAN RUNNING  
SLOW

# So far



# No-op()



```
PS D:\BrainVault\sources\CadaiHub\tfTraining\temp> terraform plan
local_file.turing: Refreshing state... [id=759a14a307f9d59c0f92e83e5e81c1edc1a74fa6]
```

**No changes.** Your infrastructure matches the configuration.

Terraform has compared your real infrastructure against your configuration and found no differences, so no changes are needed.

# Generating the execution plan

```
| terraform plan -out plan.out
```

```
| terraform show -json plan.out > plan.json
```

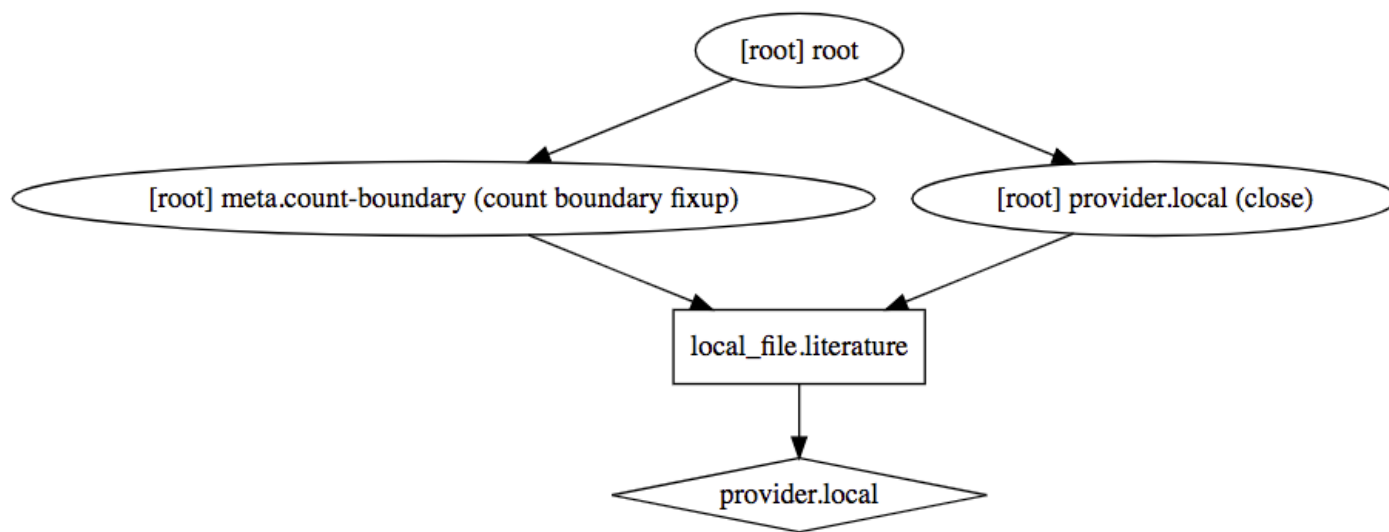


# Visualizing the plan

```

digraph {
  compound = "true"
  newrank = "true"
  subgraph "root" {
    "[root] local_file.turing (expand)" [label = "local_file.turing", shape = "box"]
    "[root] provider[\"registry.terraform.io/hashicorp/local\"]" [label = "provider[\"registry.terraform.io/hashicorp/local\"]", shape = "diamond"]
    "[root] local_file.turing (expand)" -> "[root] provider[\"registry.terraform.io/hashicorp/local\"]"
    "[root] meta.count-boundary (EachMode fixup)" -> "[root] local_file.turing (expand)"
    "[root] provider[\"registry.terraform.io/hashicorp/local\"] (close)" -> "[root] local_file.turing (expand)"
    "[root] root" -> "[root] meta.count-boundary (EachMode fixup)"
    "[root] root" -> "[root] provider[\"registry.terraform.io/hashicorp/local\"] (close)"
  }
}

```



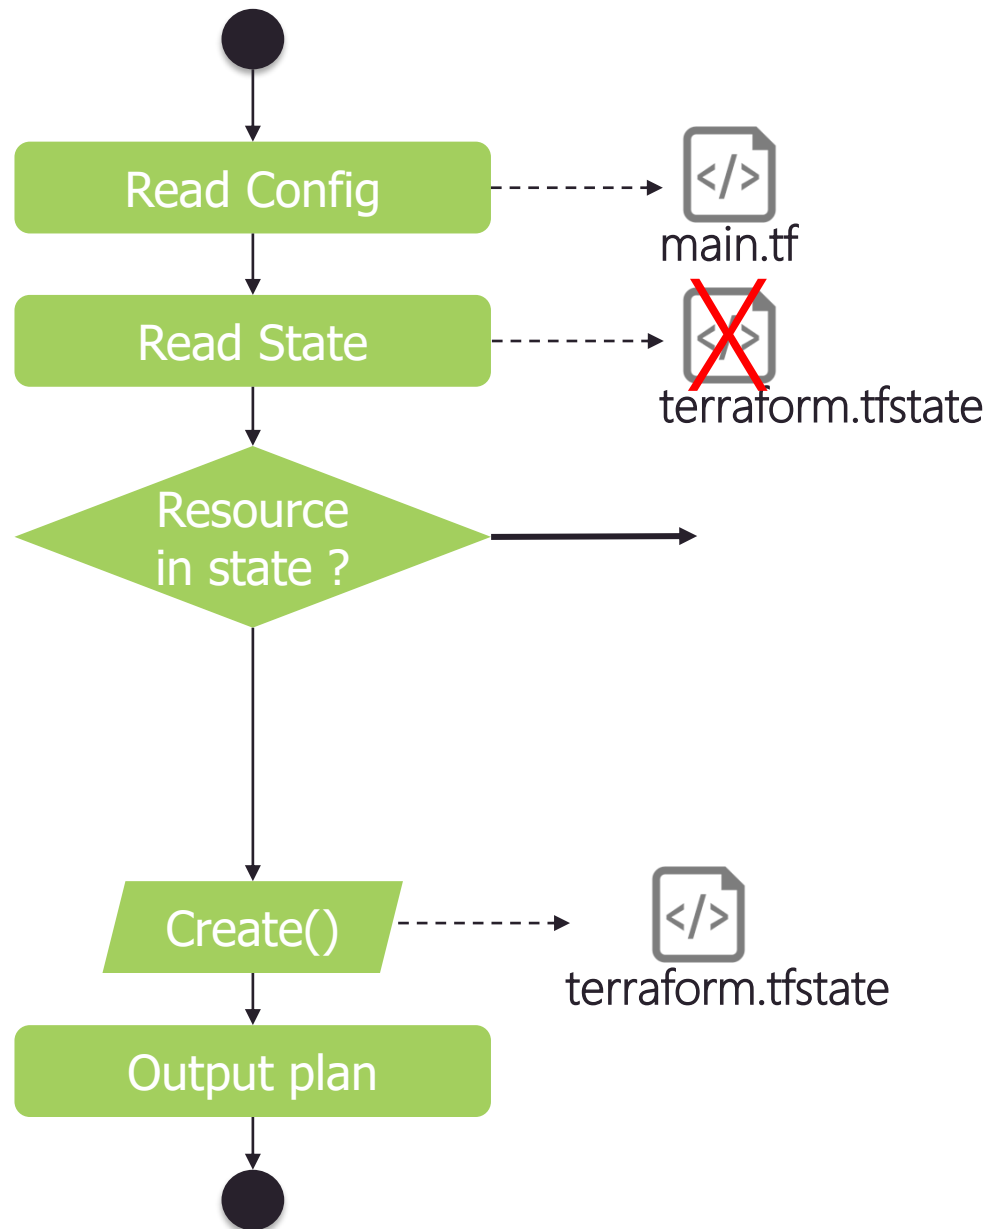
# Creating the resource

```
| terraform apply "plan.out"
```

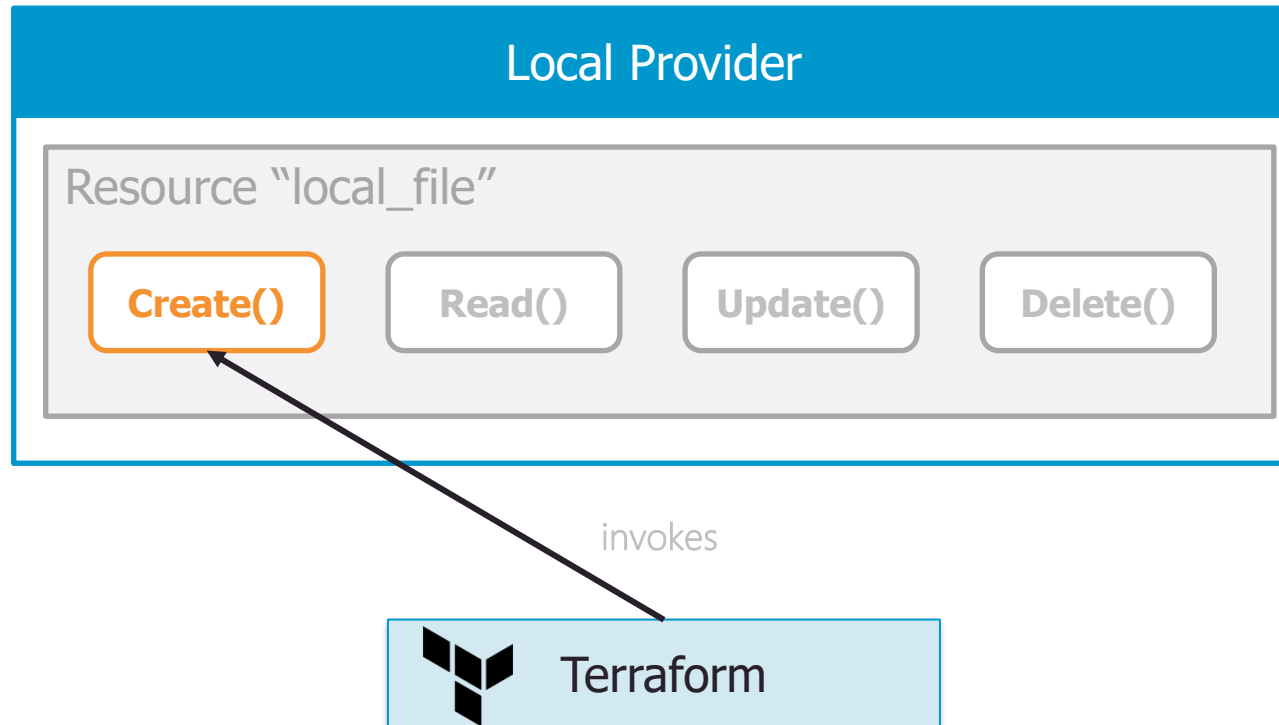
or

```
| terraform plan -out plan.out && terraform apply "plan.out"
```

# So far



# Behind the scene - Terraform apply





# No-op() - redo

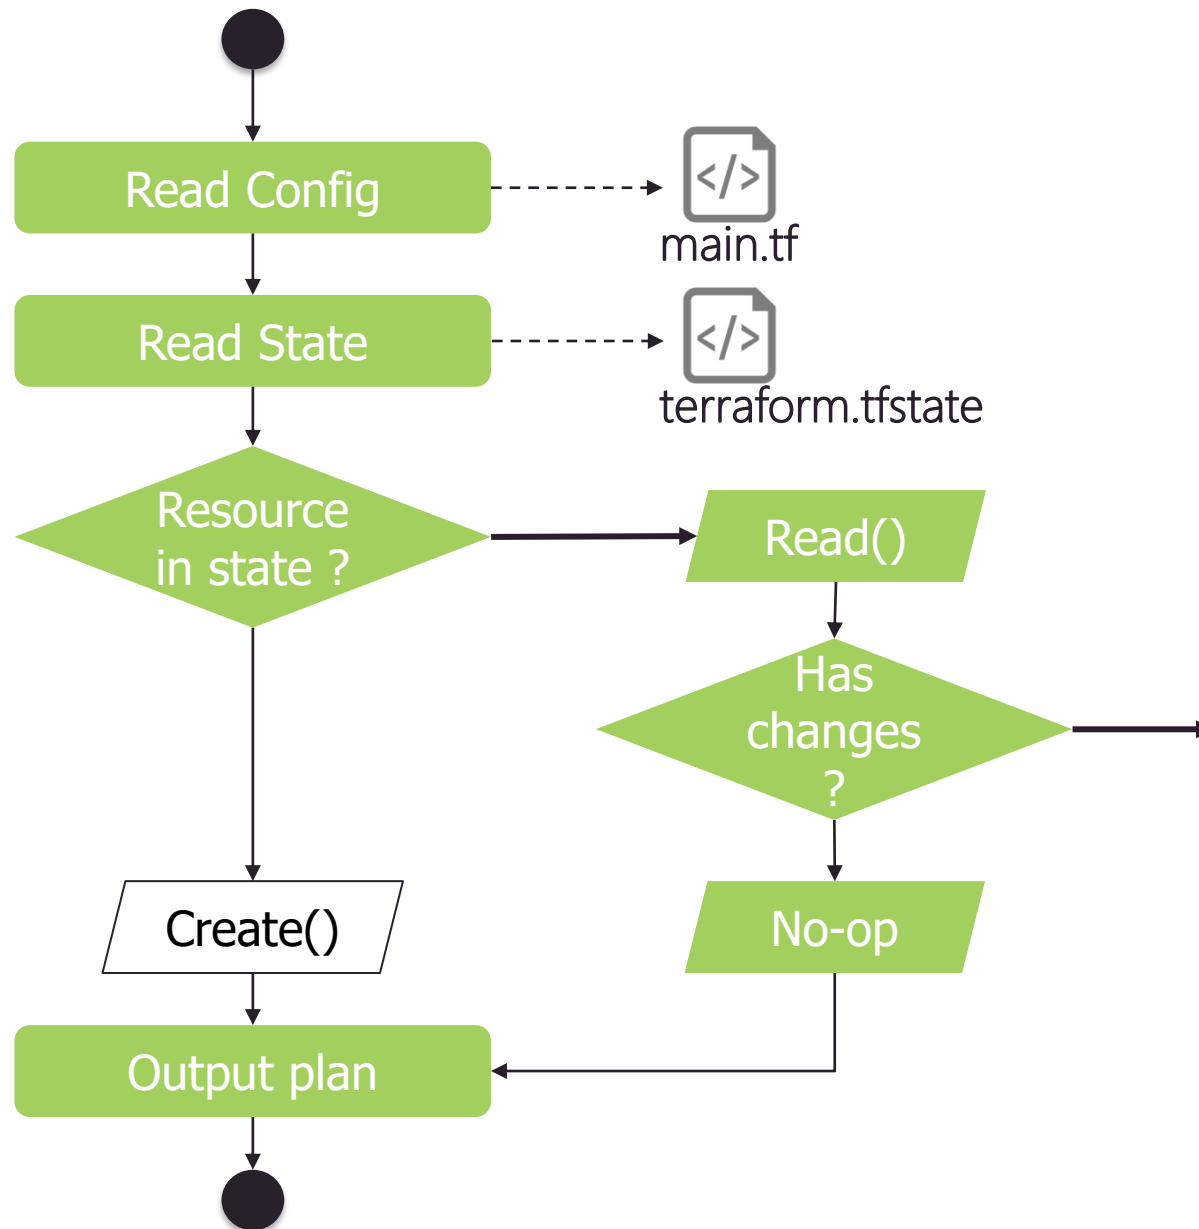
| terraform plan

```
maximedehaut@Maximes-MacBook-Pro ex03 % terraform plan
local_file.Britney: Refreshing state... [id=a42df82844eb946abad149069c0785dab26eac21]
```

**No changes.** Your infrastructure matches the configuration.

Terraform has compared your real infrastructure against your configuration and found no differences, so no changes are needed.

# So far



# Terraform states

- | Store stateful information about only three configuration blocks

- | `terraform state list`

- | `terraform show`

- | But what about

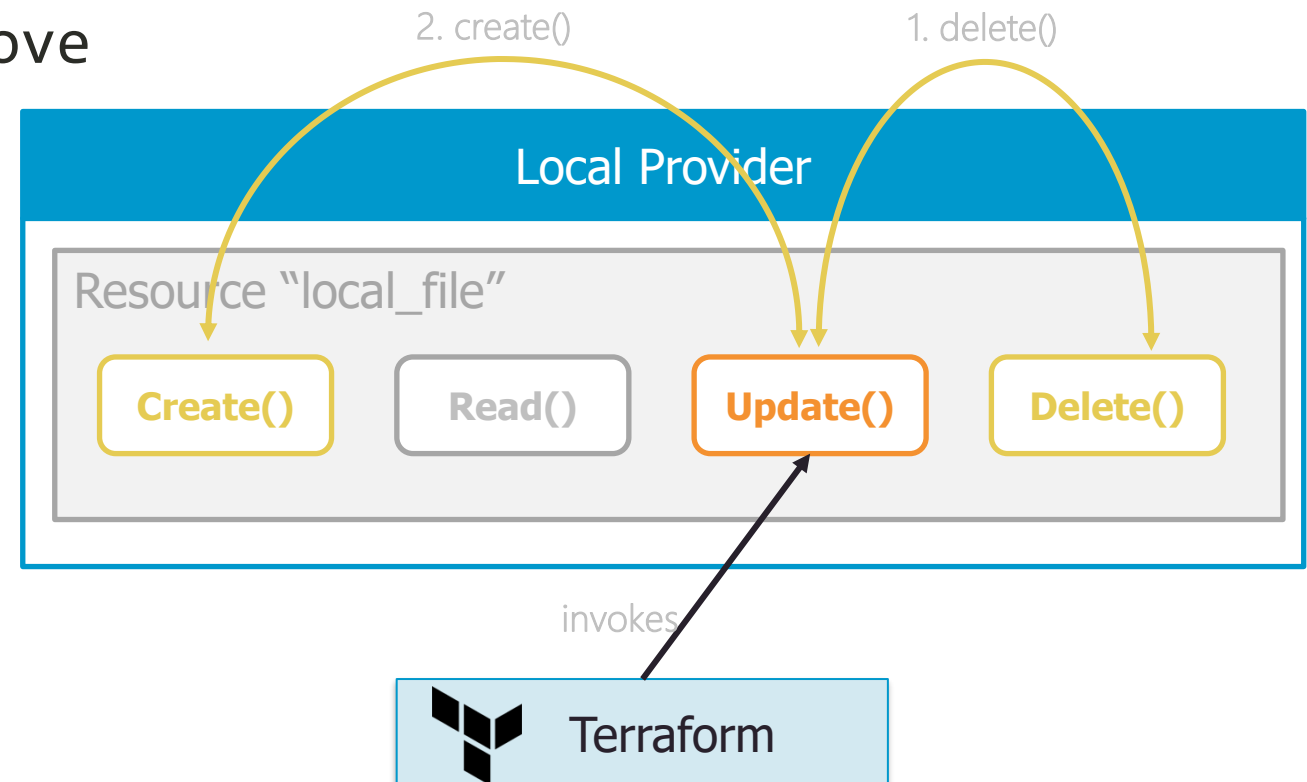
- | Shared storage

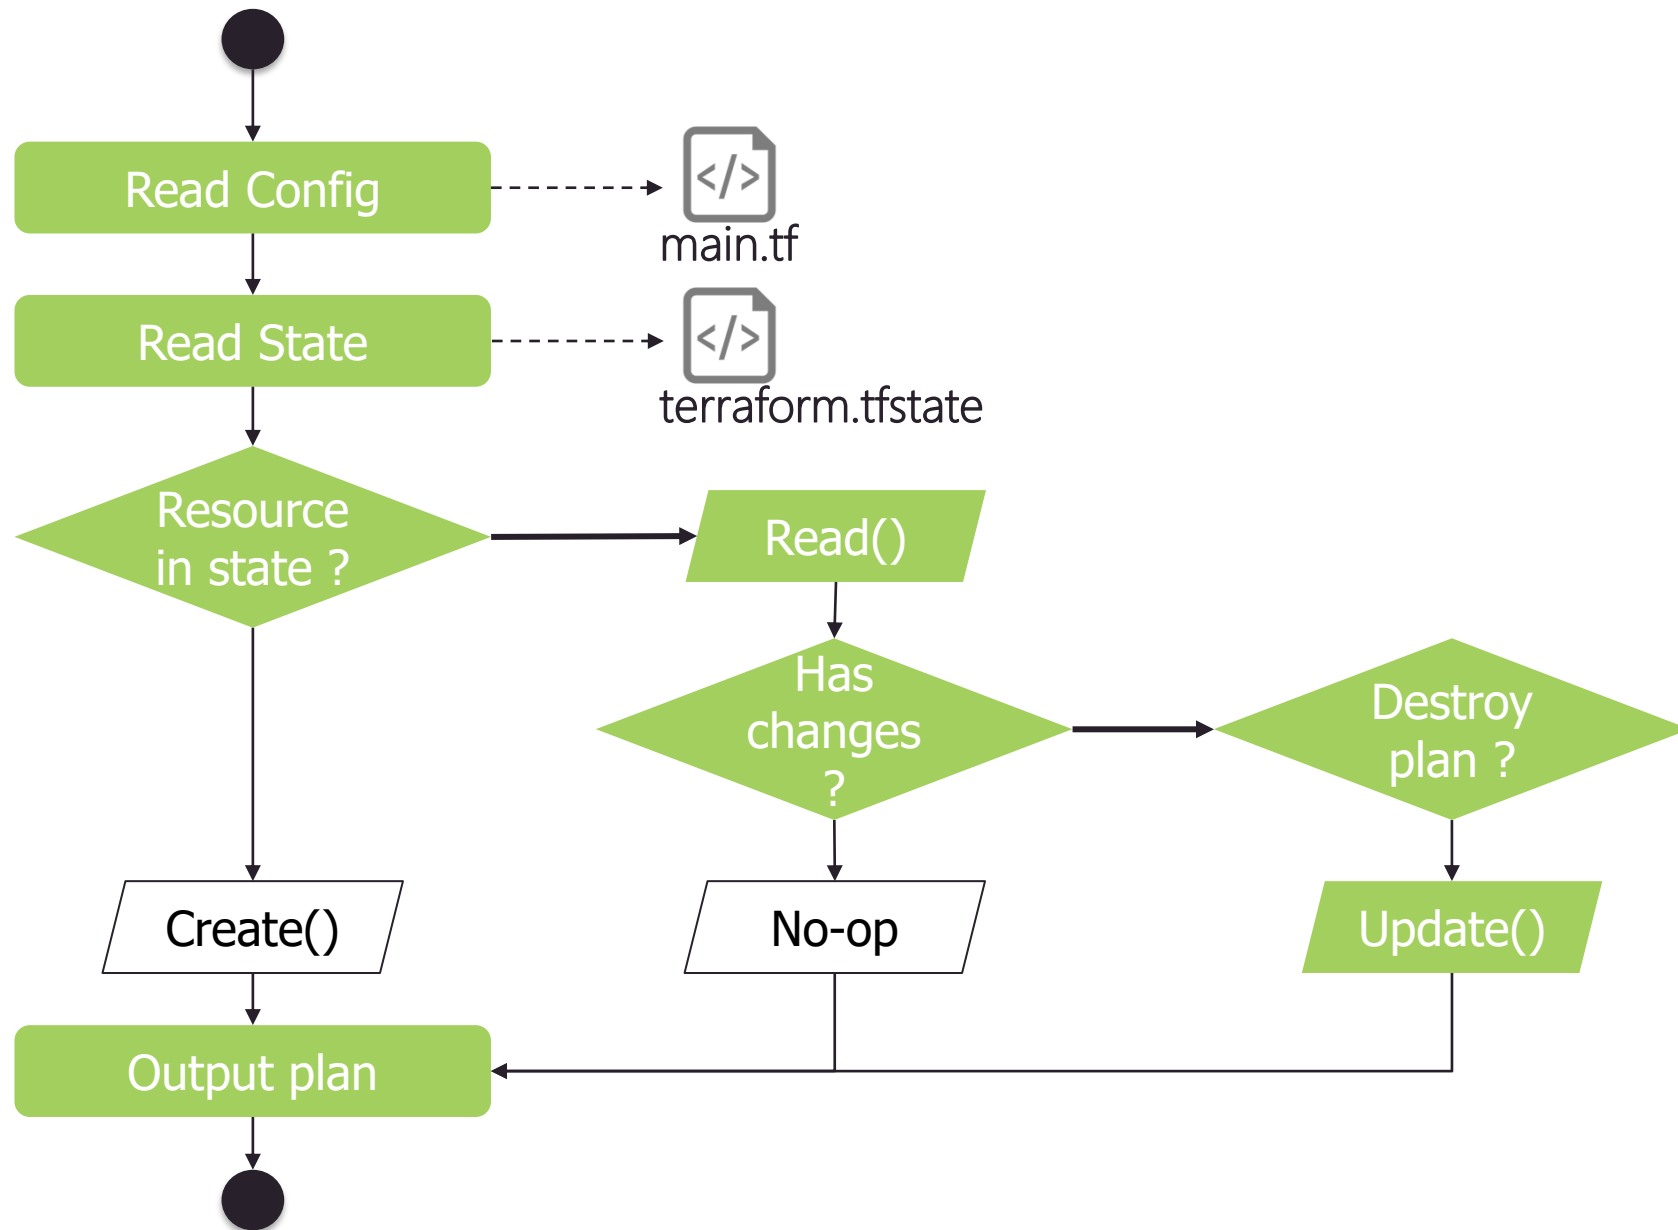
- | Locking

- | Isolation

# Immutable vs mutable update

- | Update main.tf code or pass a variable
- | terraform plan
- | terraform apply -auto-approve





# Configuration drift

| Change content of text file

| `terraform plan`

| `terraform show`

| `terraform apply -refresh-only (= terraform refresh)`

| Existing resource deleted

| New resource created

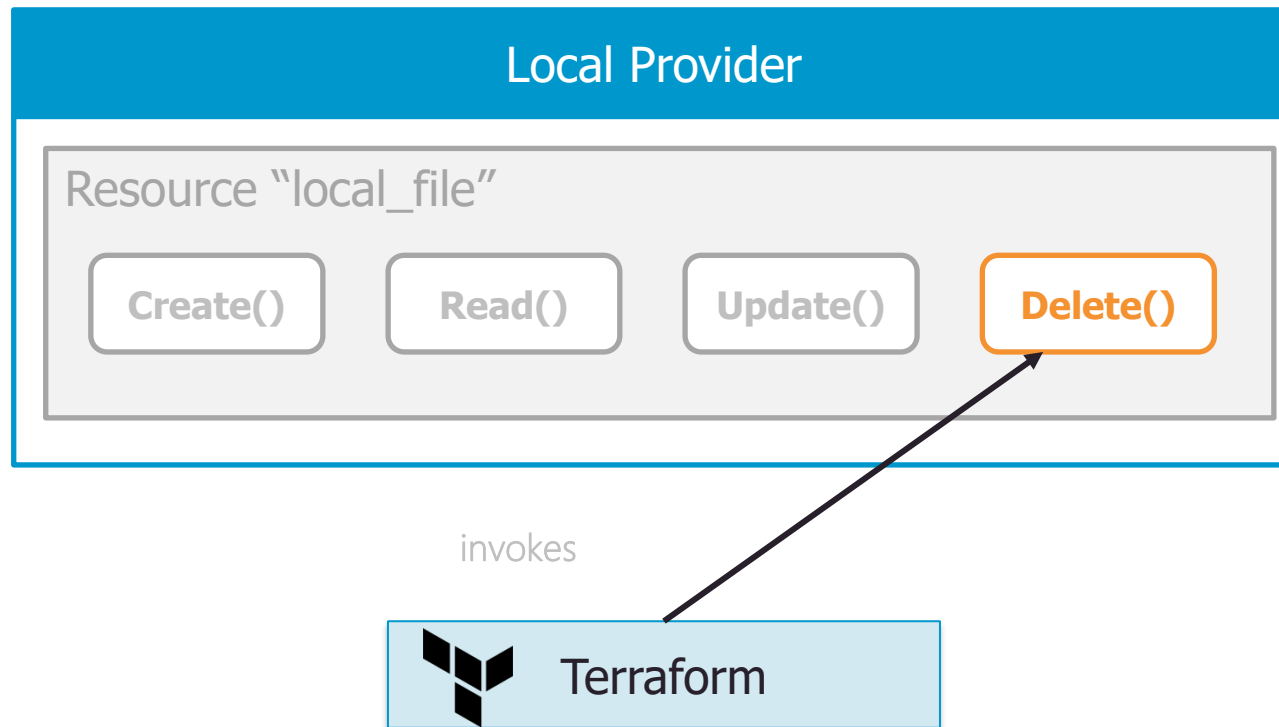


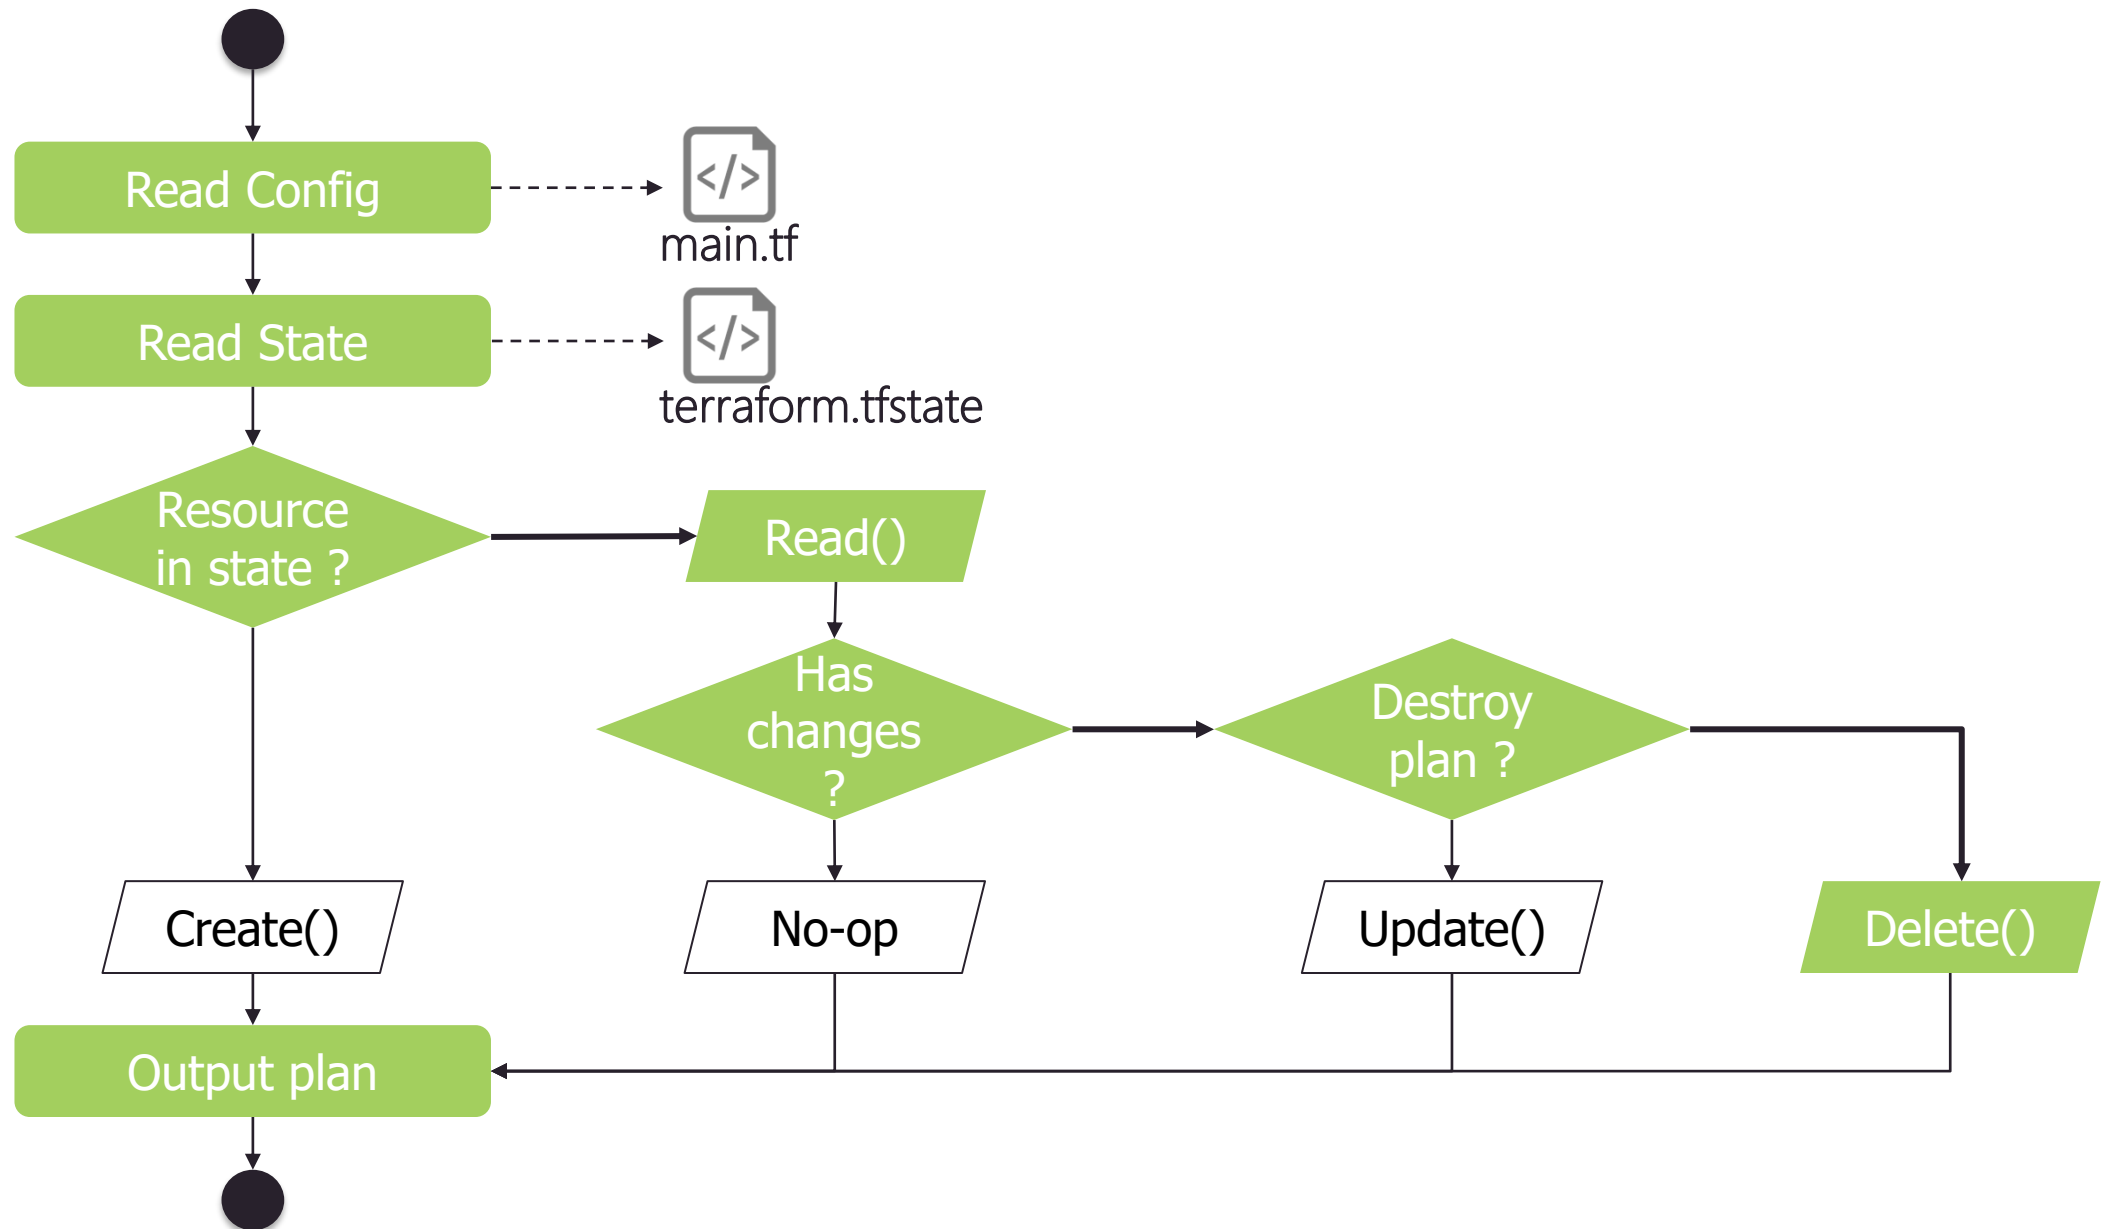
Show the state of  
each resource

reconcile the  
state

# Deleting a resource

| `terraform destroy -refresh-only -auto-approve`







# Note: Azure Authentication

- | Authenticating to Azure using

- | The Azure CLI

- | Managed Service Identity

- | A Service Principal and

- | A client certificate

- | A client secret

## Exercise 2.05 – Lifecycle in Azure

Check README file

| Note: for Windows users, graph command might generate some grey hair, therefore this step is optional

## Exercise 2.06 – Configuration drift in Azure

Check README file

# Getting Started with Terraform - Summary

- Terraform is a simple state management engine
- Resources are created in sequence dictated by the execution plan
- Terraform uses the state file during a plan

# 3. Functional Programming



# Content



Input variables, local values, and output values



Making Terraform more expressive with functions and for expressions



Incorporating two new providers: Random and Archive



Templating with templatefile()



Scaling resources with count

# Functional programming ?

- | Declarative programming paradigm
- | Aggregation of modular functions
- | Function's attributes:
  - | Pure functions
  - | First-class and high-order functions
  - | Immutability

# Procedural s Functional

## PROCEDURAL

```
const numList = [1, 2, 3, 4, 5]
let result = 0;
for(let i = 0; i < numList.length; i++) {
  if(numList[i] % 2=== 0) {
    result += (numList[i] * 10)
  }
}
```

## FUNCTIONAL

```
const numList = [1, 2, 3, 4, 5]
const result = numList
  .filter(n=>n % 2=== 0)
  .map(a=>a * 10)
  .reduce((a, b) =>a + b)
```

## TERRAFORM

```
locals {
  numList = [0, 1, 2, 3, 4, 5]
  result  = sum([for x in local.numList : 10 * x if x % 2 == 0])
}
```



# Local value

- | Assigns a name to an expression
- | Allows multiple repetition

Element declaration

```
┌  
locals {  
  ...  
}
```

**BUT**

```
local.
```

When invoked

# Input variable – Declaration

| The syntax of a variable block is

```
      Element
    {
      variable "environment" {
        ...
      }
    }
```

Diagram illustrating the syntax of a variable block. The word "Element" is above the opening curly brace. The word "Name" is below the closing curly brace of the "environment" block. The word "variable" is above the opening curly brace of the "environment" block. The word "environment" is above the closing curly brace of the "environment" block. The word "..." is below the opening curly brace of the "environment" block.

| Variable values can be accessed via

When invoked: `var.environment`

Within a string: `${var.environment}`

| `terraform apply -var variable_name="value"`

# Input variable – Arguments

- | default
- | description
- | type
  - | Primitive: string, integer, bool
  - | Complex: list, map, set, object, tuple
- | validation

```
variable "environment" {  
    default = ...  
    description = ...  
    type = ...  
    validation {}  
}
```

# Input variable – Primitive types

```
variable "environment" {  
    type = string  
    default = "prod"  
}
```

```
variable "size" {  
    type = number  
    default = 123  
}
```

```
variable "refresh_pwd" {  
    type = bool  
    default = false  
}
```

## Exercise 3.01 – Azure Resource Group Name

Check README file

# Input variable – Collection types

```
variable "roles" {  
    type = list(string)  
    default = ["admin", "user"]  
}
```

} role = var.roles[0]

```
variable "plans" {  
    type= map  
    default = {  
        "basic" = "1xCPU-1GB"  
        "heavy" = "1xCPU-2GB"  
    }  
}
```

} sizing = var.plans["basic"]  
or  
size = lookup(var.plans, "basic")

# Input variable – Structural types

```
variable "user" {  
    type = object({  
        login = string          # a required attribute  
        name  = optional(string) # an optional attribute  
    })  
}
```

# Validating Variables

| Validation block

```
variable "login" {  
    type = string  
  
    validation {  
        condition = (var.login == var.login)  
        error_message = "This is an error message."  
    }  
}
```



# Validating Variables - Example

```
validation {  
    condition = (length(var.login) >= 8)  
    error_message = "Login does not match expected length."  
}
```

```
validation {  
    condition = (length(var.accepted_envs) <= 3)  
    error_message = "Nbr of environment is too high."  
}
```

```
validation {  
    condition = length(var.rgName) >= 6 && substr(var.rgName, 0, 3) == "rg-"  
    error_message = "Must start with a 'rg-' and contains at least 6 chars."  
}
```

...

## Exercise 3.02 – Validation rule

Check README file

# Functions

- | Terraform functions are expressions that transform inputs into outputs.
- | Restricted to built-in functions.
- | Terraform extension is done through customised provider

Function name  
`function_name( param1, param2, ... )`  
Param(s)

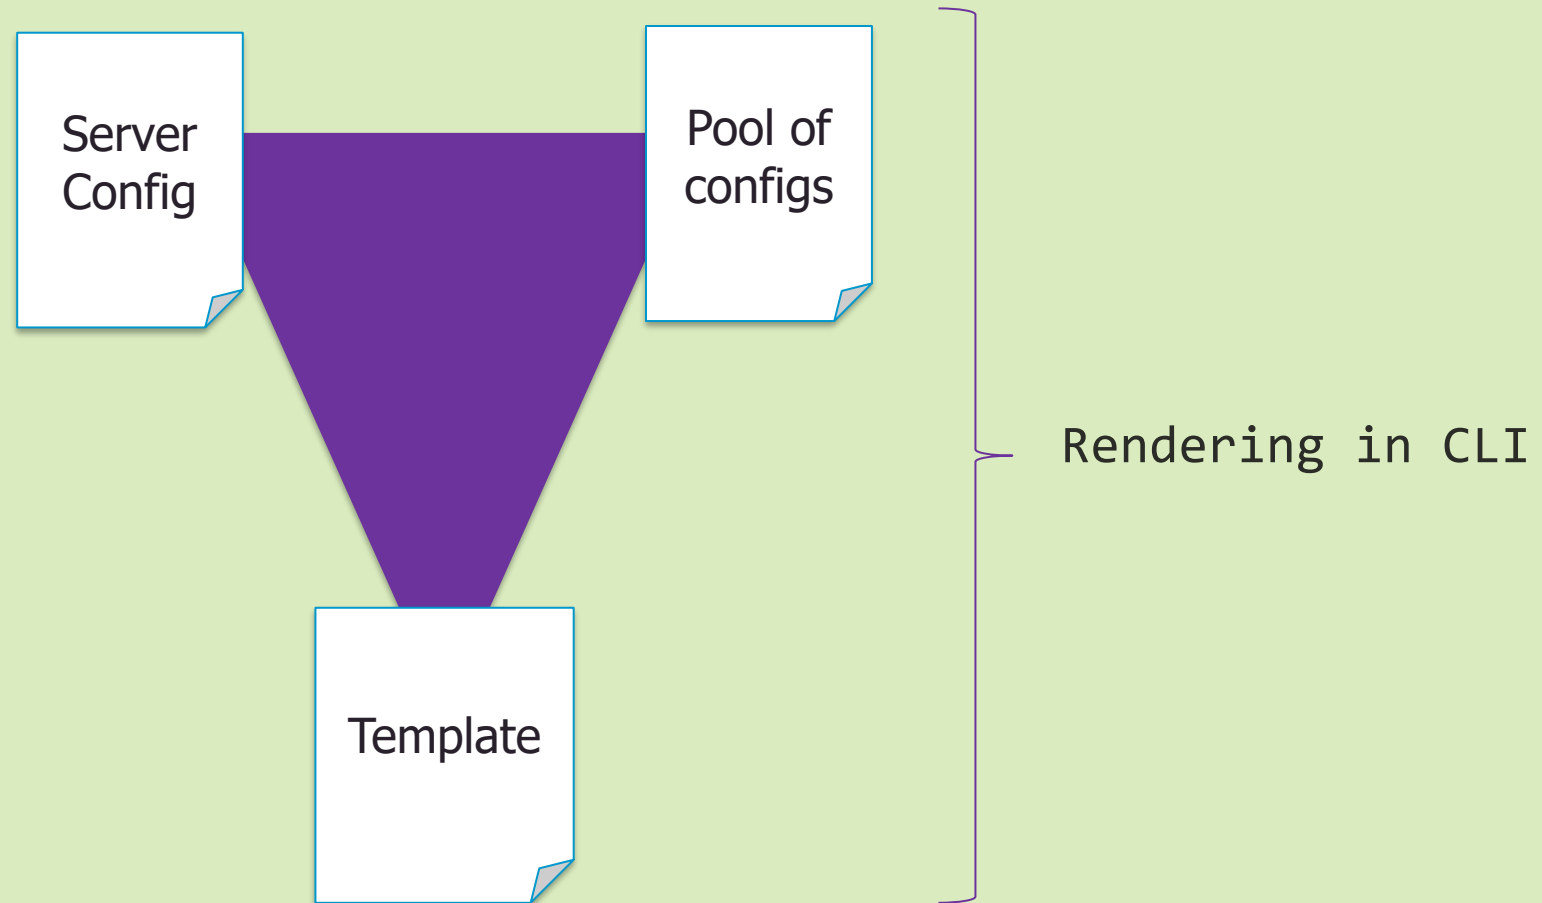
`length(...)`  
`contains(..., ...)`

# Function – templatefile()

| templatefile() used to replace placeholder values in a template file

Function name  
templatefile( "templates/configs.txt" , {os = ["ubuntu", "windows"] ... } )  
Path  
Templates variables


## Example 3.03 – Description



# Server Config(s)



## TYPE COERCION



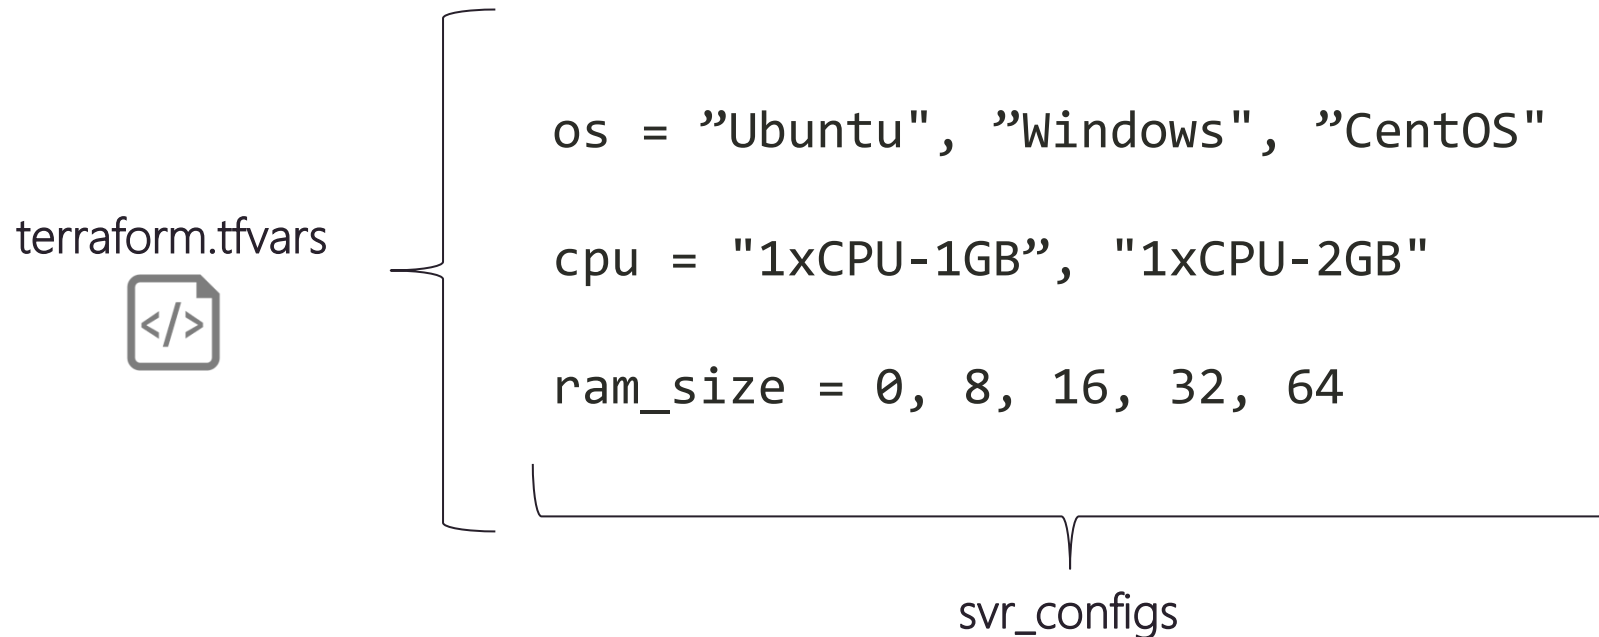
```
terraform {  
  required_version = ">= 0.15"  
}
```

```
variable "svr_configs" {  
  description = "A list of svr config"  
  type = object ({  
    os = list(string),  
    cpu = list(string),  
    ram_size = list(number),  
  })  
}
```

**Exercise – Add a validation rule to check that there is at least 3 ram\_sizes**

# Assigning Values with a Variable Definition File

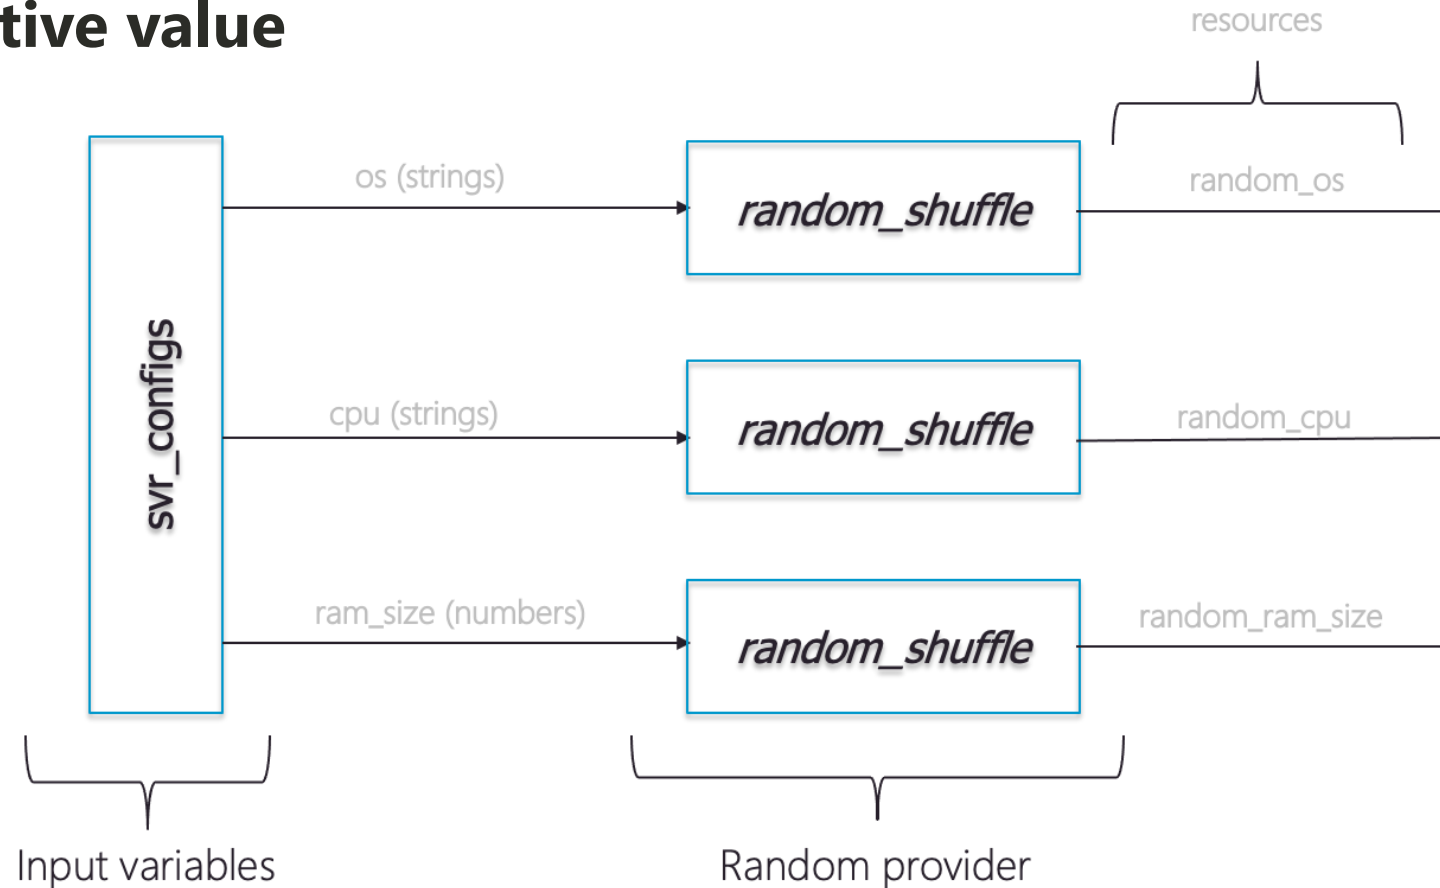
| File ending in either .tfvars or .tfvars.json



| **Exercise – Create a pool of configs file**

# Server Config

**Exercise – Create the resources and set the ‘input’ attribute with the respective value**



```

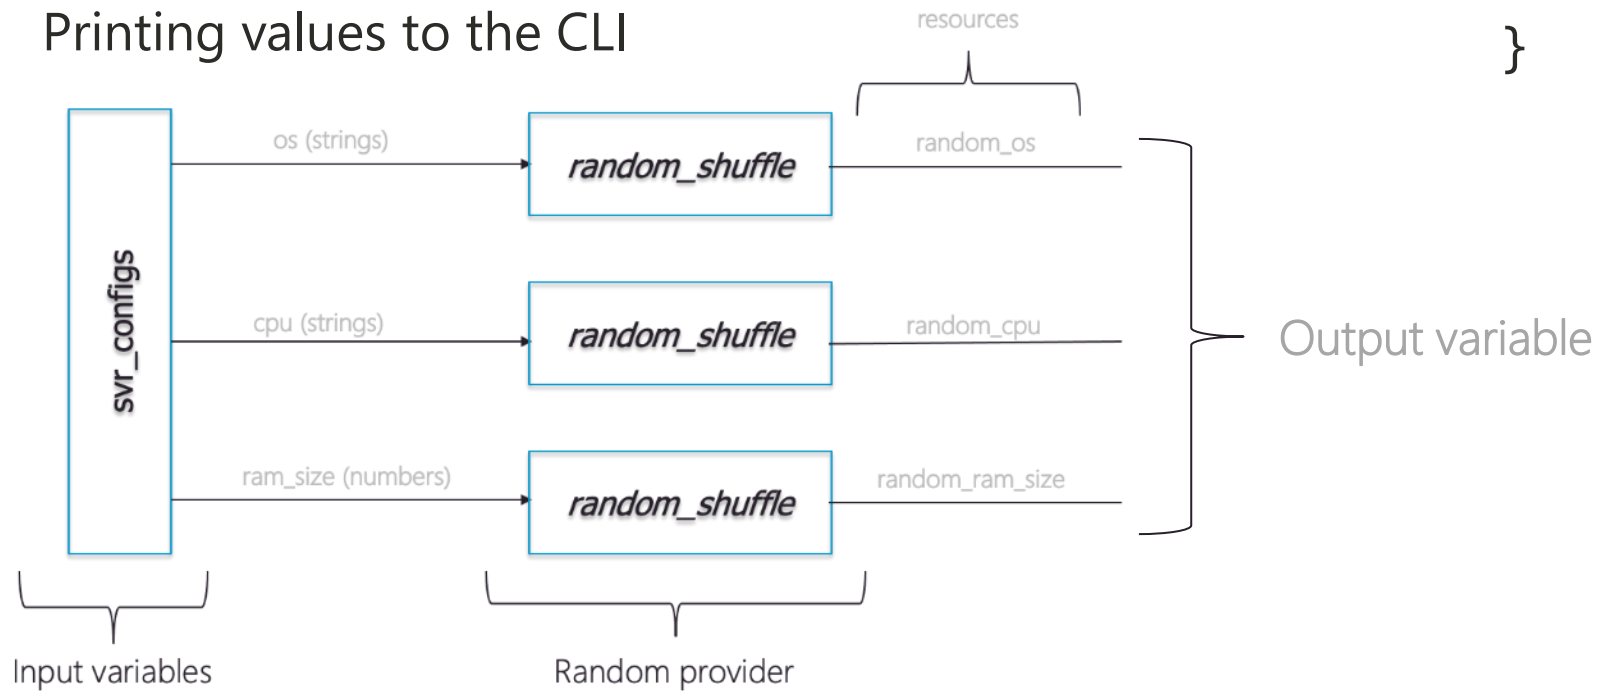
Identifier
├── Element
├── Type
├── Name
└── Attribute
  {
    resource "docker_image" {
      name = "nginx:latest"
      keep_locally = true
    }
  }
  
```



# Output variable

Output values are for doing one of two things:

- Passing values between modules
- Printing values to the CLI



Element

```

output "environment" {
  ...
}

```

Name

**Exercise – Create the output variable**

# So far

```
terraform {  
  required_version = ">= 0.15"  
}  
  
variable "svr_configs" {  
  description = "A list of server config"  
  type = object ({  
    ...  
  })  
  validation {  
    ...  
  }  
}  
  
resource "random_shuffle" "random_os"{  
  input = ...  
}  
  
resource "random_shuffle" "random_cpu"{  
  input = ...  
}  
  
resource "random_shuffle" "random_ram_size"{  
  input = ...  
}  
  
output "out_cfg" {  
}
```

Server  
Config

Pool of  
configs

Template

```
svr_configs = {  
  os = [ "ubuntu", "windows", "centos" ]  
  cpu = [ "1xCPU-1GB", "1xCPU-2GB" ]  
  ram_size = [0, 8, 16, 32, 64]  
}
```

# Template file

- | Terraform syntax based

- | **Exercise – Create a new directory called templates**

- | **Exercise – In this directory, create a “typical\_svr.json” file representing a typical server: os, cpu, ram\_size**

# So far

```

terraform {
  required_version = ">= 0.15"
}

variable "svr_configs" {
  description = "A list of server config"
  type = object ({
    ...
  })
  validation {
    ...
  }
}

resource "random_shuffle" "random_os"{
  input = ...
}

resource "random_shuffle" "random_cpu"{
  input = ...
}

resource "random_shuffle" "random_ram_size"{
  input = ...
}

output "out_cfg" {
  value = ...
}

```

Server  
Config

Pool of  
configs

```

svr_configs = {
  os = [ "ubuntu", "windows", "centos" ]
  cpu = [ "1xCPU-1GB", "1xCPU-2GB" ]
  ram_size = [0, 8, 16, 32, 64]
}

```

Template

```

{
  "operating_system": "${os[0]}",
  "cpu_config": "${cpu[0]}",
  "ram_config": ${ram_size[0]}
}

```

# Random provider

**Exercise – Declare the required provider within the terraform element**

```
required_providers {  
    random= {  
        source= "hashicorp/random"  
        version = "~> 3.0"  
    }  
}
```

# Let's roll

| terraform init && terraform apply -auto-approve

| What happens if you apply twice ? Why ?

## Example 3.04 – Generating multiple configs

- | Comment output element
- | Create 2 additional templates files

# Local file – Prepare

- | Read all the template files from the template folder into a set

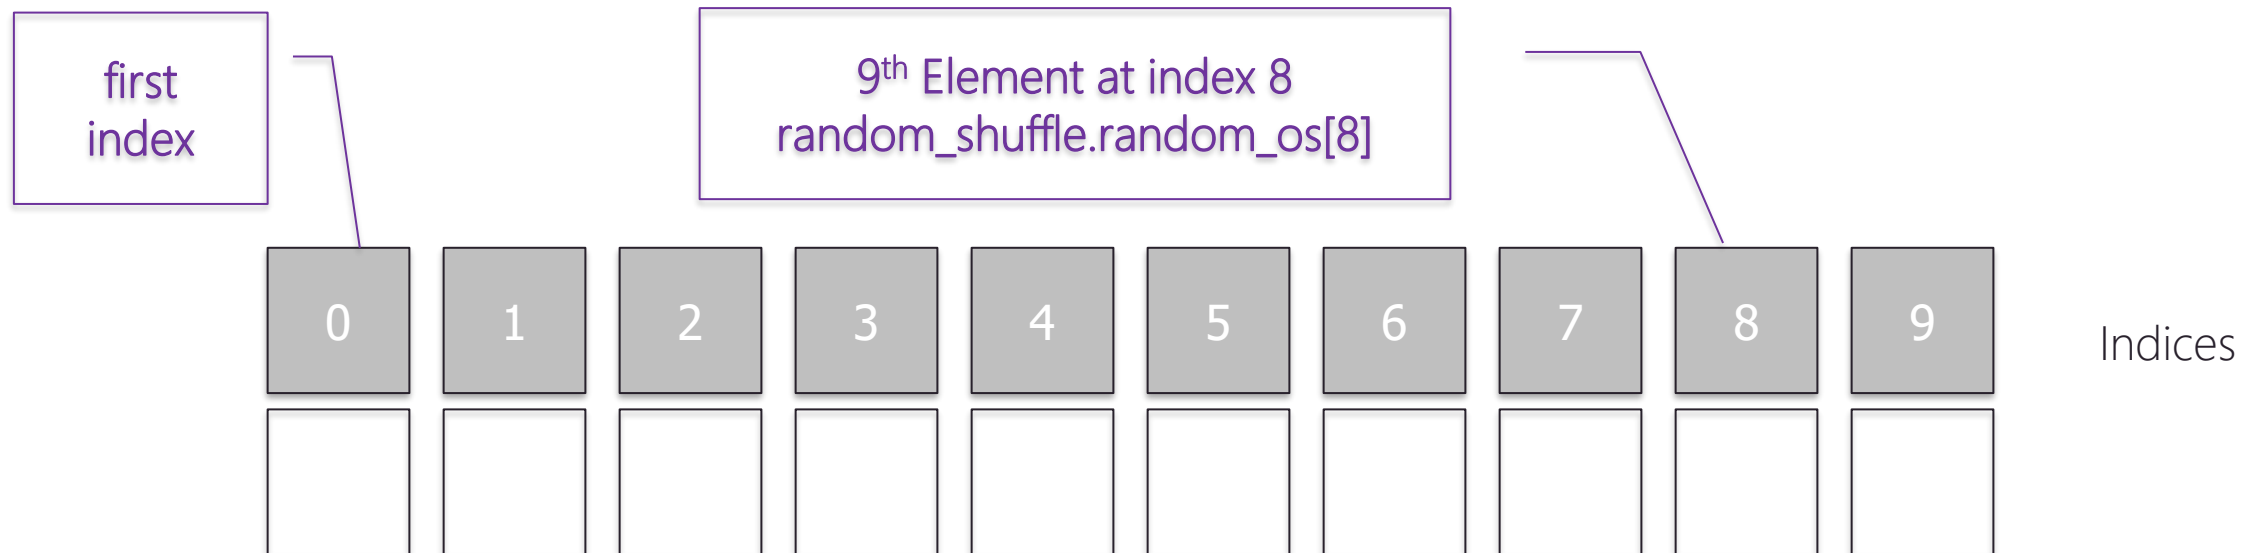
```
fileset(path.module, "templates/*.json"))
```

- | Convert set into a list and declare it in locals

```
locals {  
    templates = tolist(fileset(path.module, "templates/*.json"))  
}
```



# Counter parameter



**Exercise – Add a new variable named `var.num_files` having type number and a default value of 10**

**Exercise – Reference this variable to dynamically set the count meta argument on each of the `shuffle_resources` & the file name**

# Archive the files

## | Exercise – Declare the required provider within the terraform element

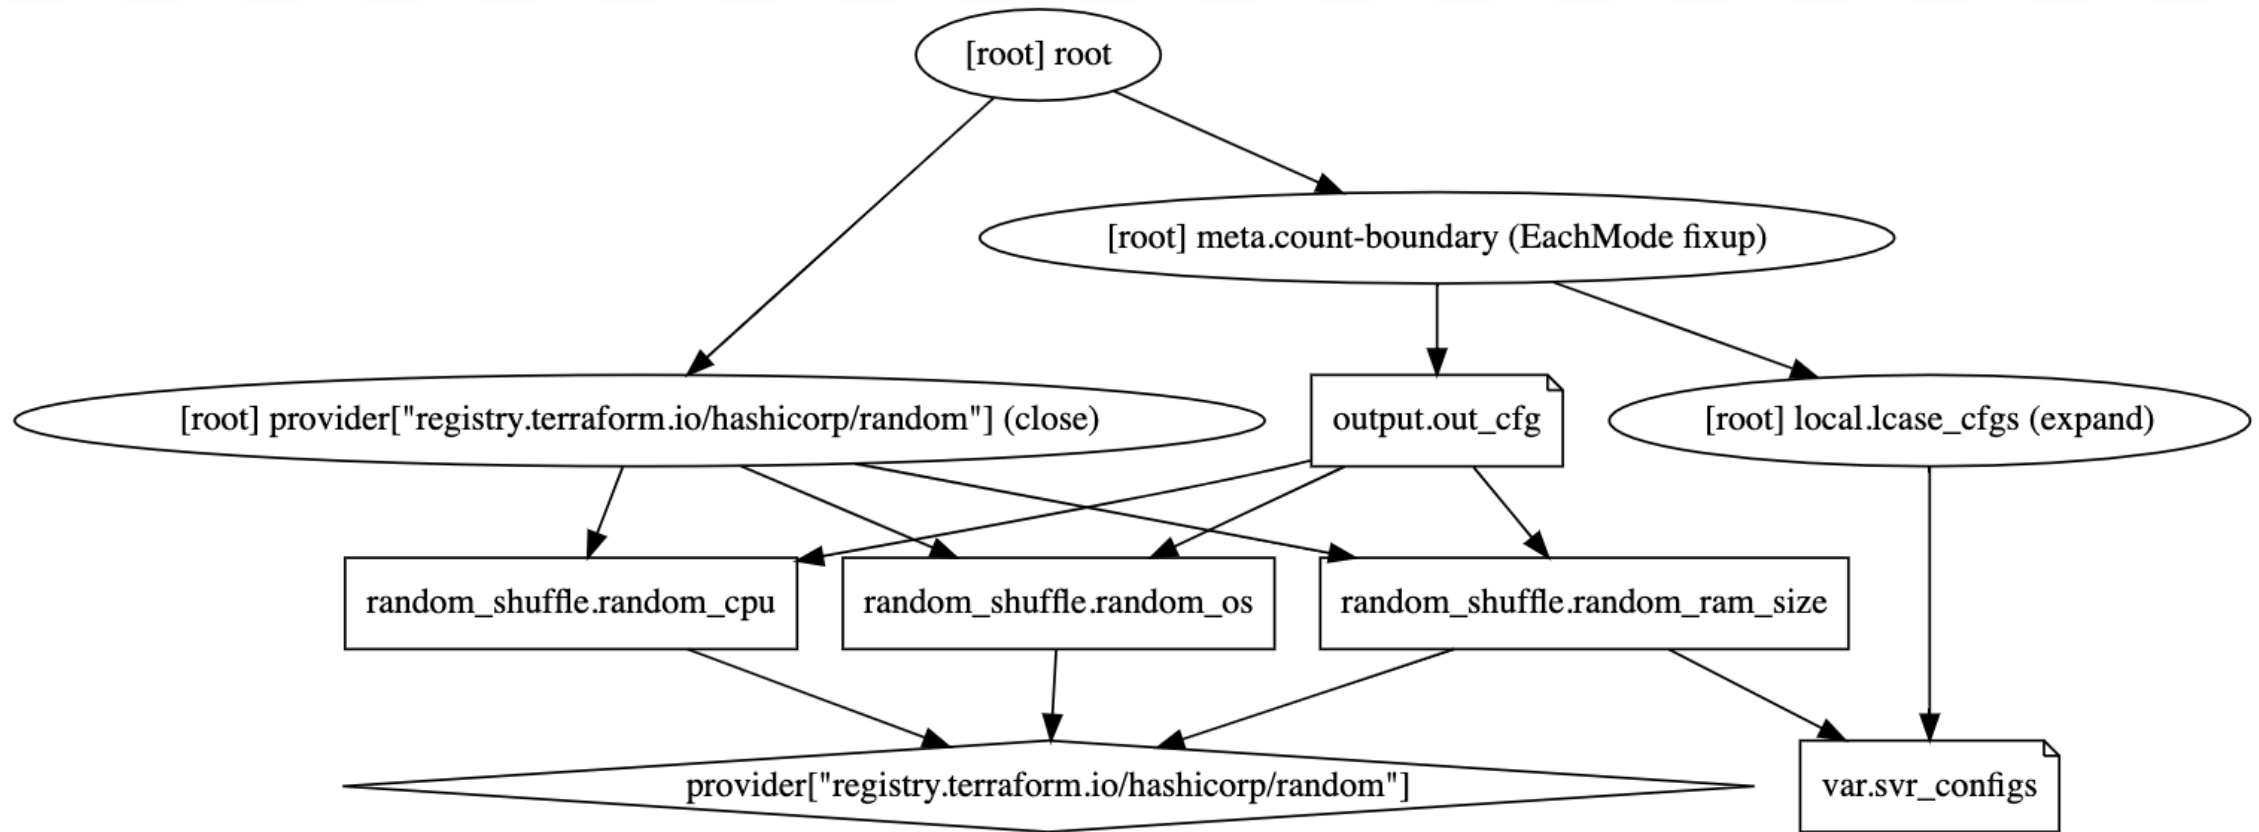
```
required_providers {  
    archive= {  
        source= "hashicorp/archive"  
        version = "~> 2.0"  
    }  
}
```

# Archive the files – cont'd

## | Exercise – Declare a data element representing the archive file


```
data "archive_file" "zip_cfg" {  
  type = "zip"  
  source_dir = "${path.module}/configs"  
  output_path = "${path.cwd}/config.zip"  
  depends_on = [local_file.out_cfg]  
}
```

# Note: Implicit dependencies



# Note: Conditional expressions

| Conditional expressions hurt readability a lot, so avoid using them if you can



```
locals {  
    v = length(var.svr_configs["os"])>=1 ? var.svr_configs["os"] : [][0]  
}
```

## Exercise 3.05 – Conditional expression

Check README file

# So far

- | Input variables, local values, output values
- | For expressions
- | Randomness must be constrained
- | Zip at runtime – explicit dependency
- | `templatefile()`
- | Count meta argument

# Terraform expressions

Name	Description	Example
Conditional Expressions	Use the value of a boolean expression to select one of two values	<code>condition ? true_value : false_value</code>
Function Calls	Transform and combine values	<code>&lt;FUNCTION NAME&gt;(&lt;ARG 1&gt;, &lt;ARG2&gt;)</code>
For Expression	Transform one complex type to another	<code>[for s in var.list : upper(s)]</code>
Split Expressions	Shorthand for some common use cases that could otherwise be handled by for expressions	<code>var.list[*].id</code> equivalent for expression: <code>[for s in var.list : s.id]</code>
Dynamic Blocks	Construct repeatable nested blocks within resources	<pre>dynamic "ingress" {   for_each = var.service_ports   content {     from_port = ingress.value     to_port = ingress.value     protocol = "tcp"   } }</pre>
String Template Interpolation	Embed expressions in a string literal	<code>"Hello, \${var.name}!"</code>
String Template Directives	Use conditional results and iteration over a collection within a string literal	<code>%{ for ip in var.list.*.ip } server \${ip} %{ endfor }</code>



# 4. Application Deployment



# Content



Deploying a multi-tiered web application in Azure with Terraform



Setting project variables in a variable's definition files



Organizing code with nested modules

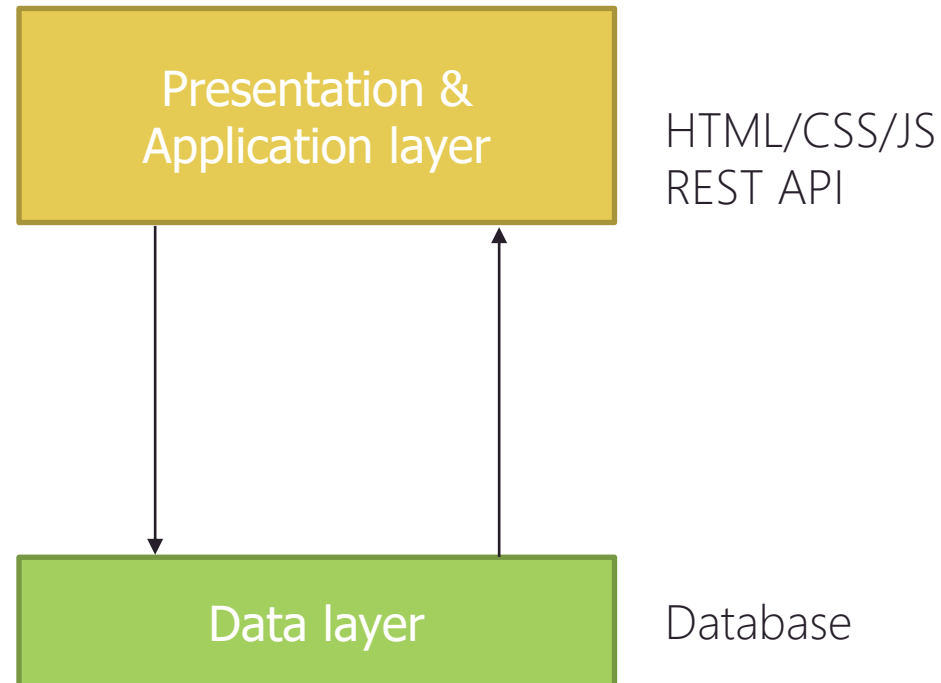


Leveraging modules from the public module registry

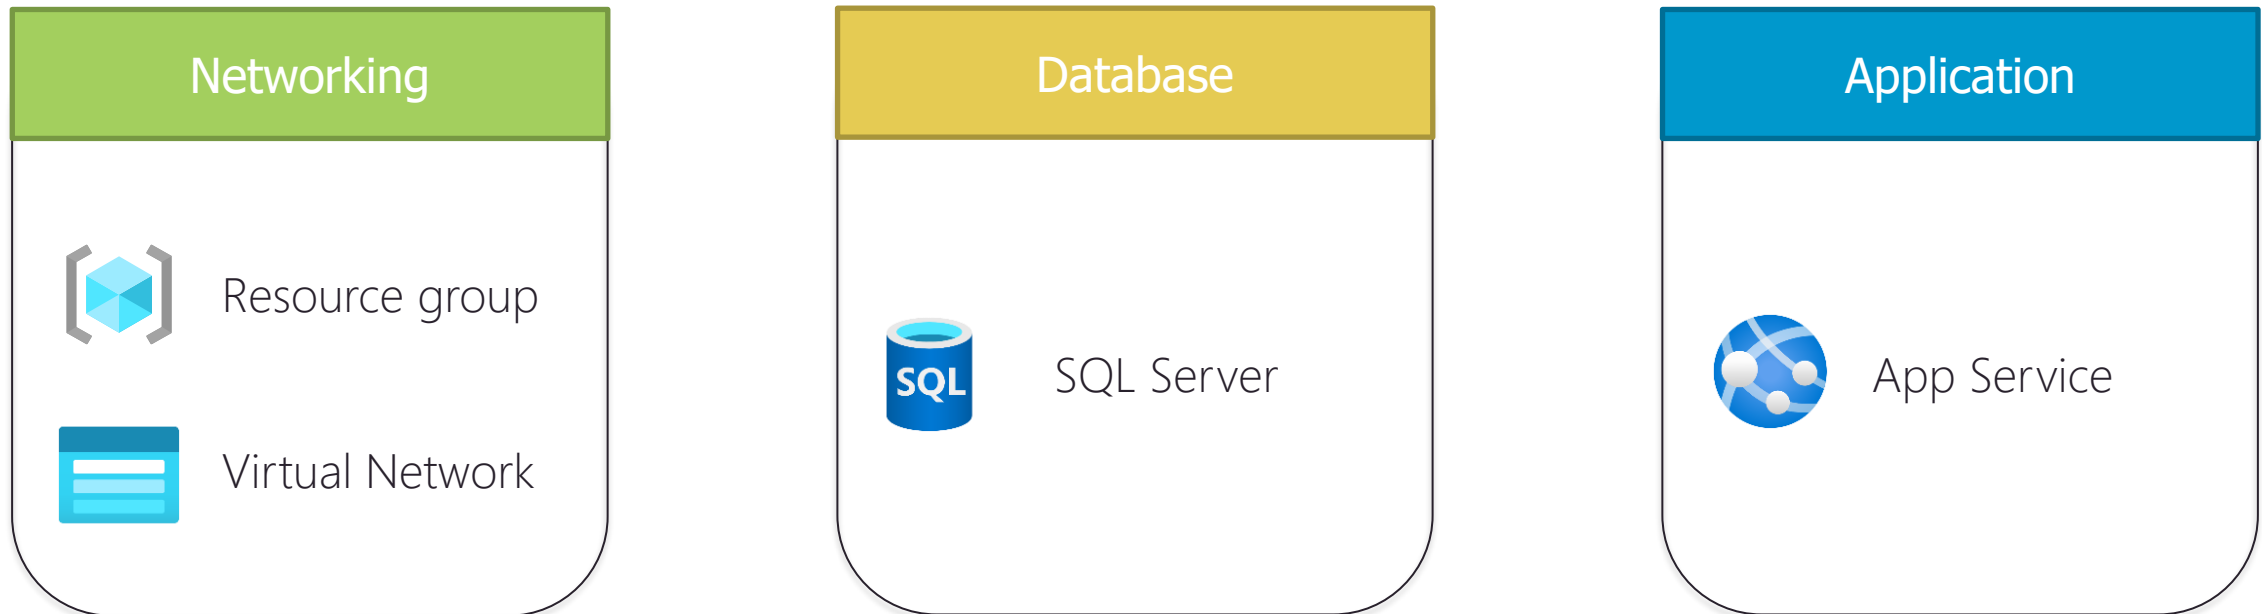


Passing data between modules using input variables and output values

# Multi-Tiered Web Application in Azure



# Architecture



# Terraform Module

## | Definition

- | Self-contained packages of terraform code
- | Consume inputs, produces outputs
- | Allow/Use for code reuse and software abstraction

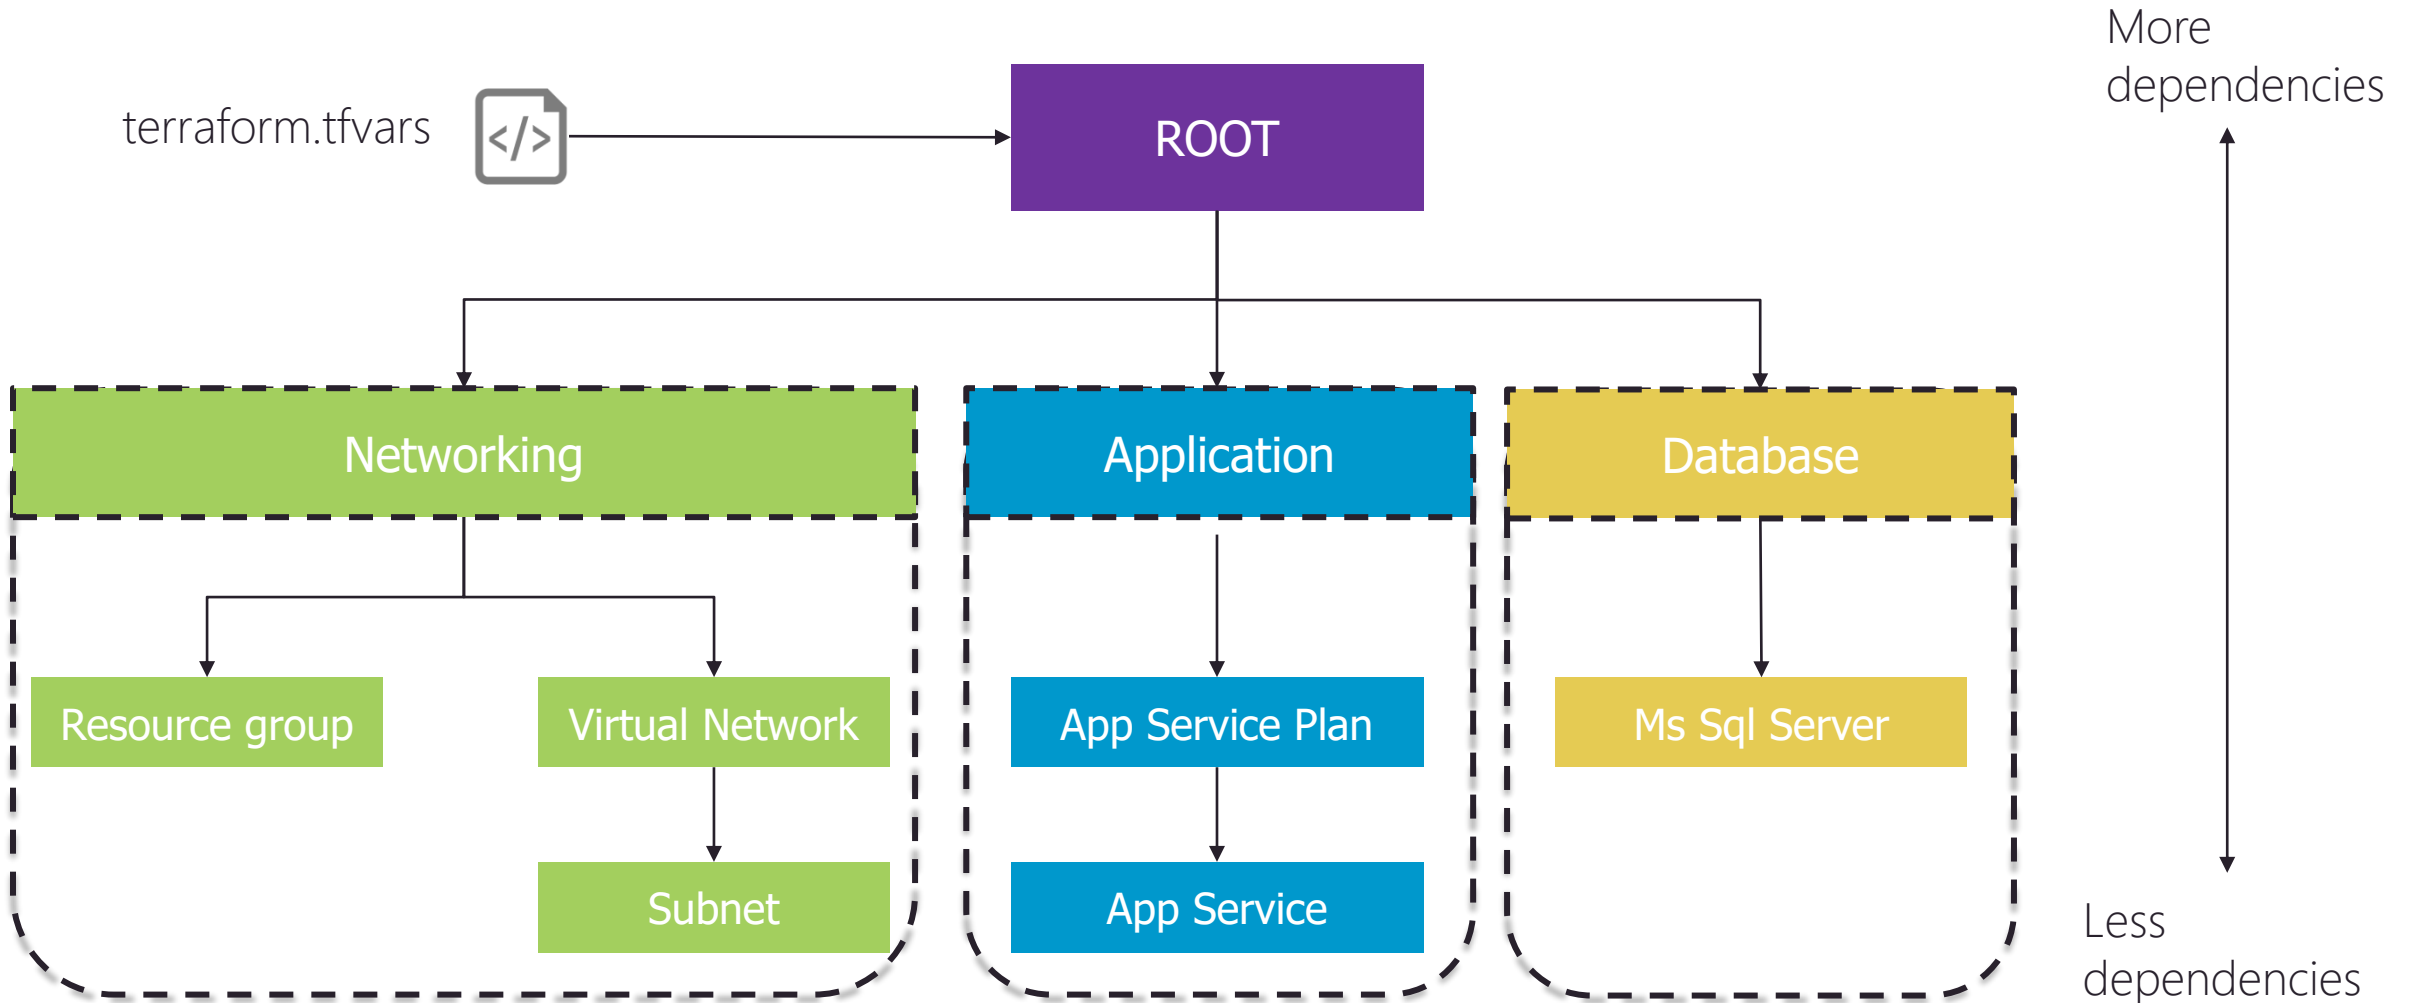
## | Syntax

```
module "networking" {  
    source = "path"  
    version = "..."  
    arg_one = arg_value  
}
```

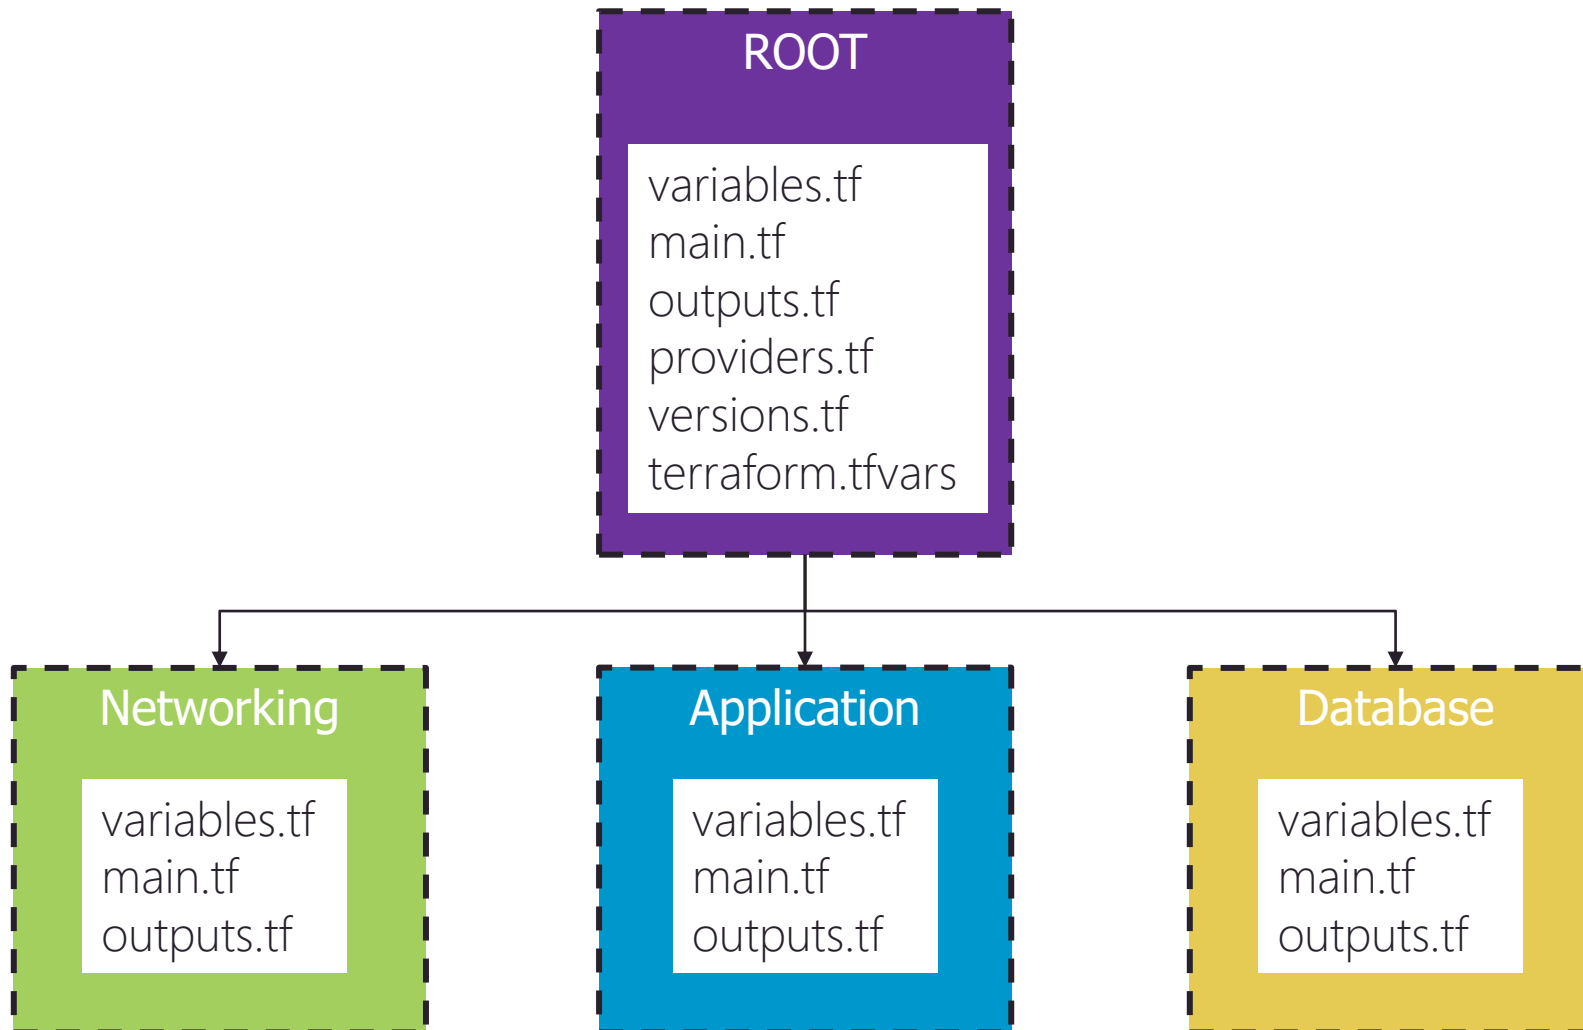
Diagram annotations:

- module name**: points to the module identifier "networking".
- meta args**: points to the `source` and `version` arguments.
- module's input variables**: points to the `arg_one` argument.

# Approach by Terraform



# Recommended structure



## Exercise 4.01 – Define root module

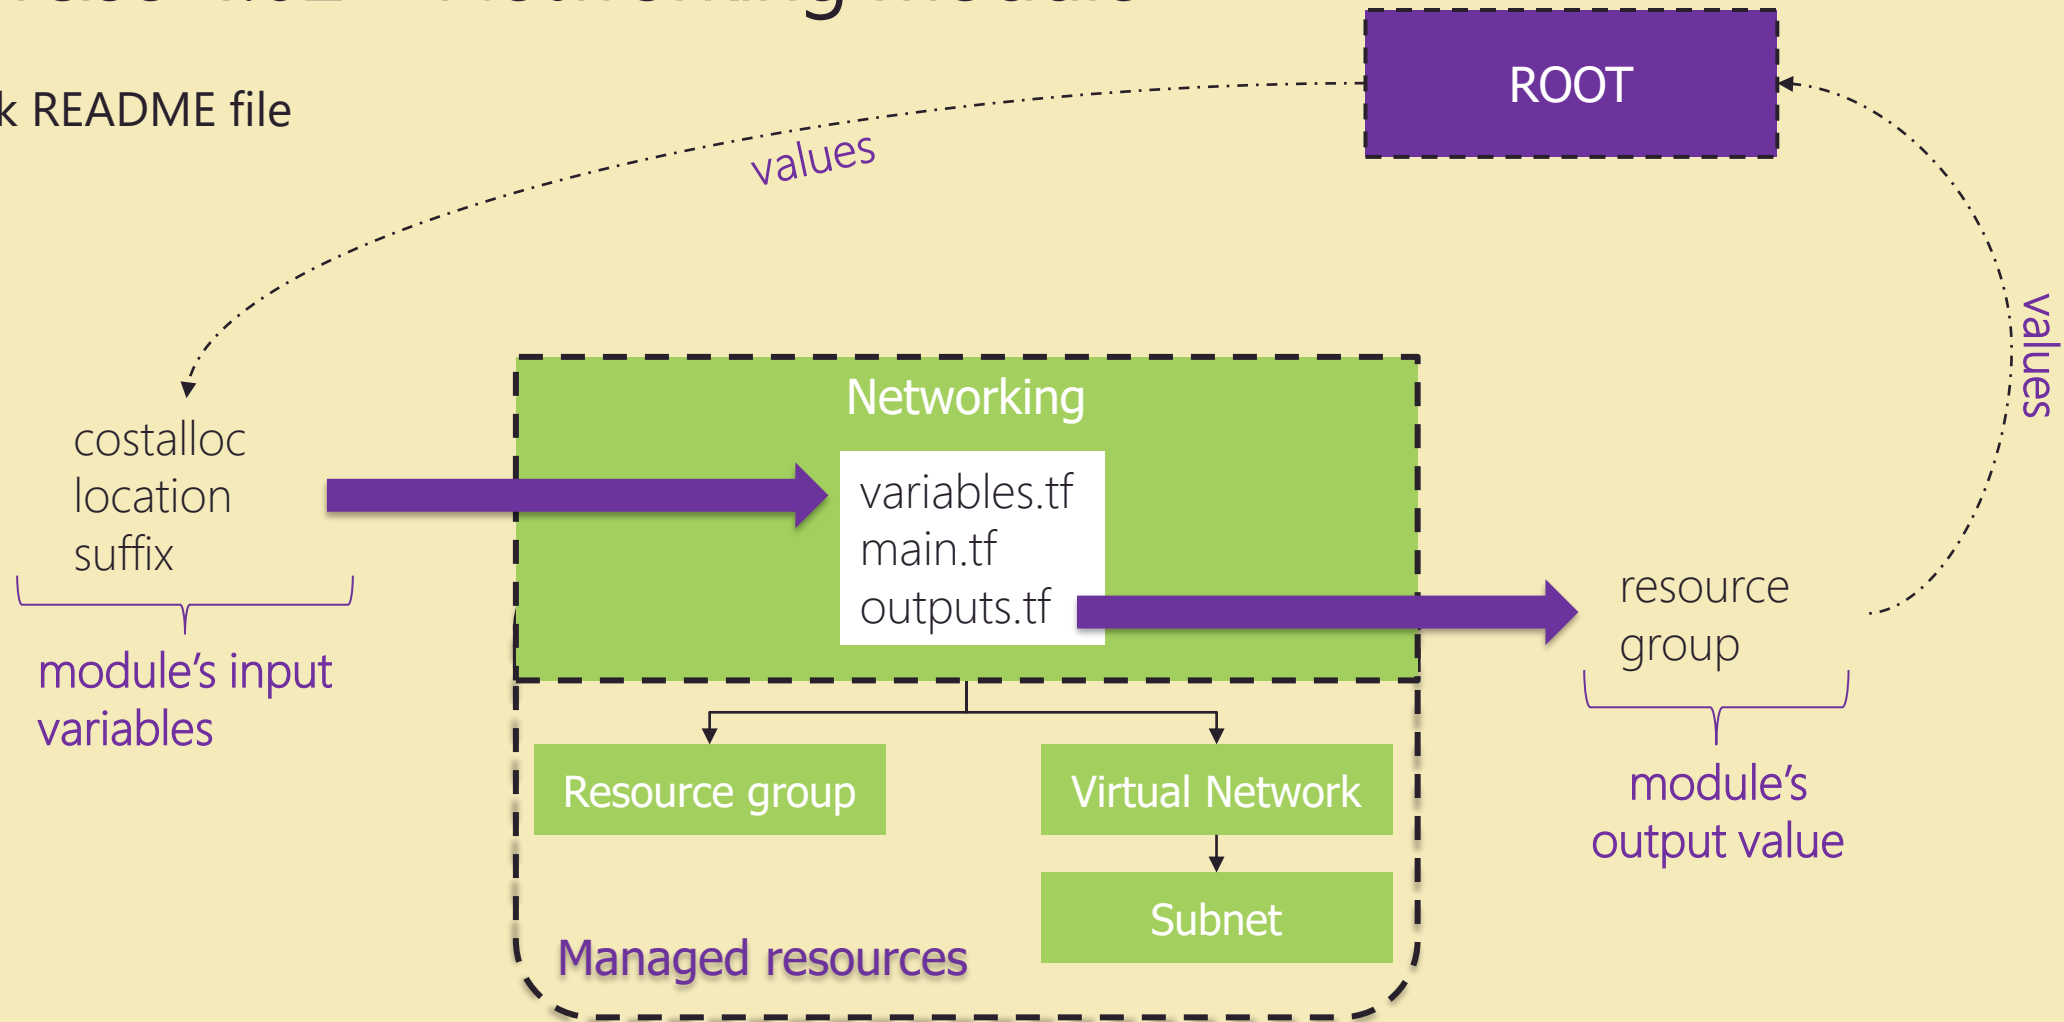
Check README file



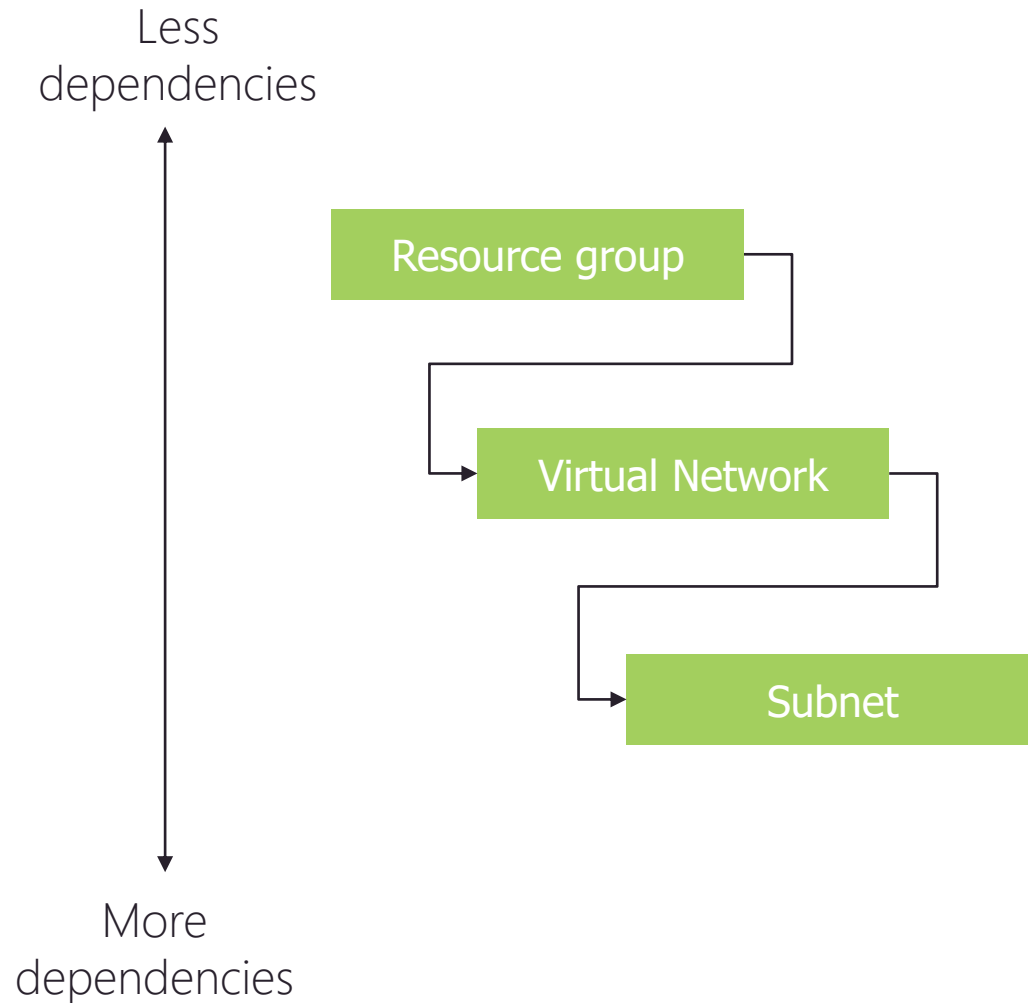


## Exercise 4.02 – Networking module

Check README file



# Networking module – Structure



# Example 4.03 – Software Componentization

## ROOT

```
variables.tf  
main.tf  
providers.tf  
versions.tf  
terraform.tfvars
```

| Existing module

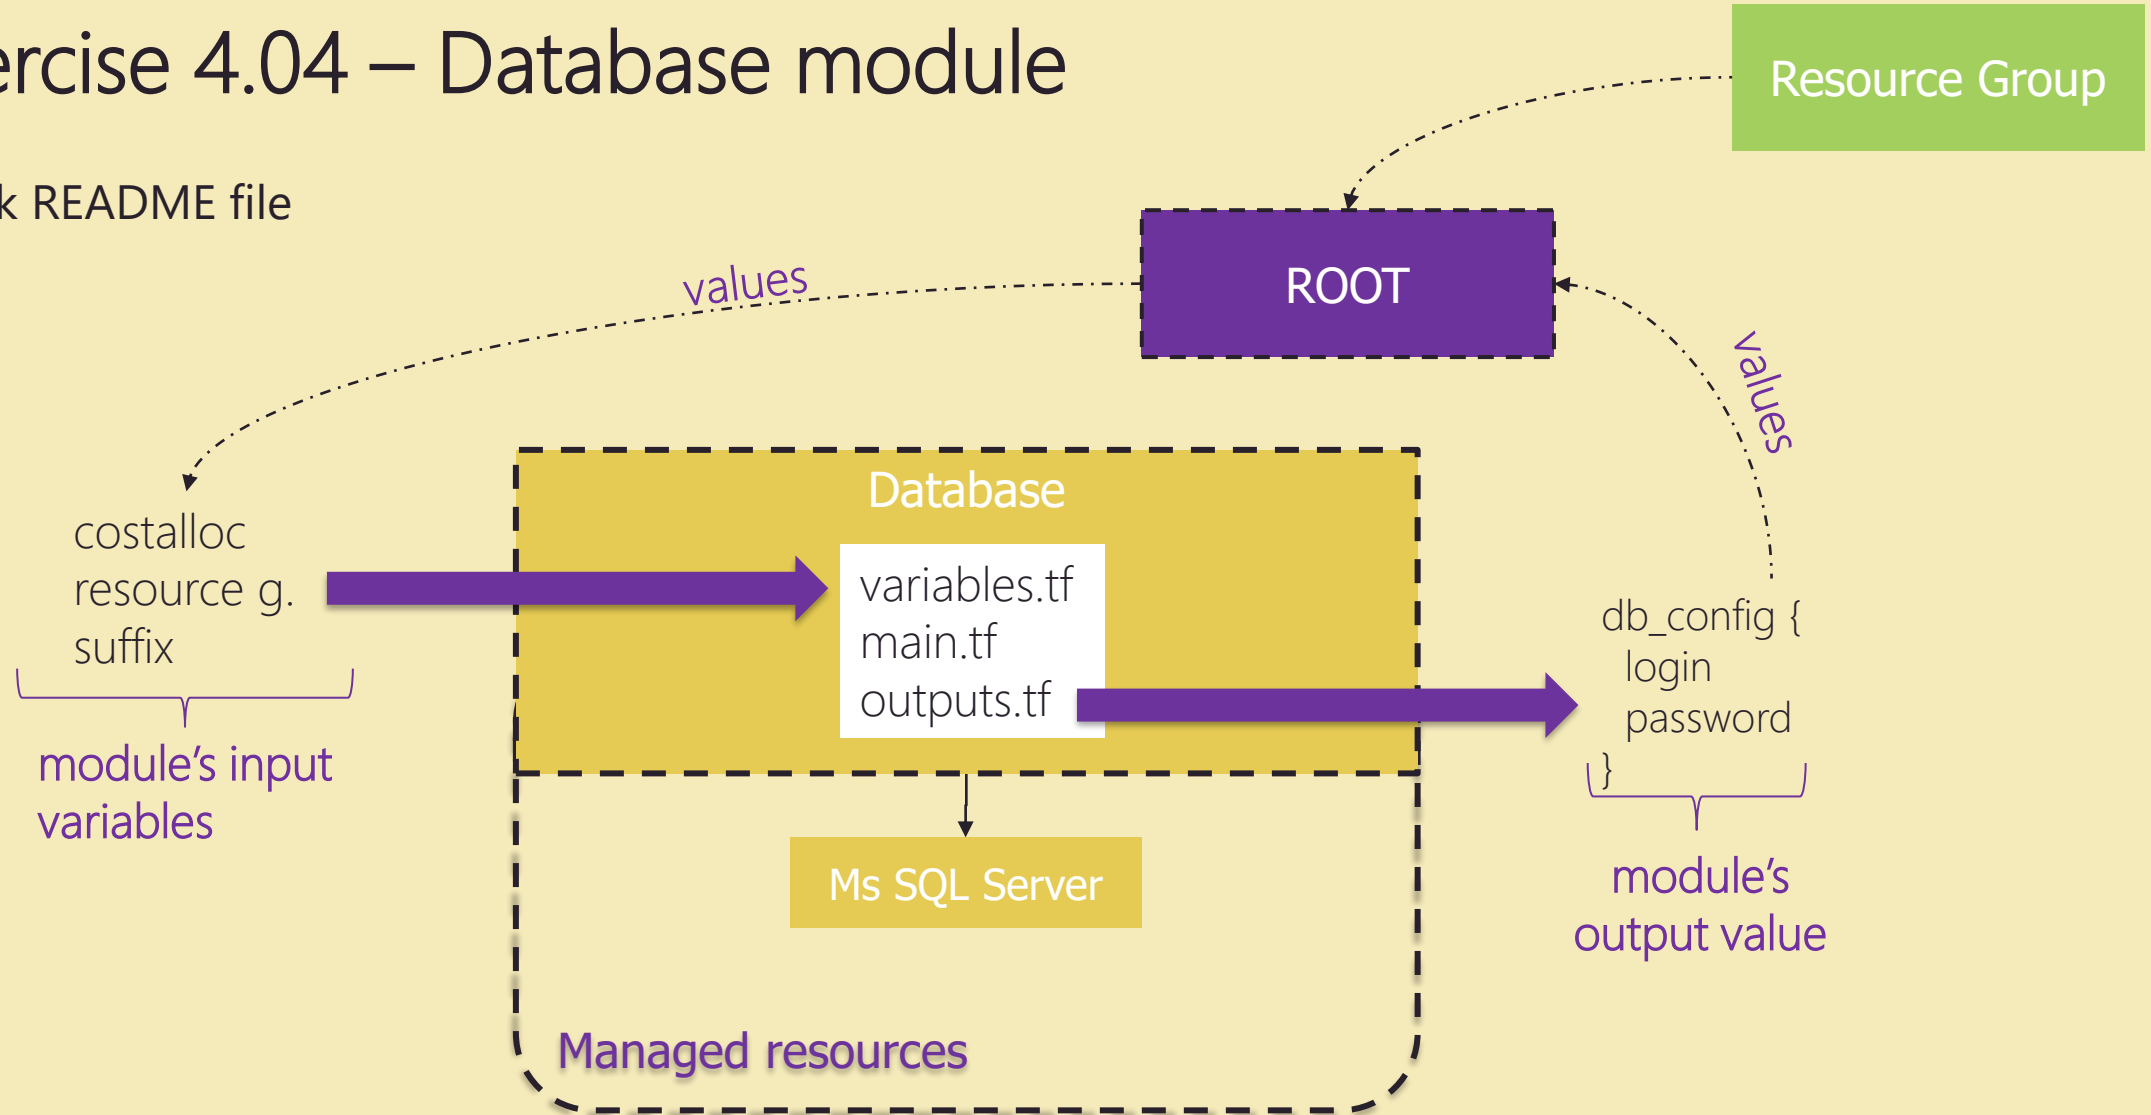
| [Azure/vnet/azurerm | Terraform Registry](#)

| Source code (see in Section 5):

| [Azure/terraform-azurerm-vnet: Terraform module to create/provision Azure vnet \(github.com\)](#)


## Exercise 4.04 – Database module

Check README file



# Database module – Ex 4.04 (bonus)

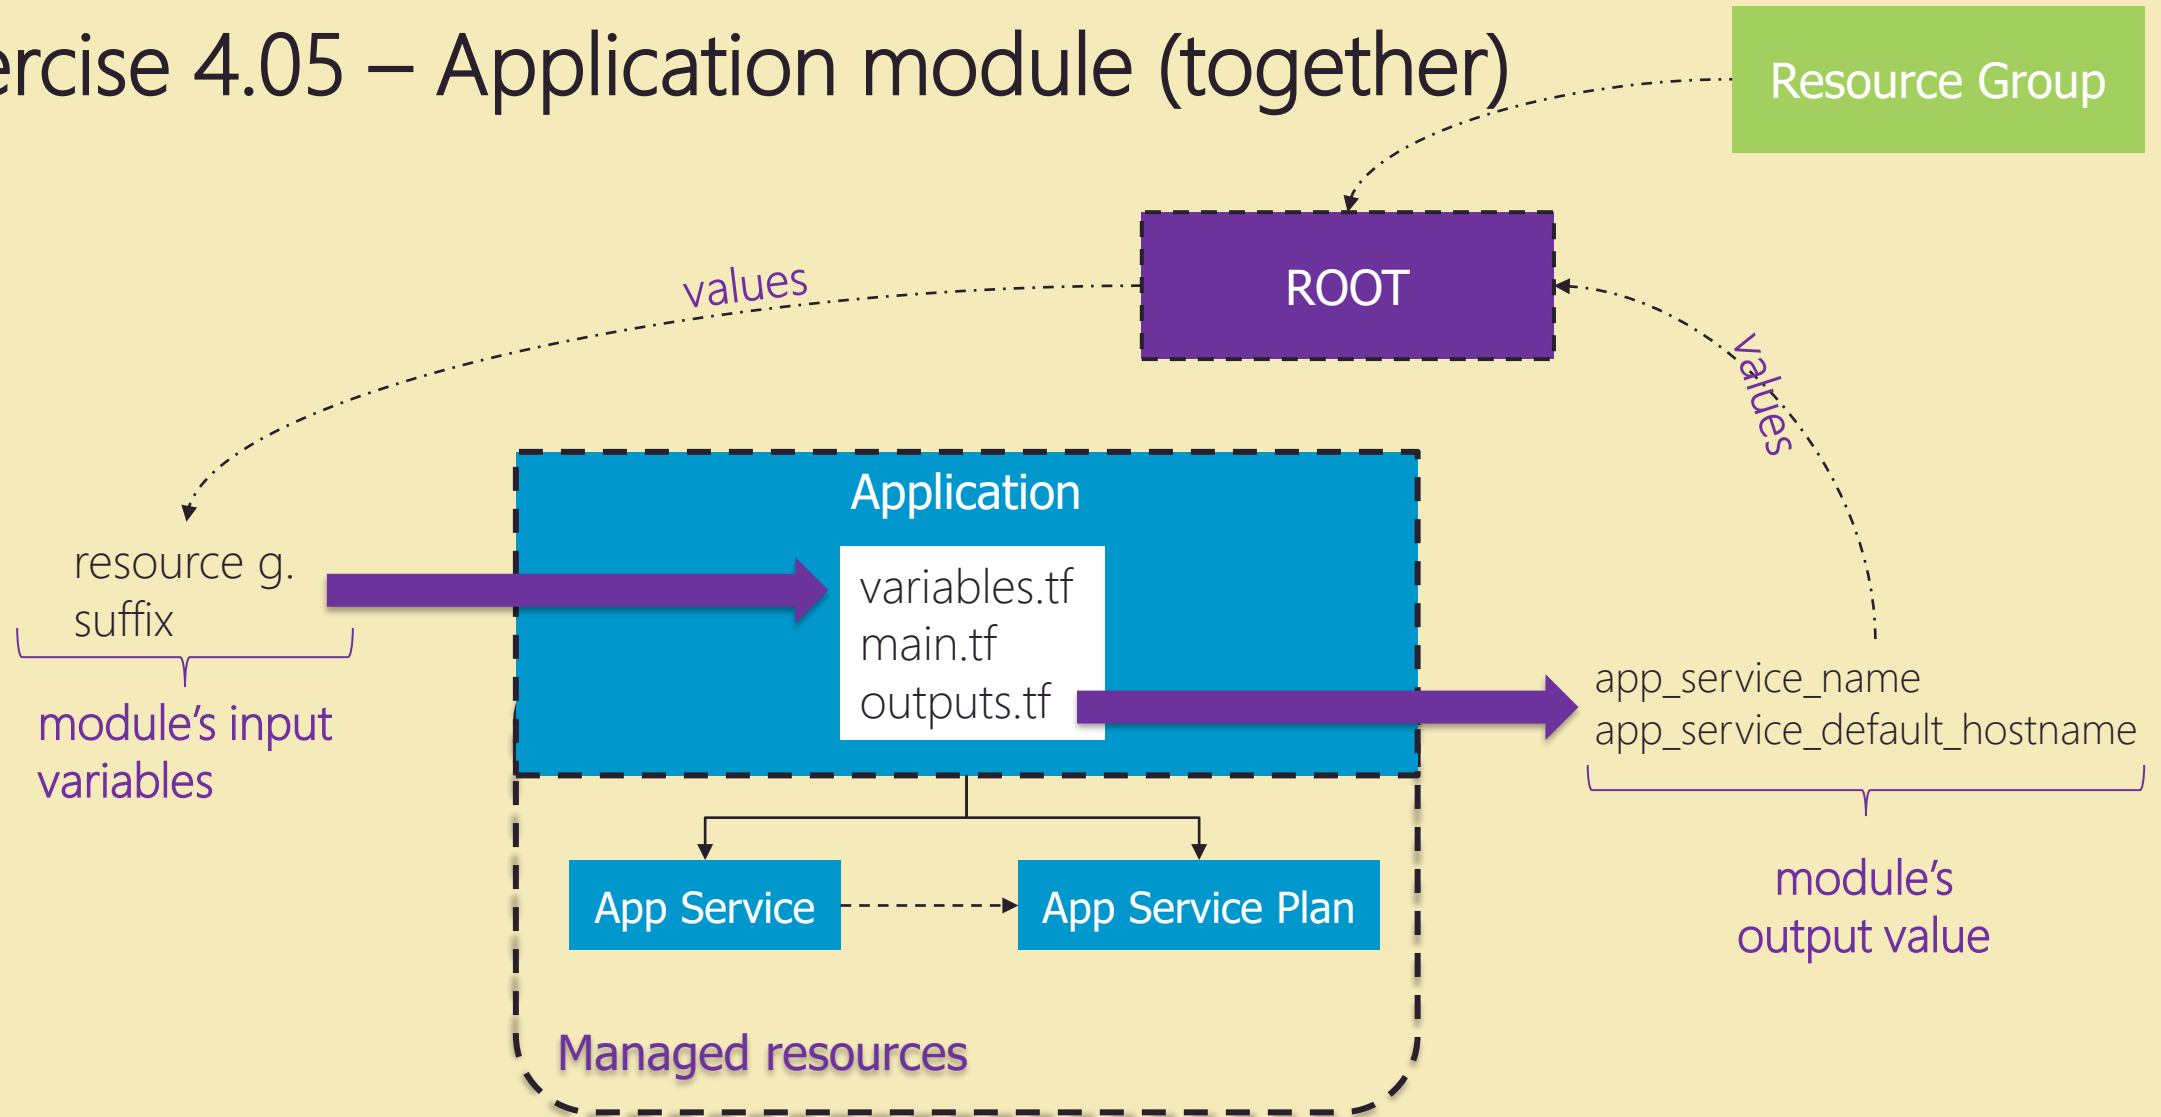
- | Add an MsSQL database
- | Add a storage account
- | *Check in Terraform Registry*



# Display output values

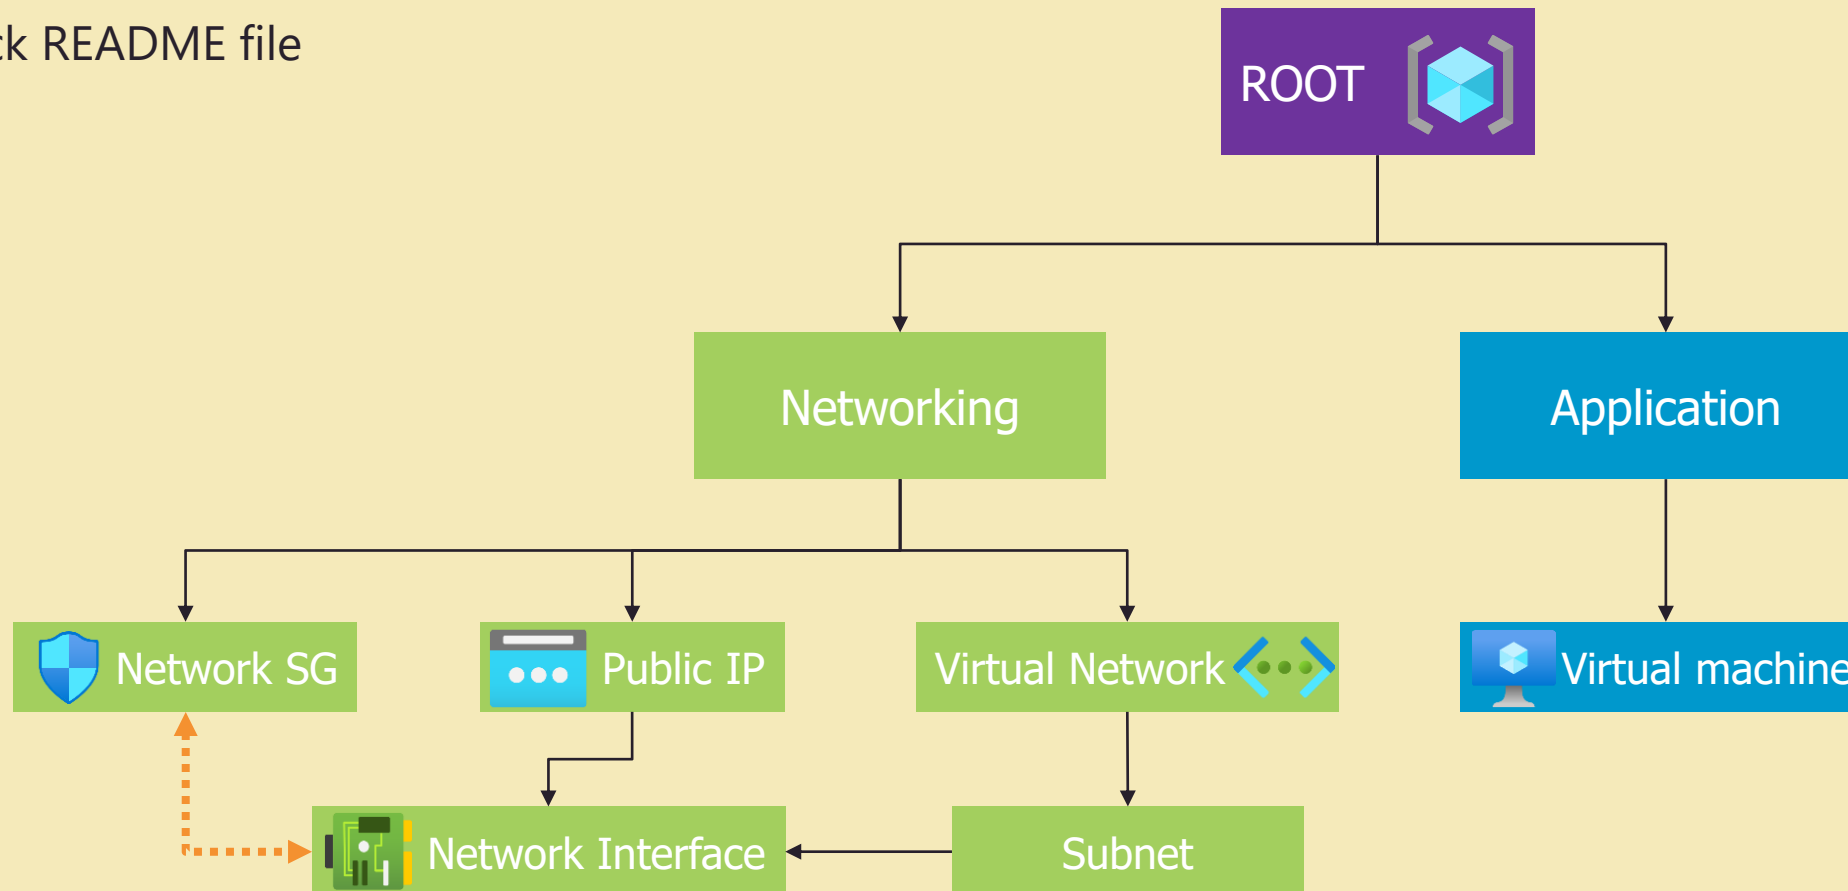
```
| terraform output «output value»
```

## Exercise 4.05 – Application module (together)



## Exercise 4.06 – Typical infrastructure

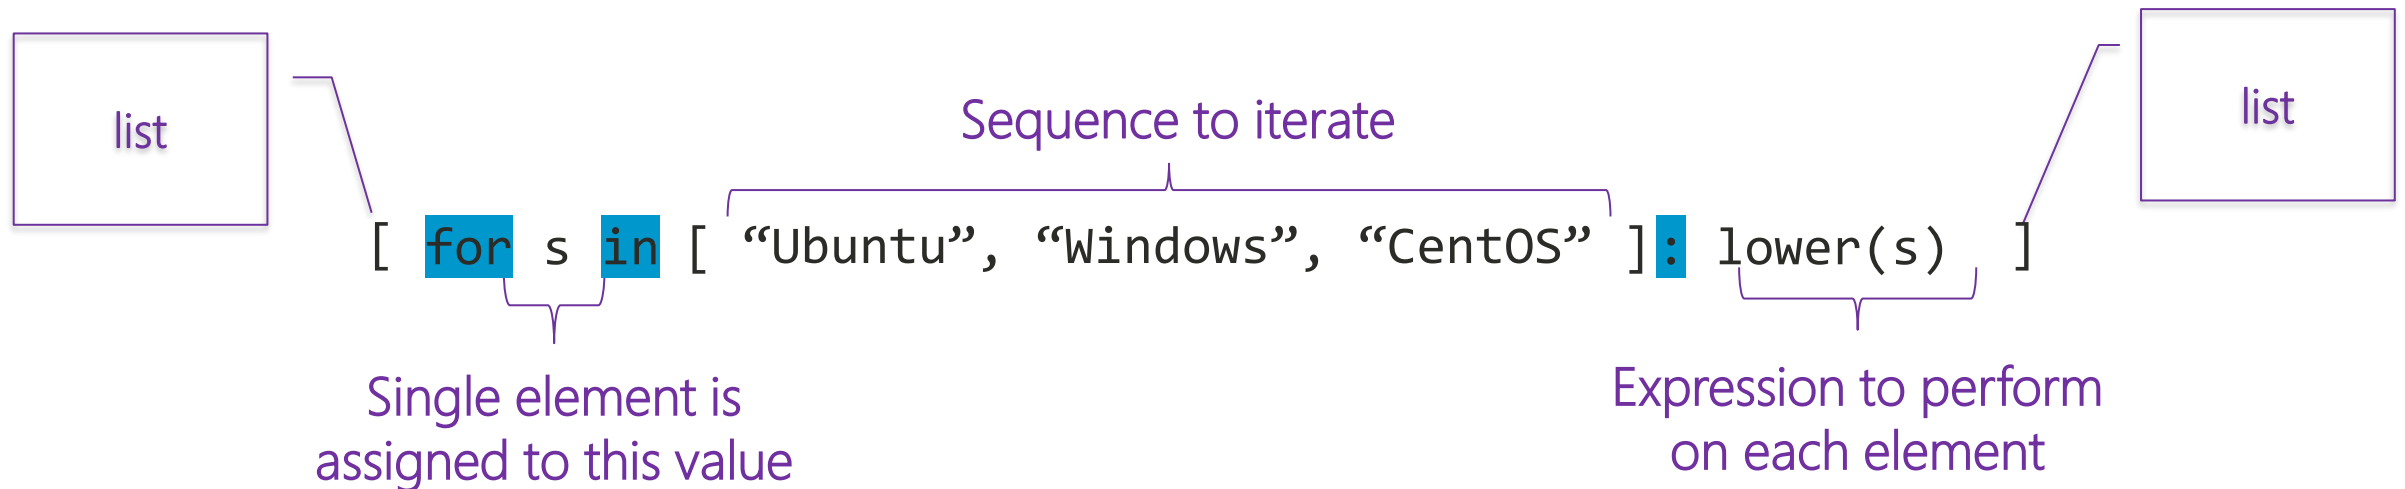
Check README file



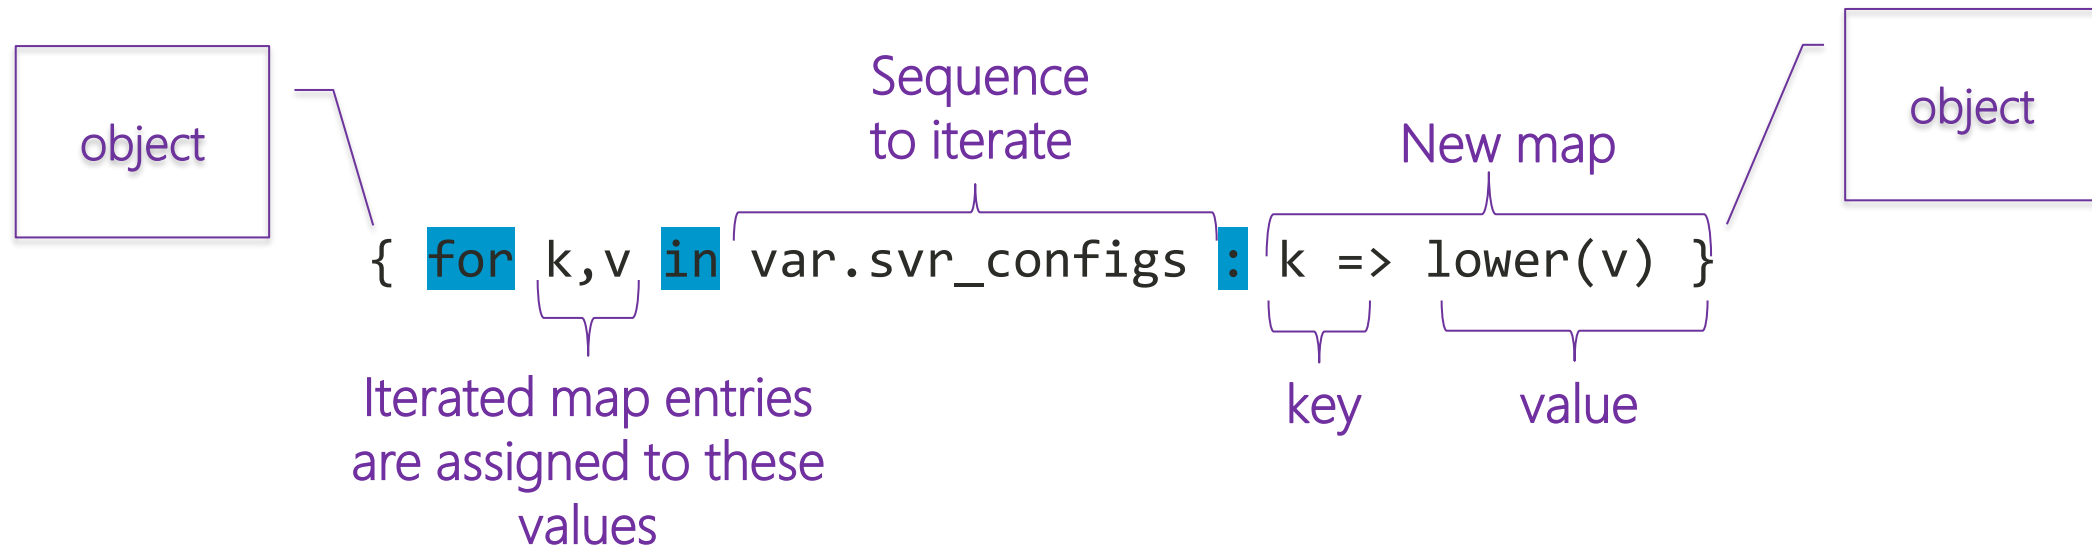


# "For" Expressions

- | Are anonymous functions that can transform one complex type into another
- | Simple expressions can be composed to construct higher-order functions
- | They use lambda like syntax and are comparable to lambda expressions and streams



# "For" Expressions – cont'd



# "For" Expressions – cont'd

```
{ for k,v in var.svr_configs : k => v }
```

[ for s in v

k	v
os	Ubuntu, windows, ...

{ for k,v in

v	s
Ubuntu, Windows, ...	ubuntu, windows, ...
1xCPU-1GB, 1xCPU-2GB	1xcpu-1gb, 1xcpu-2gb
0, 8, 16, 32, 64	0, 8, 16, 32, 64

s) ] }

# For-each + each.key/value

| For\_each can be defined within a resource configuration

- | Accepts as input either a map, or a set of strings
- | Outputs an instance for each entry in the data structure

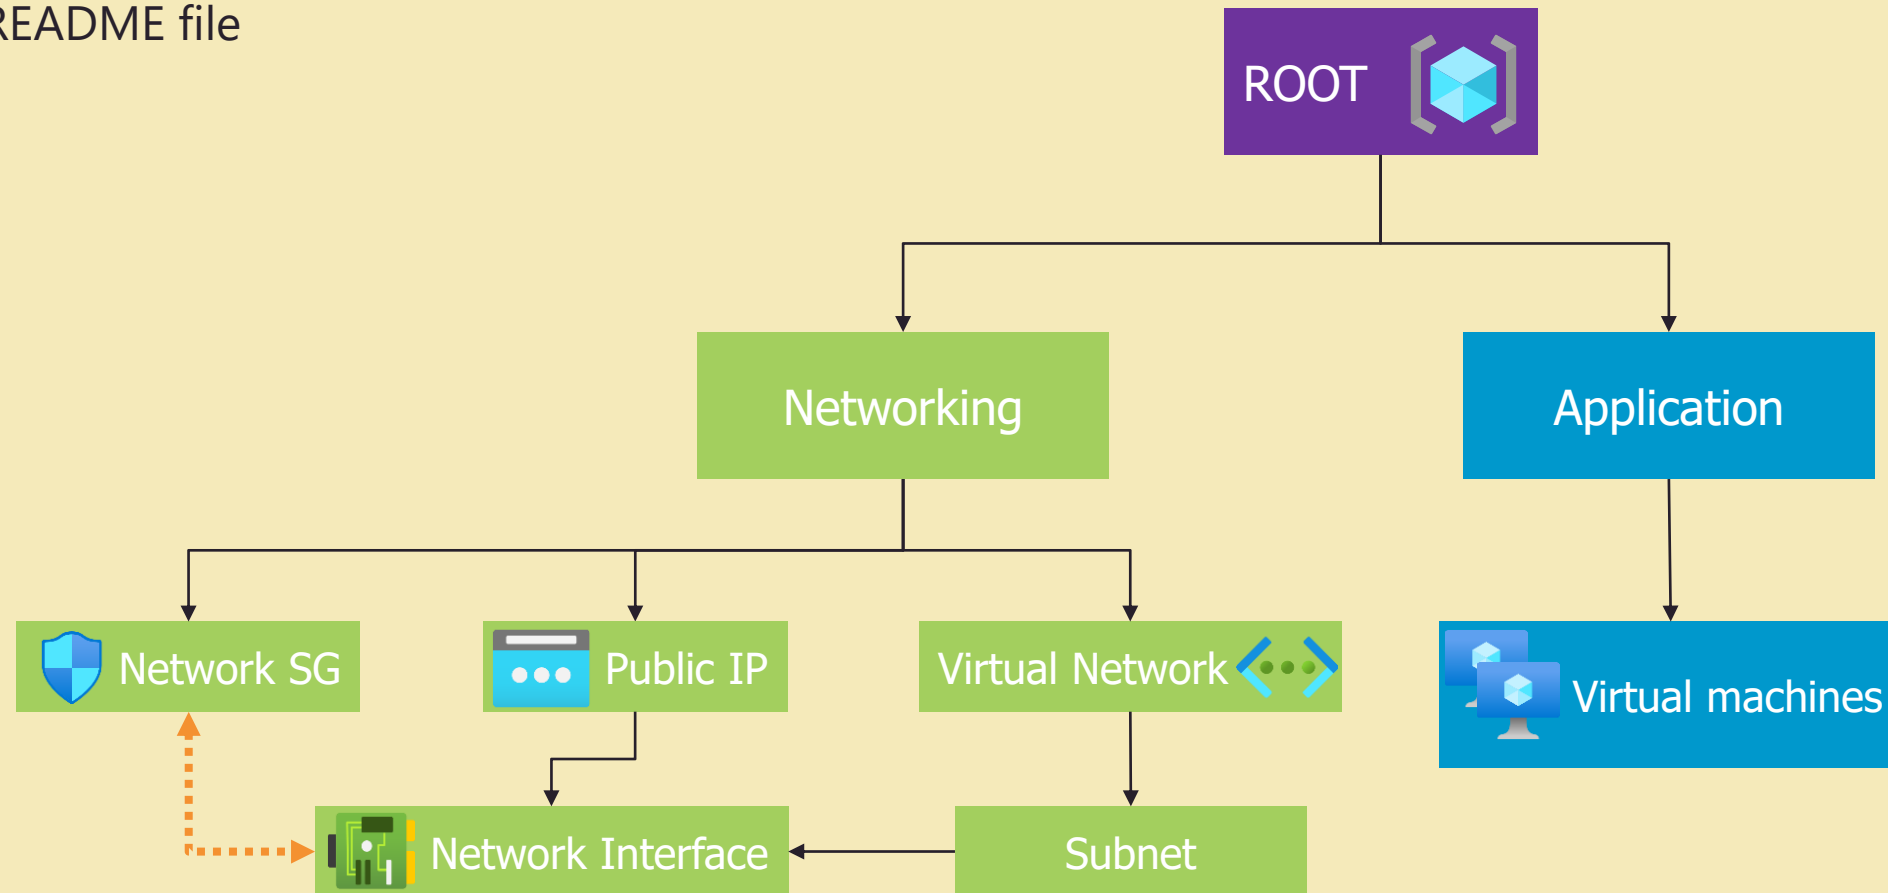
| Benefits

- | Intuitive
- | Less verbose
- | Ease of use

```
resource "azurerm_public_ip" "training" {  
  for_each = toset(var.listOfMachines)  
  name = "vm-${each.value}"  
  ...  
}  
...  
azurerm_public_ip.training[each.key].id  
...
```

## Exercise 4.07 – For each virtual machine

Check README file



# Cloud\_init\_config

- | Fresh install but still no software provisioning
- | cloud-init as a shell script

```
cloudinit = {  
    source = "hashicorp/cloudinit"  
    version = "~> 2.1"  
}
```

## Exercise 4.08 – Configure deployed server

Check README file

# Advanced config (.yaml)

```
#cloud-config
write_files:
  - path: /etc/server.conf
    owner: root:root
    permissions: "0644"
    content: |
      {
        "user": "${user}",
        "password": "${password}",
        "database": "${database}",
        "netloc": "${hostname}:${port}"
      }

runcmd:
  - curl -sL https://.../releases/latest | jq -r ".assets[].browser download url" | wget -qi -
  - unzip deployment.zip
  - ./deployment/server
...
```



# Advanced config (.yaml) – cont'd


```
data "cloudinit_config" "config" {
  gzip = true
  base64_encode = true

  part {
    content_type = "text/cloud-config"
    content = templatefile("${path.module}/config.yaml", var.infra_config)
  }
}
```

# So far

- | Complex projects are RELATIVELY easy to design & deploy thanks to TF modules
- | Root module is the main entry point
- | Nested modules is the practice of organizing code
- | Public module registry
- | Data passed using a bubble-up and trickle-down techniques

# Example 4.09 – Software Componentization



Source code

| [Github.com](https://github.com)

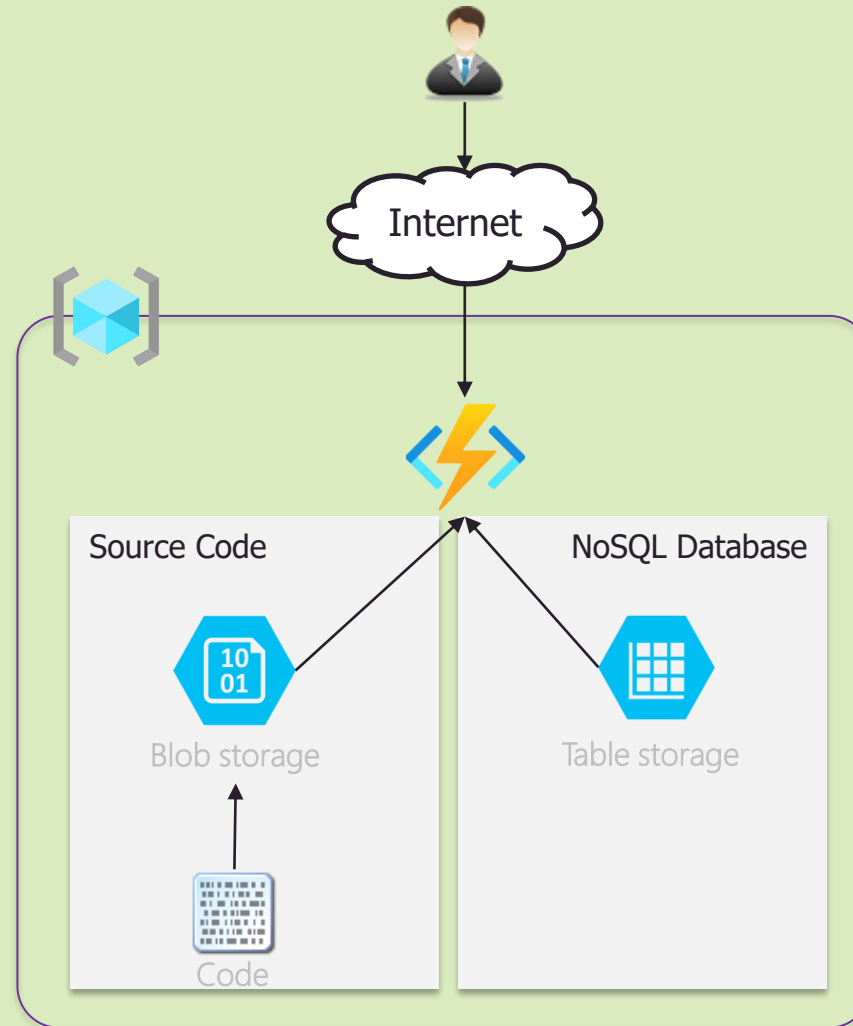
Shared module

| [Terraform Registry](https://registry.terraform.io/)

Main.tf

| module "tweetish"

# Example 4.10 – Serverless

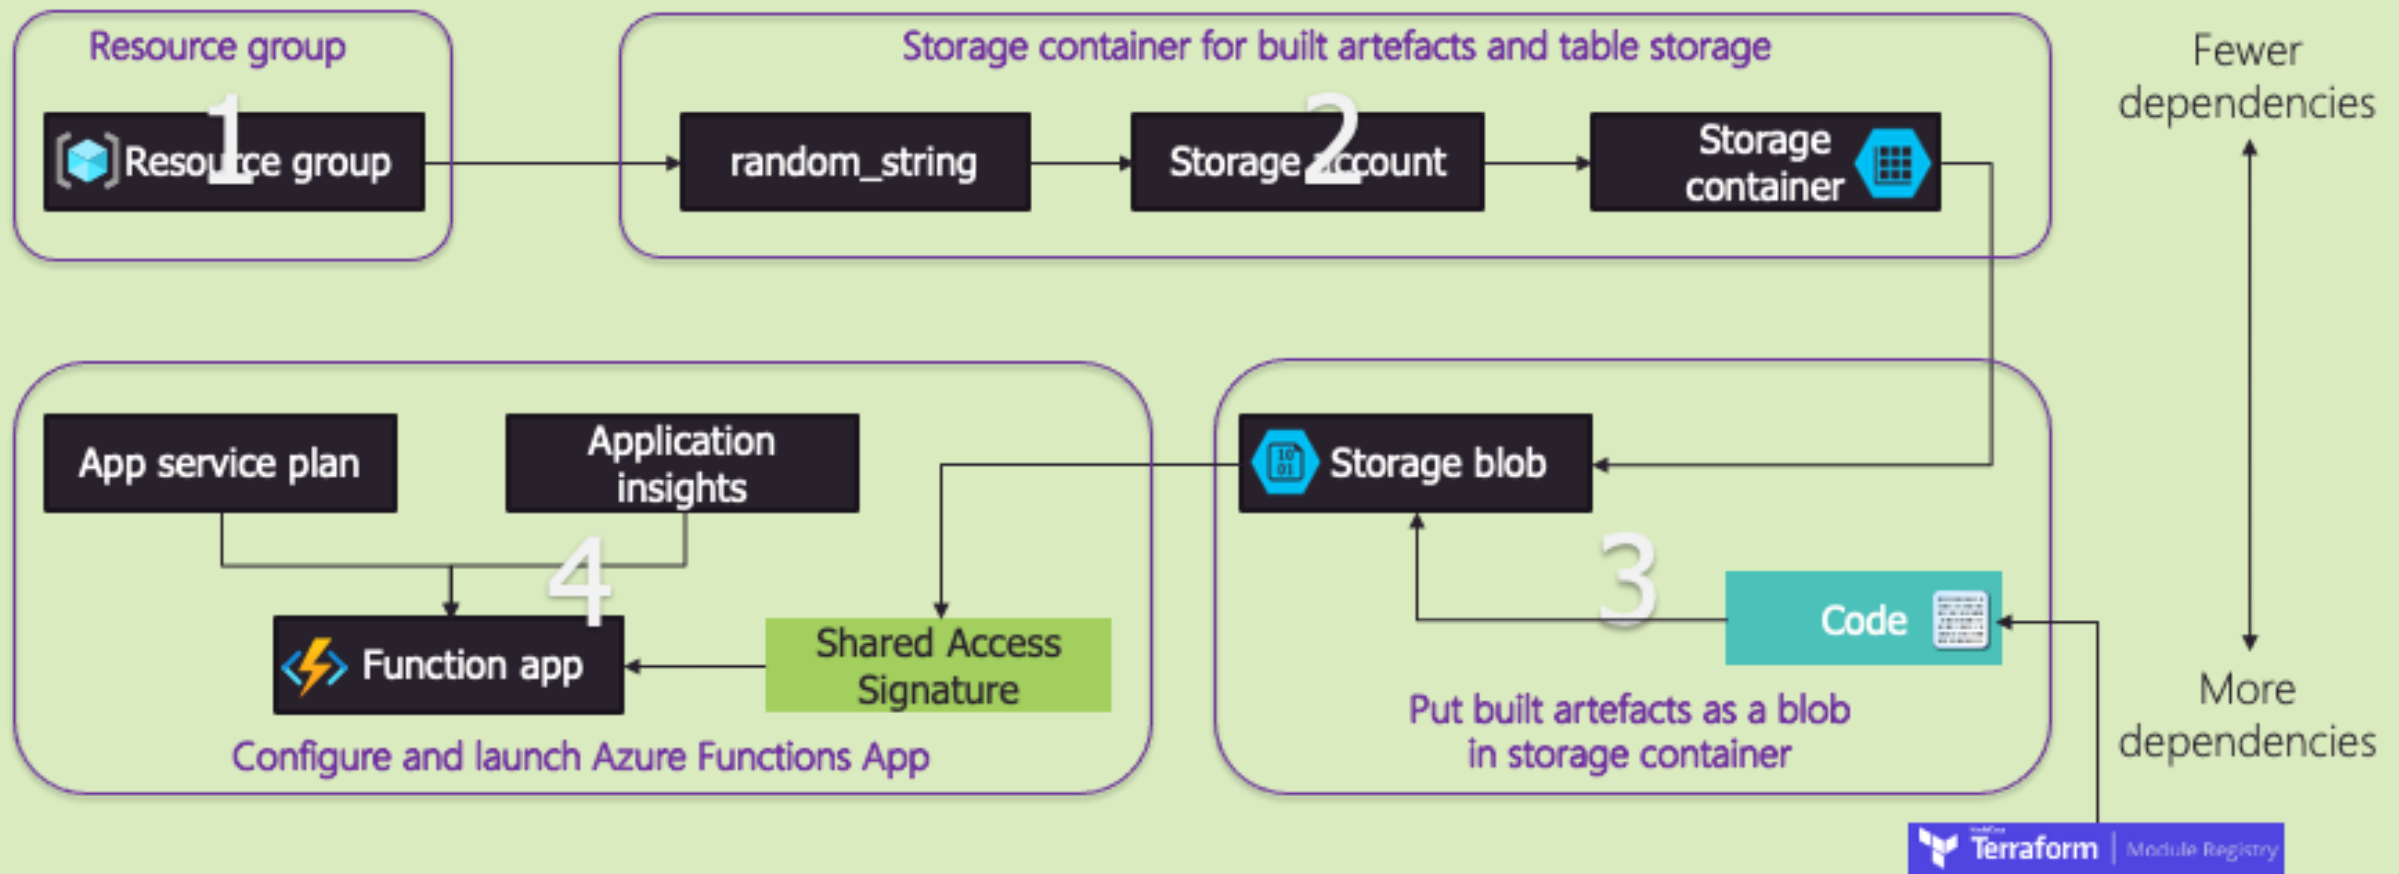


# Example 4.10 – Serverless

Resource

Module

Data source



# Note: Groups matter more than size

- | No more than a few hundred lines of code per TF file
- | Grouping resources that belong together
- | Organize your code in a sensible manner (NSS\*)

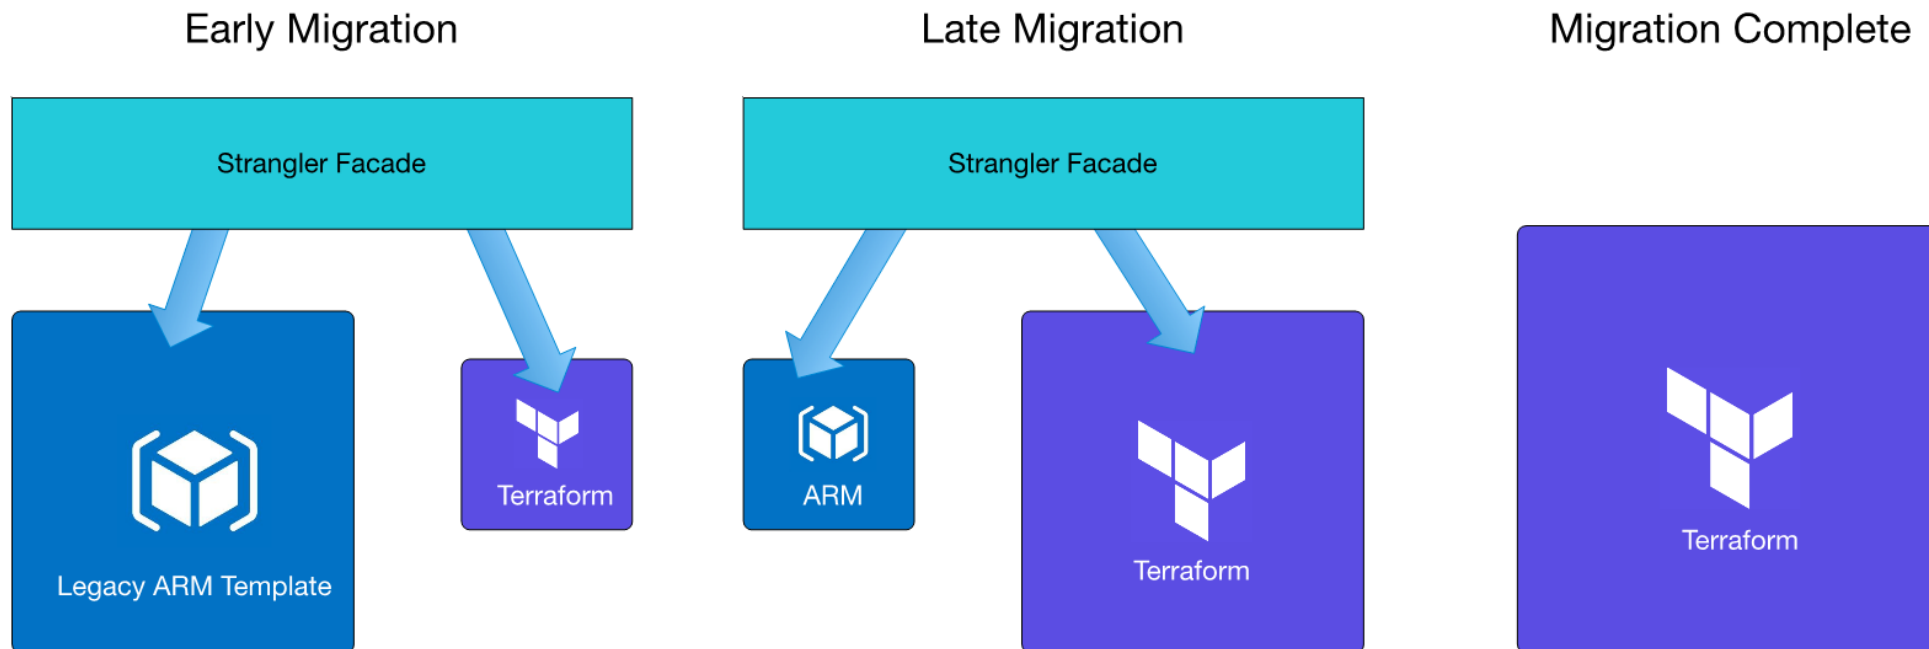
# ARM and Terraform

- | Legacy use cases where ARM is still useful
  - | Deploying resources that aren't yet supported by Terraform
  - | Migrating legacy ARM code to Terraform
  - | Generating configuration code

```
resource "azurerm_template_deployment" "extension" {  
  name = "extension"  
  resource_group_name = azurerm_resource_group.rg-app.name  
  template_body = file("ARM_siteExtension.json")  
  
  parameters = {  
    appserviceName = azurerm_app_service.app.name  
    extensionName = "AspNetCoreRuntime.2.2.x64"  
    extensionVersion = "2.2.0-preview3-35497"  
  }  
  
  deployment_mode = "Incremental"  
}
```

# Migrating from Legacy Code

## | Strangler façade





# Generating Configuration Code

The screenshot shows the Microsoft Azure portal interface. The left sidebar contains navigation options like 'Home', 'Dashboard', 'All services', 'FAVORITES', 'Resource groups', 'All resources', 'App Services', 'SQL databases', 'Azure Cosmos DB', 'Virtual machines', 'Load balancers', 'Storage accounts', 'Virtual networks', 'Azure Active Directory', 'Monitor', 'Advisor', 'Security Center', 'Cost Management + Billing', and 'Help + support'. The main content area is titled 'ballroominaction - Export template' and shows a list of resources under the 'Resources (11)' section. The 'Template' tab is selected, displaying a JSON configuration code for the template. The code includes parameters for 'sites\_ballroominaction\_23sr1wf\_name', 'serverfarms\_ballroominaction\_23sr1wf\_name', 'components\_ballroominaction\_23sr1wf\_name', 'storageAccounts\_23sr1wfme7vb8eatlr4ly\_name', and 'actionGroups\_Application%20Insights%20Smart%20Detection\_name'.

```
1 {
2   "$schema": "https://schema.management.azure.com/schemas/2015-01-01/deploymentTemplate.json#",
3   "contentVersion": "1.0.0.0",
4   "parameters": {
5     "sites_ballroominaction_23sr1wf_name": {
6       "defaultValue": "ballroominaction-23sr1wf",
7       "type": "String"
8     },
9     "serverfarms_ballroominaction_23sr1wf_name": {
10      "defaultValue": "ballroominaction-23sr1wf",
11      "type": "String"
12    },
13    "components_ballroominaction_23sr1wf_name": {
14      "defaultValue": "ballroominaction-23sr1wf",
15      "type": "String"
16    },
17    "storageAccounts_23sr1wfme7vb8eatlr4ly_name": {
18      "defaultValue": "23sr1wfme7vb8eatlr4ly",
19      "type": "String"
20    },
21    "actionGroups_Application%20Insights%20Smart%20Detection_name": {
22      "defaultValue": "Application%20Insights%20Smart%20Detection",
23      "type": "String"
24    }
25  },
26  "variables": {},
27  "resources": [
28    {
29      "type": "microsoft.insights/actionGroups",
```

# Summary

- | Terraform orchestrates serverless deployments with ease.
- | Code organization is paramount when designing Terraform modules.
- | Any files that are in a Terraform module are downloaded as part of terraform init or terraform get.
- | Azure Resource Manager (ARM) is an interesting technology that can be combined with Terraform to patch holes in Terraform

# 5. Collaborating with Terraform



# Content



Developing a remote backend module



Publishing modules via GitHub and the Terraform Module Registry



Switching between workspaces with the greatest of ease

# Terraform states

| Why not version control ?

- | Manual error
- | Locking
- | Secrets

| Remote state = backend

# Backend

- | Definition

- | Tasks:

- | Synchronize access to state files via locking
- | Store sensitive information securely
- | Keep a history of all state file revisions
- | Override CLI operations

- | Standard vs Enhanced backend


- | Access key can be used to increase "security"

# Exercise 5.01 - Azure Storage as Terraform backend

Check README file

# Workspace

- | Handle different subsets of the infrastructure
- | Handle more than one state file for the same configuration
- | There is always a “default” workspace
- | Every workspace has its own variables definition






## Exercise 5.02 - Workspaces

Check README file

# Summary

- | Remote backend is probably the best option for collaboration.
- | Workspaces allow you to deploy to multiple environments. The configuration code stays the same, the only thing that changes is the variables, and the state file.
- | Modules can be shared through a variety of means including: Azure Storage Container, GitHub repos, and the Terraform Module Registry. You can also implement your own Private Module Registry, if you're feeling adventurous.

# 6. Continuous Integration/Deployment



# Content



Two-stage deployments for separating static and dynamic infrastructure



Dynamic blocks



Implicit vs. explicit providers and provisioners



Use Azure DevOps to deploy a terraform-based infrastructure


# CI/CD Pipeline



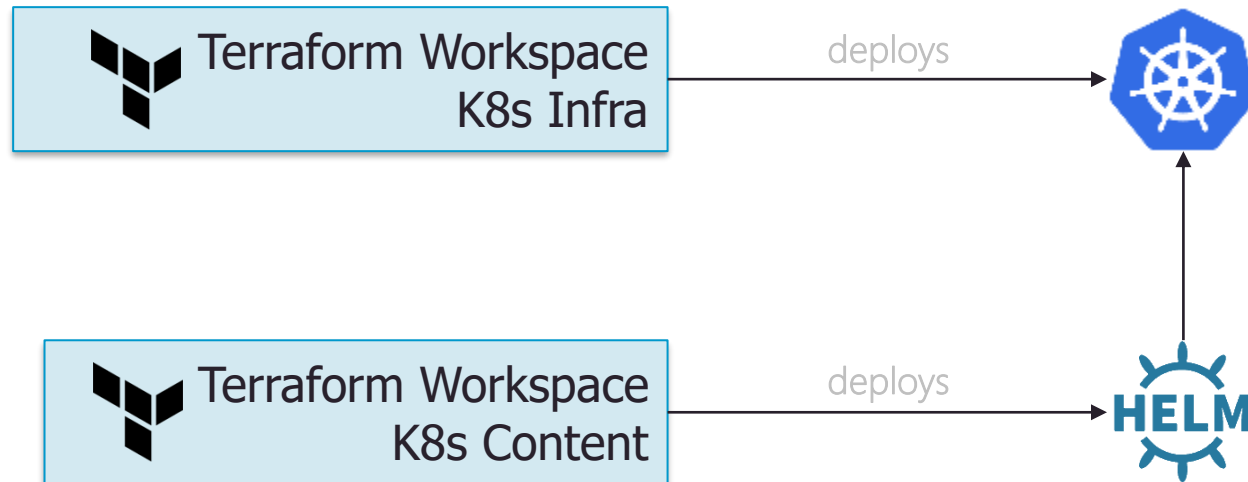
# Static Infrastructure

| So far,

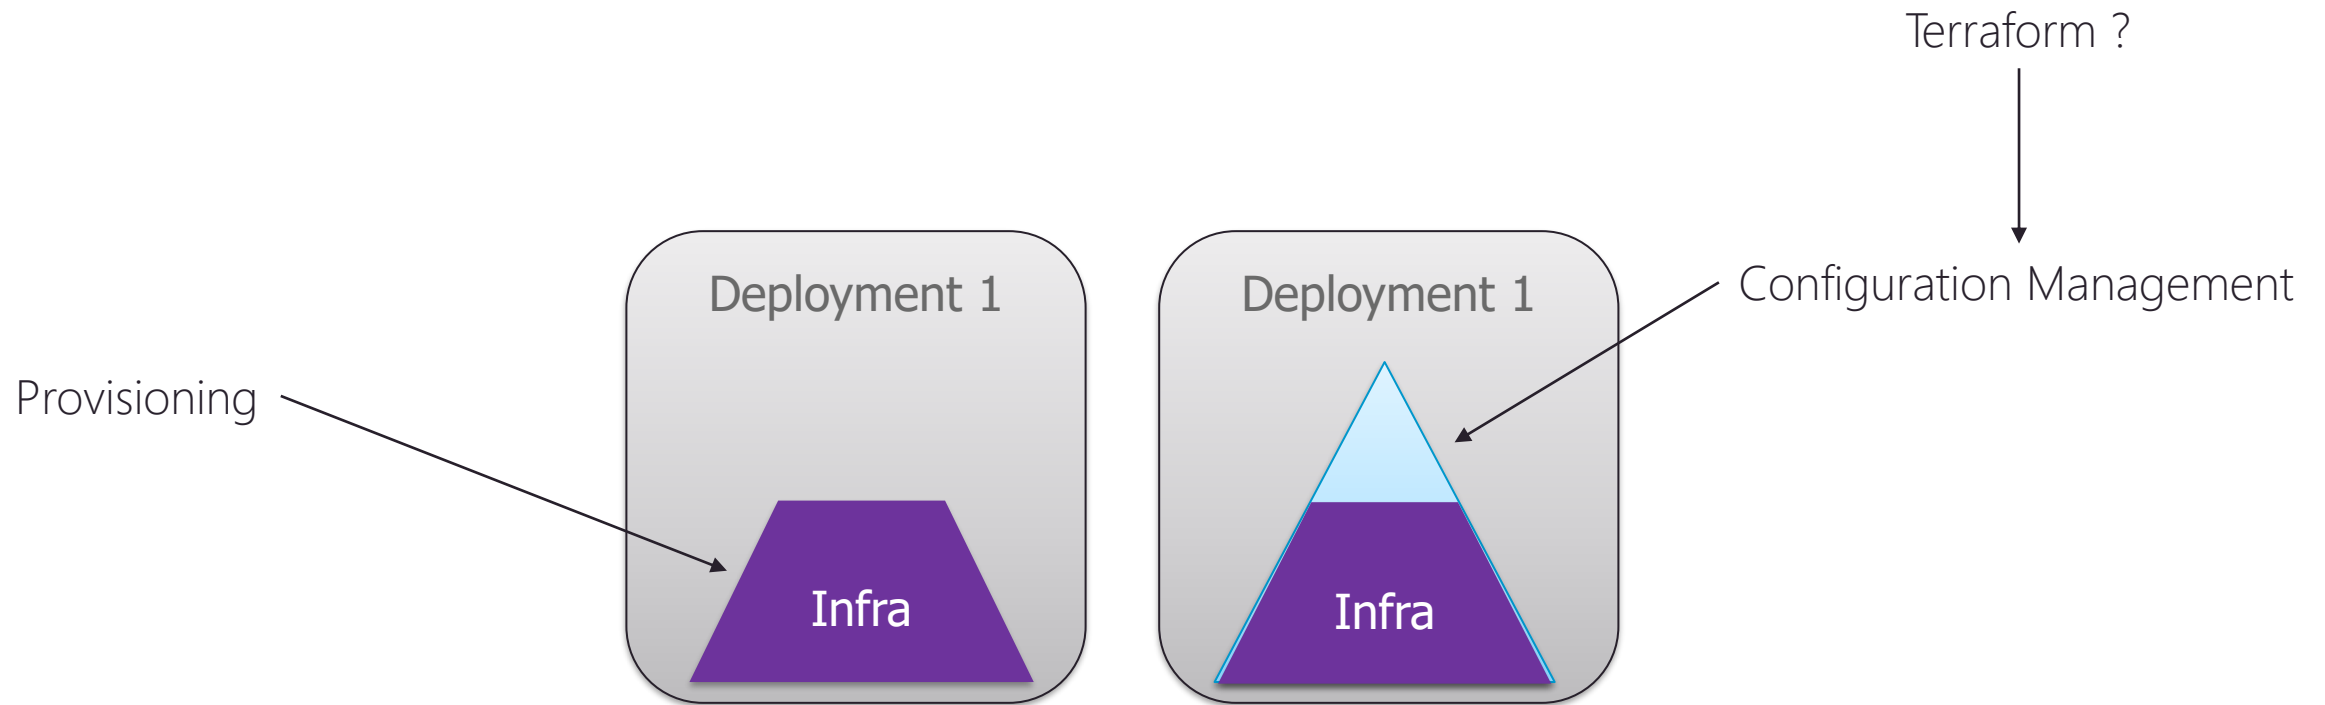
- | Terraform not well suited for managing frequent changes
- | All-in-one deployment



# What about multi-stage deployment ?



# Dynamic infrastructure

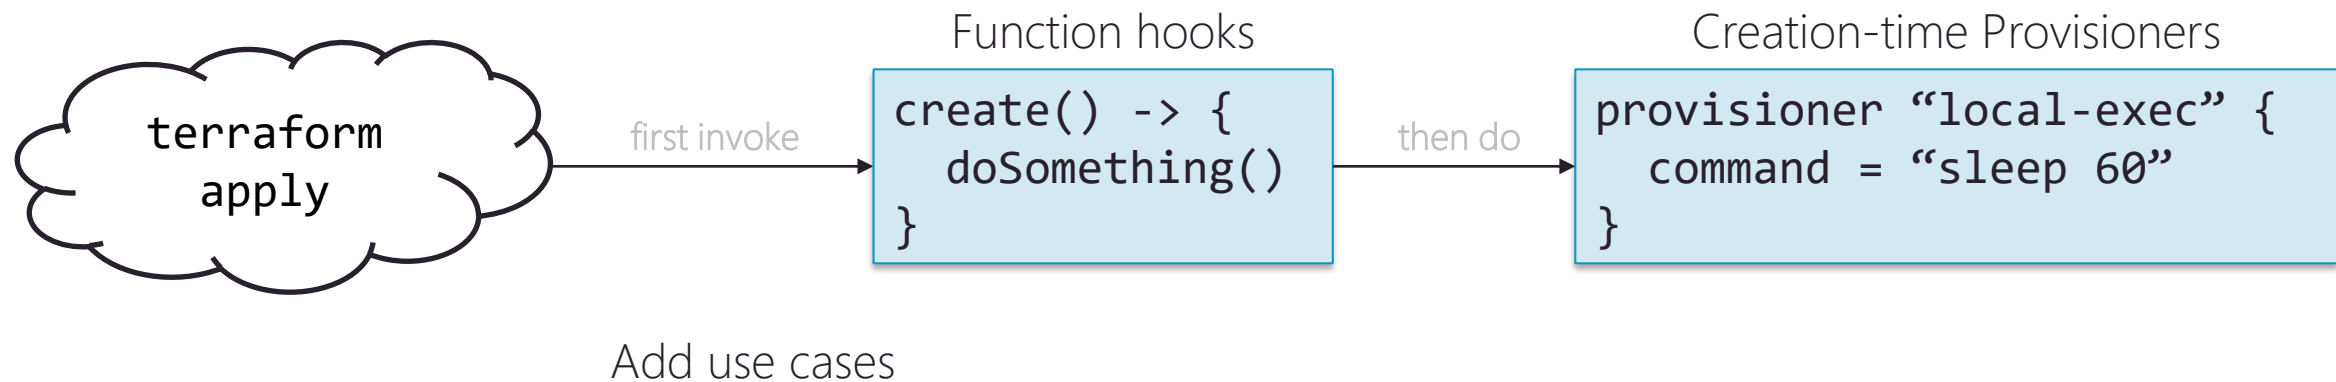




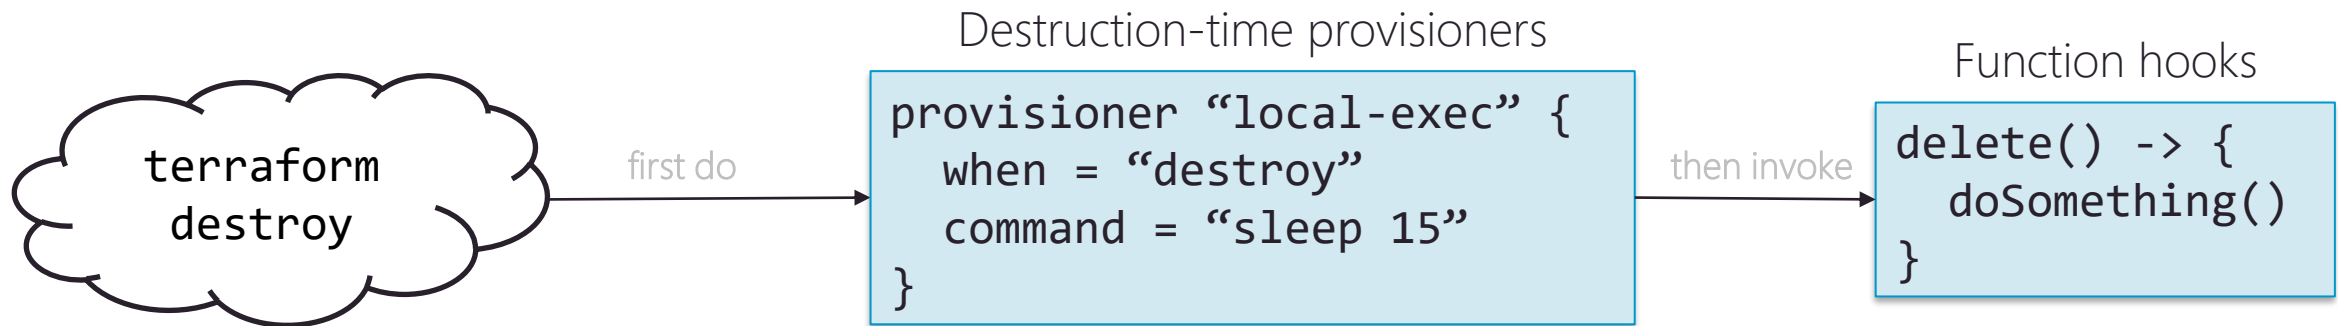
# Resource Provisioner

- | Race conditions - Timing is everything
- | Allow to execute scripts on local or remote machine during the creation or destruction of a resource.
- | Provisioner can be attached to any resource.
- | There are three categories for the default provisioners: file operations, script execution, and configuration management/provisioning.

# Creation-Time Provisioner



# Destruction-Time Provisioner



# Null resource

- | File provisioner
- | Script execution
  - | remote-exec
  - | local-exec
- | Configuration Management/Provisioning
  - | [chef](#), [habitat](#), [puppet](#), and [salt-masterless](#)

```
resource "null_resource" "upload" {  
  provisioner "file" {  
    ...  
  }  
}  
  
resource "null_resource" "azure-cli" {  
  provisioner "local-exec" {  
    command = "ssl-script.sh"  
  }  
}
```

# Example 6.01 – Null resource & provisioner

# Remote-exec example

```
provisioner "remote-exec" {  
  inline = [  
    "echo \"nameserver 8.8.8.8\" | sudo tee -a /etc/resolv.conf",  
    "sudo yum update -y",  
    "sudo yum install epel-release -y",  
    "sudo yum install puppet-agent -y",  
    "sudo /opt/puppetlabs/bin/puppet agent --version"  
  ]  
  connection {  
    user = "centos"  
  }  
}
```

# Dynamic Blocks

- | Use to dynamically create a nested configuration block;
- | Can *only* be used within other blocks, and *only* when the use of repeatable configuration blocks is supported;
- | See them as *for* expressions

```
dynamic "security_rule" {  
  for_each = var.ngs_rules  
  content {  
    name = security_rule.value["name"]  
    priority = security_rule.value["priority"]  
    ...  
  }  
}
```

The diagram illustrates the components of a dynamic block configuration. Annotations include:

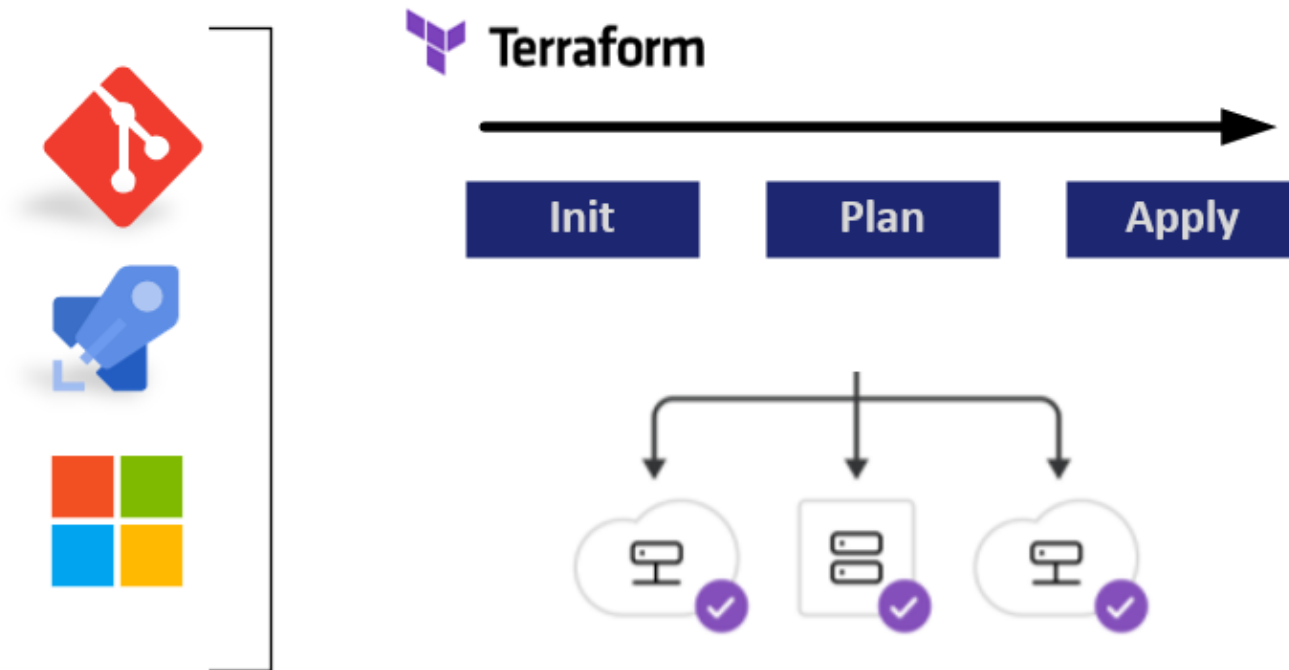
- Name of the block**: A bracket points to the `"security_rule"` string in the `dynamic` statement.
- Complex value to iterate**: An arrow points to the `var.ngs_rules` assignment in the `for_each` line.
- Current value accessor**: A bracket points to the `security_rule.value["..."]` syntax used within the `content` block.

## Exercise 6.02 – Dynamic block

Check README file



# Azure DevOps – CI/CD



3 “requirements”:

| Backend

| Azure DevOps Project

| Azure Service Principal

# Exercise 6.03 – Azure DevOps & Terraform

# Exercise 6.04 – Dynamic backend

# Summary

- | CI/CD pipeline can easily be supported by Azure DevOps
- | Resource provisioner complement the Terraform runtime.
- | Dynamic block is not often used by can speed up the elaboration of scripts and facilitate their reading
- | Combining Azure DevOps with Terraform increase the robustness of your CI/CD methodology

## 7. Deployment



# Content



Customizing resource lifecycles with the `create_before_destroy` flag



Performing Blue/Green deployments with Terraform

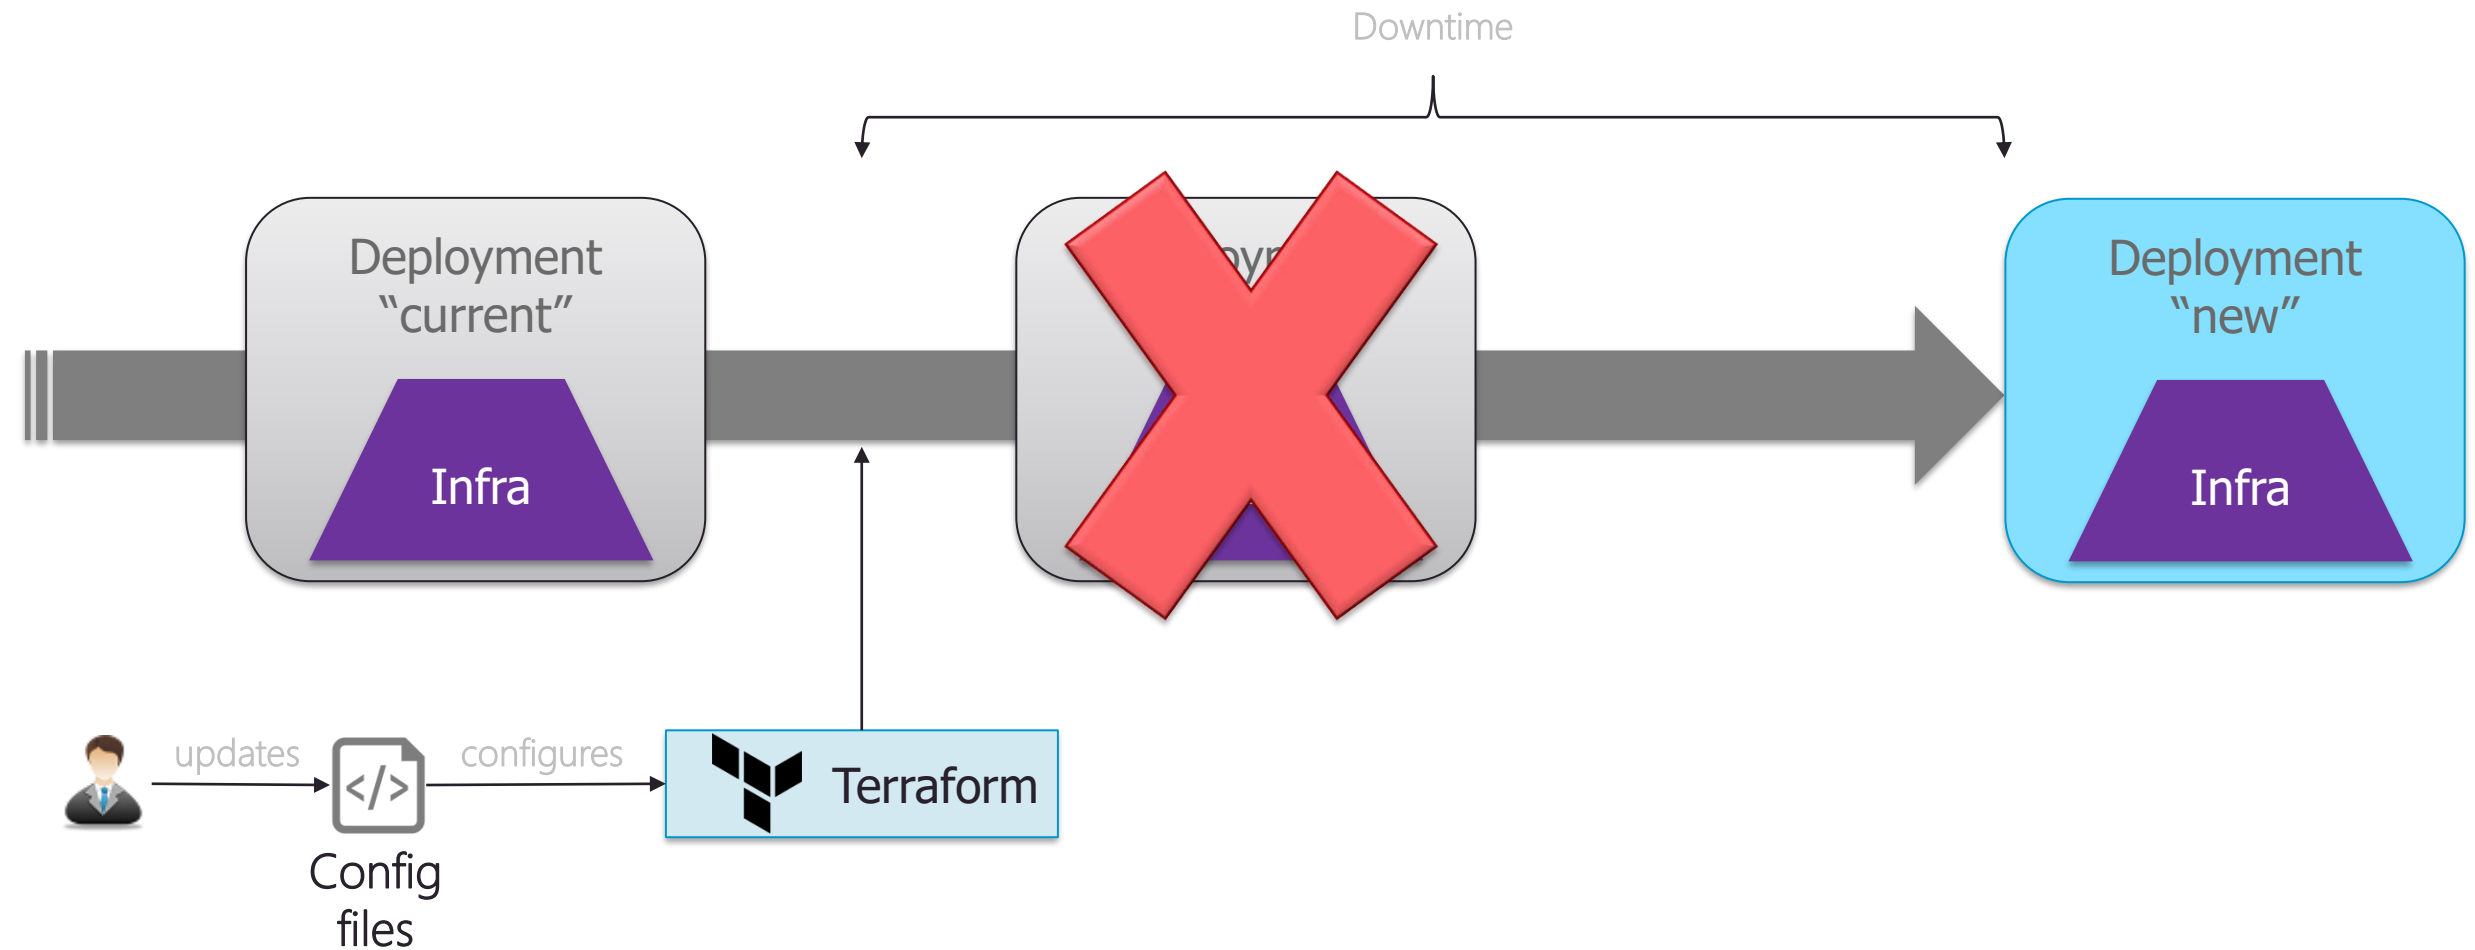


Installing software on virtual machines with `remote-exec` provisioners

# Zero Downtime Deployment

- | Definition: Practice of keeping services always running and available
- | Terraform's approaches:
  - | Use "create-before-destroy" meta-attribute
  - | Blue/Green deployments
  - | Responsibility transfer

# Lifecycle Customizations





# Lifecycle meta-argument

| "lifecycle" nested block in any resource

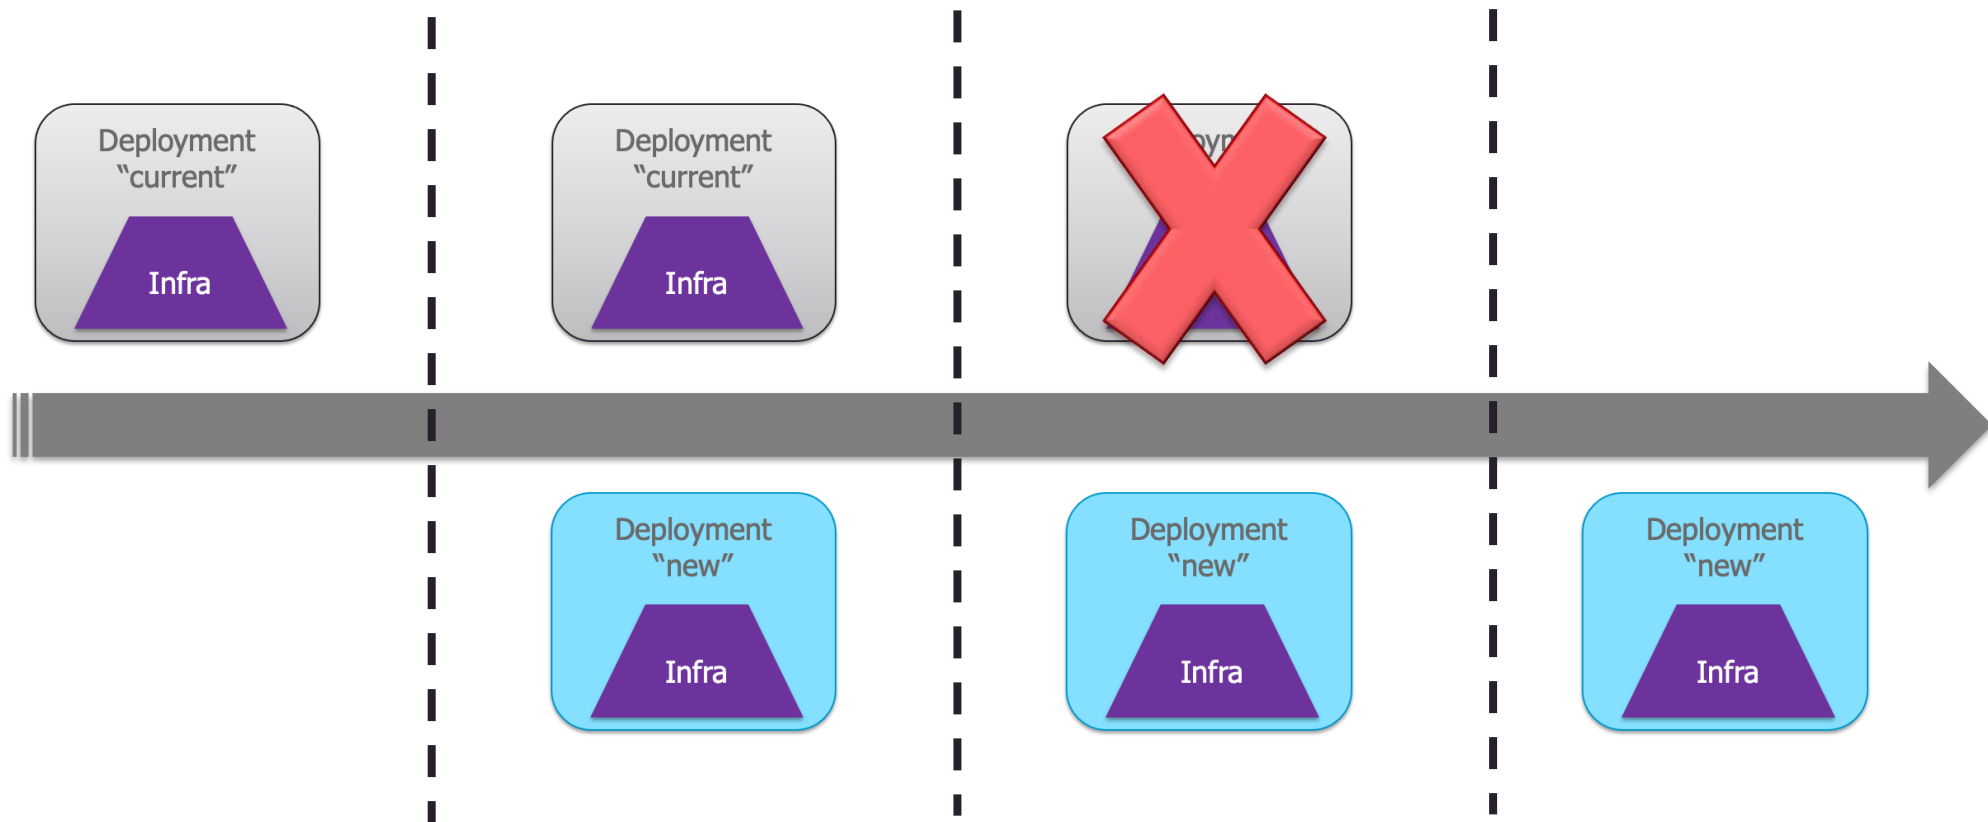
| Three flags:

- | "create\_before\_destroy" (bool)
- | "prevent\_destroy" (bool)
- | "ignore\_changes" (list of attribute names)

| Use these flags to override the default behaviour

```
resource "type" "name" {  
  lifecycle {  
    flag_one = ...  
    ...  
  }  
}
```

# Note: "create\_before\_destroy" flag



# Considerations about "create\_before\_destroy"


- | The "create\_before\_destroy" flag
  - | Can be confusing: affects the default behaviour of Terraform
  - | Is redundant: alternatives such as workspaces or modules
  - | Can create namespace collisions
  - | *"Force new" vs "updated in-place"*

## Exercise 7.01 – Lifecycle flag

Check README file

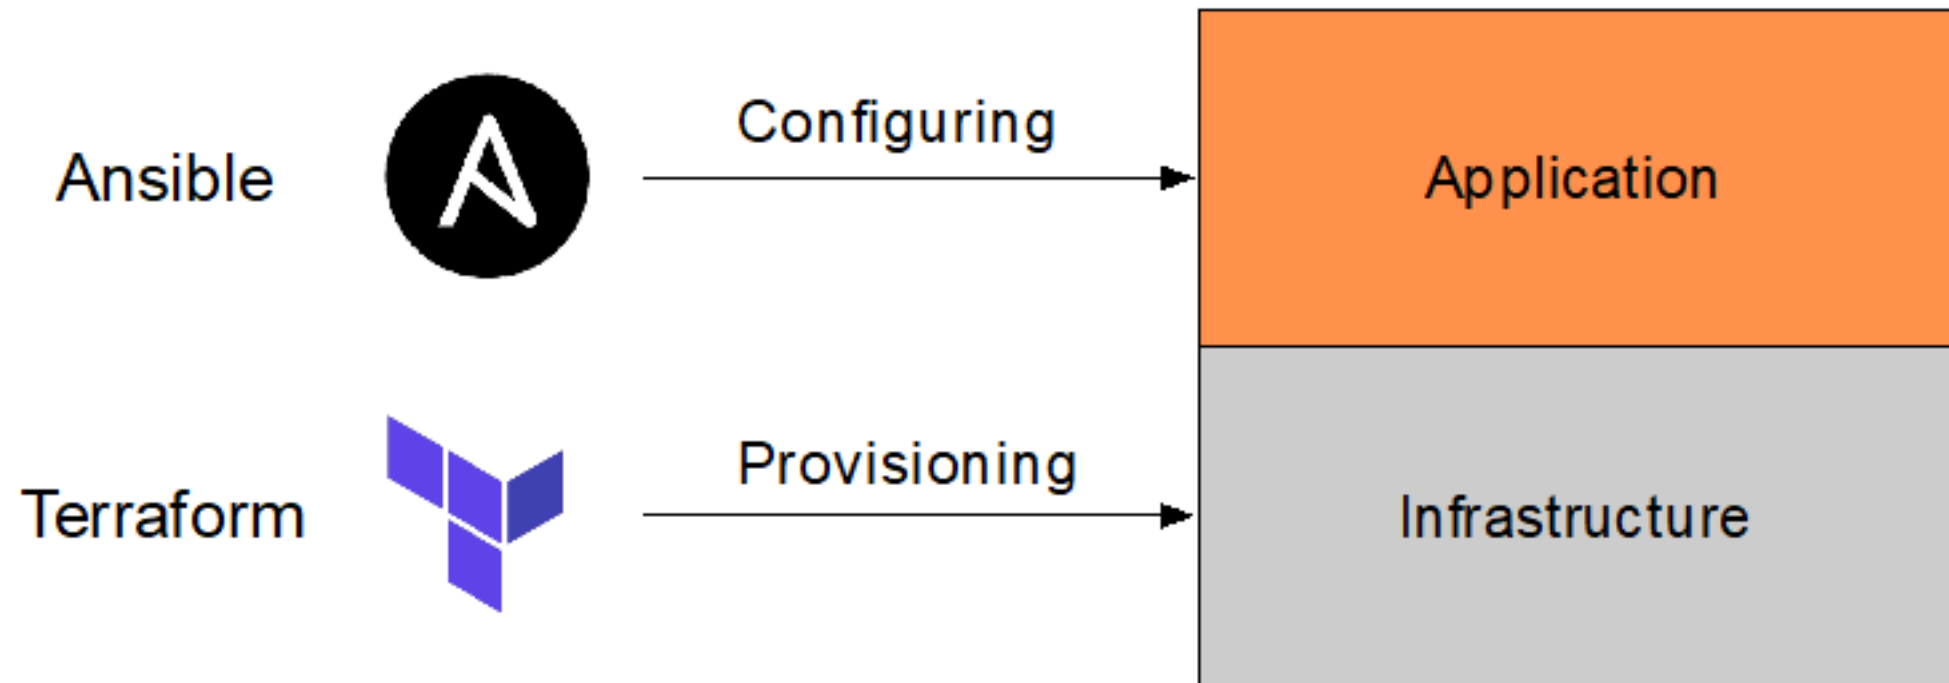
# Blue/Green Deployments

| The use of modules can help but probably not the best

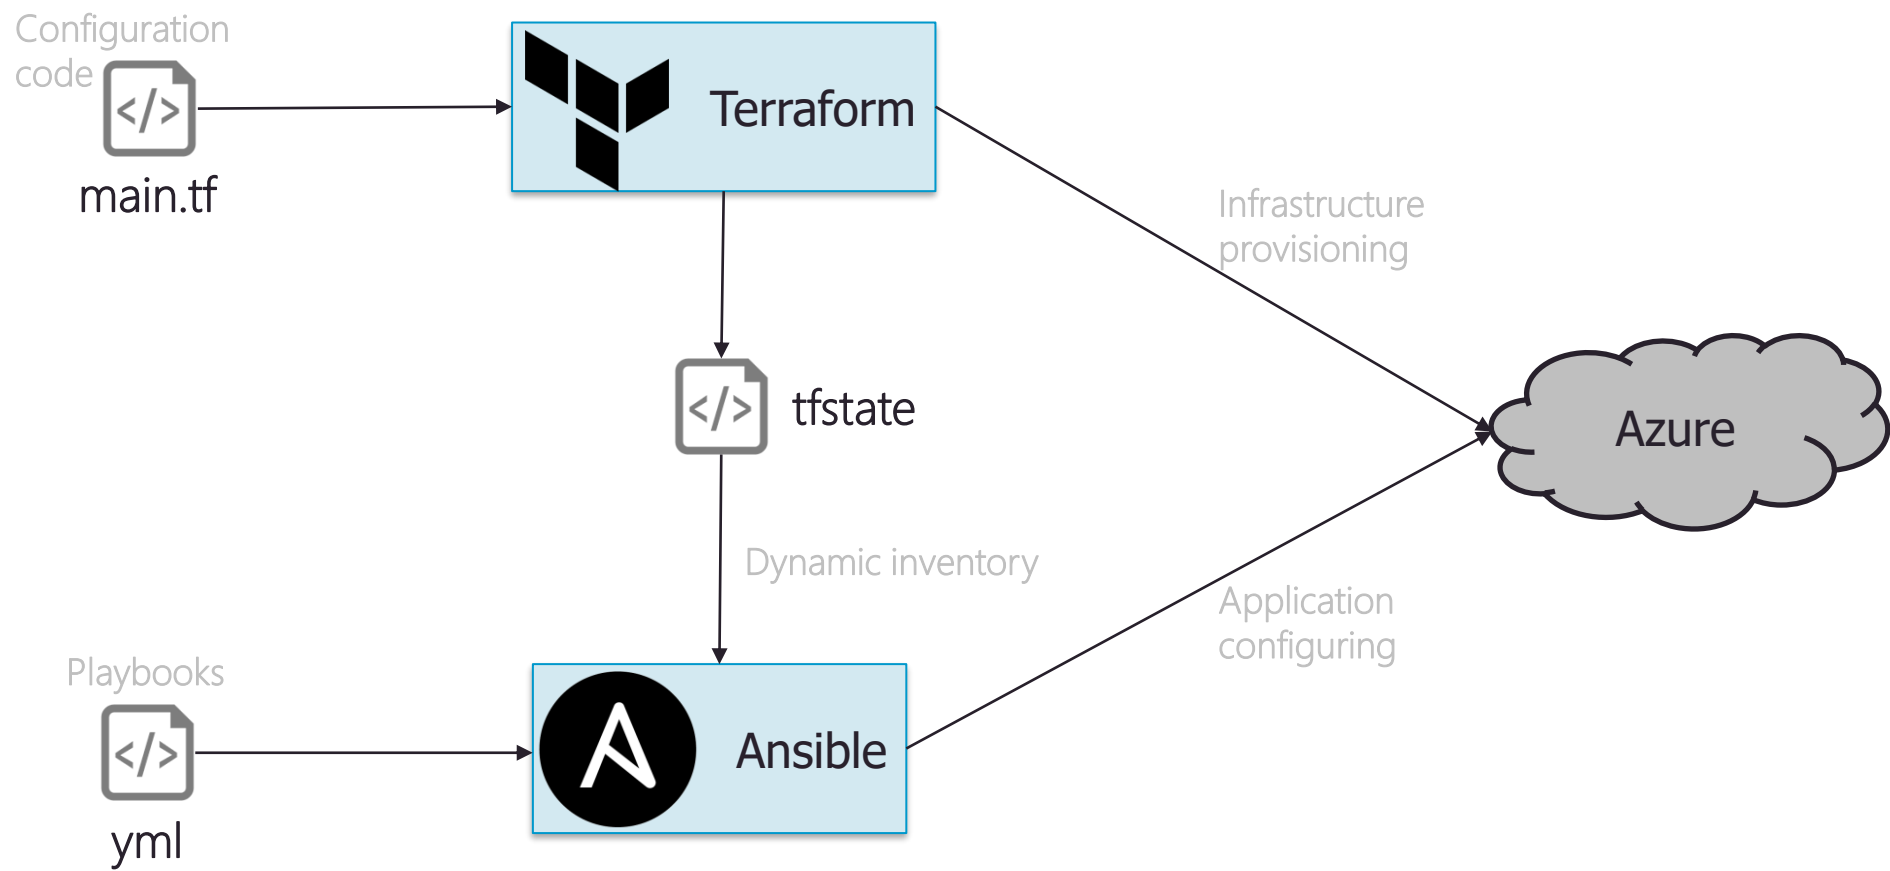


# Configuration Management

| Enables rapid software delivery onto existing server



# Terraform + Ansible



## Exercise 7.02 – Terraform & Ansible

Check README file



# Summary

- | Zero Downtime Deployment (ZDD) is the practice of keeping services always running and available to customers.
- | Performing Blue/Green deployments in Terraform is more a technique than a best practice.
- | Combination with configuration tool should always be considered

## 8. Testing and Refactoring



# Content



Resource tainting



Module expansion refactoring techniques



Migrating state with terraform mv and terraform state commands



Testing Infrastructure as Code with terraform-exec

# Refactoring

- | Definition: continuous improvement of the code design with marginal impact on behaviour
- | Refactoring goes further by strengthening maintainability, increasing extensibility and facilitating reusability

# Selective refactoring

- | **terraform taint** forces a resource to be destroyed and recreated on the next "apply" command
- | Tainted resources appears as such in the plan
- | Resources can be untainted
  
- | Use cases:
  - | Reset a resource to its initial state
  - | Force the rolling of security groups/keys
  - | Partial rebuild

# Exercise 8.01 & 8.02 – Rotate access keys

Check README file

# Securing Terraform state

- | Terraform does not treat attributes containing sensitive data any differently than it treats non-sensitive attribute.
- | Three methods for securing state files:
  - | Removing Unnecessary Secrets from Terraform State
  - | Least Privileged Access Control
  - | Encryption at Rest

## Exercise 8.03 – Azure KeyVault Secrets



# Refactoring – Modularizing Code

- | Typical deficiencies:
  - | Duplicated Code
  - | Name Collisions
  - | Inconsistency
- | The biggest refactoring improvement we can make it to put reusable code into modules.
- | Module expansions (only with TF 0.13)

## Exercise 8.04 – Modularizing Code

Check README file

# Terraform State Migration

- | Refactoring, and particularly software componentisation (modularizing code) implies to re-initialise the workspace.
- | Generating the plan will reveal that all resources will be destroyed and created during execution (apply).
- | There are 3 options:
  - | Manually editing the state file (not recommended)
  - | Moving stateful data with `terraform state mv`
  - | Deleting old resources with `terraform state rm` and reimporting with `terraform import`

# Moving resources

- | Moving an existing state from their current resource address to their final resource address.
- | `terraform state mv [options] SOURCE DESTINATION`

## Exercise 8.05 – Moving resources

Check README file

# Importing Resources


- | Migration can be performed through a DELETE + IMPORT process
- | `terraform state rm ADDRESS`
- | `Terraform import ADDRESS`

# Testing

| Testing should be at least done at three levels:

- | Unit testing
- | Integration testing
- | System testing

| terraform-exec (in GO language)



# Summary

- | “taint” is useful to rotate/refresh time sensitive resources
- | Module expansion should be favoured to flat module structure
  - | Move resources / modules
- | Unmanaged resources can be converted through import



Thank you !